

Neural Radiance Fields pt 3



made with
 nerfstudio

CS180/280A: Intro to Computer Vision and Computational Photography

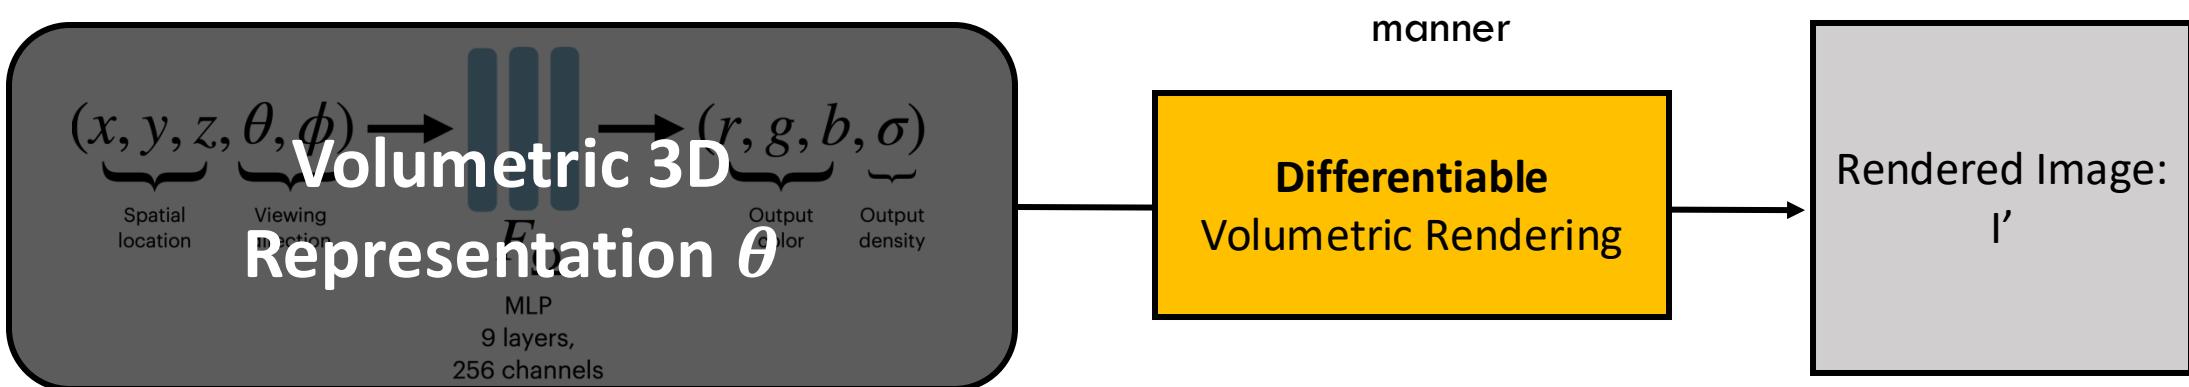
Angjoo Kanazawa and Alexei Efros

UC Berkeley Fall 2025

Lots of content from Noah Snavely and Ben
Mildenhall, Pratul Srinivasan, and Matt Tancik from
[ECCV 2022 Tutorial on Neural Volumetric Rendering
for Computer Vision](#)



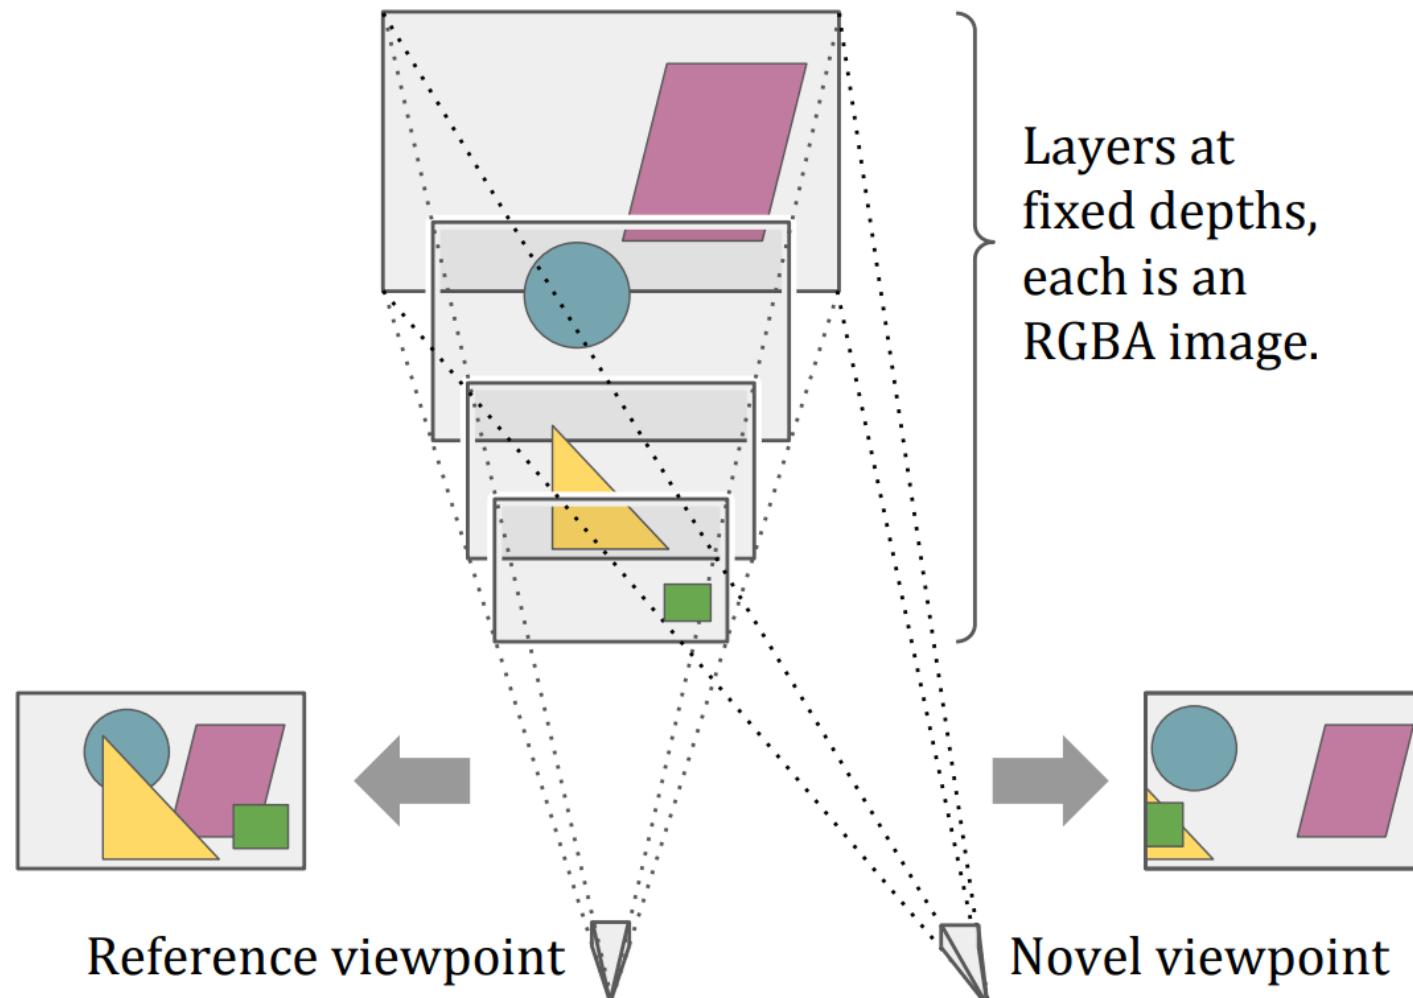
Where we are



“Training” Objective (aka Analysis-by-Synthesis):

$$\min_{\theta} \| \text{Rendered Image: } I' - \text{Observed Image: } I \|_2$$

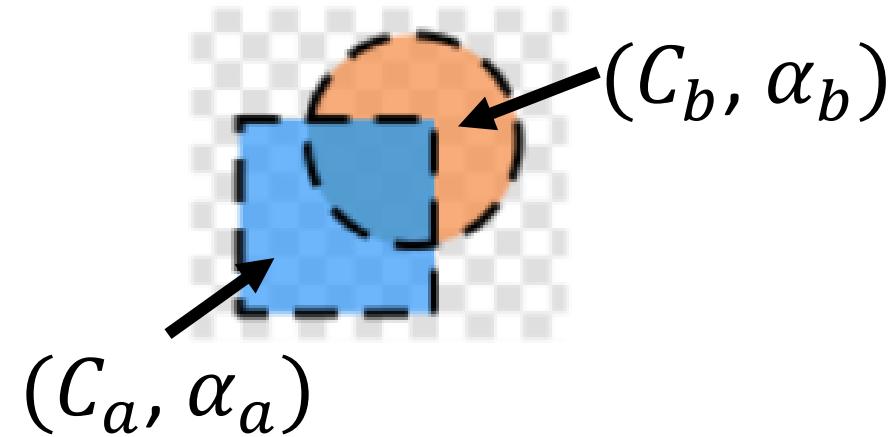
A Precursor: Multi-plane Images



Also called front-to-back compositing or “over” operation

Alpha Blending

for two image case, A and B, both partially transparent:

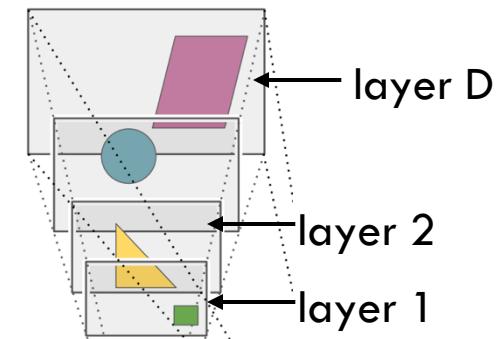


$$I = C_a \alpha_a + C_b \alpha_b (1 - \alpha_a)$$

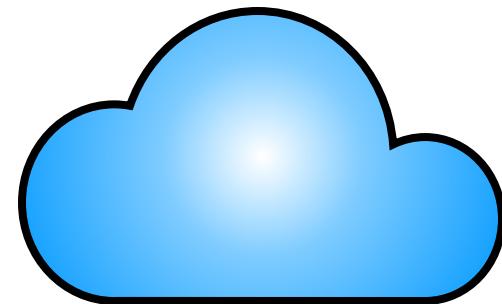
How much light is the previous layer letting through?

General D layer case:

$$I = \sum_{i=1}^D C_i \alpha_i \prod_{j=1}^{i-1} (1 - \alpha_j)$$

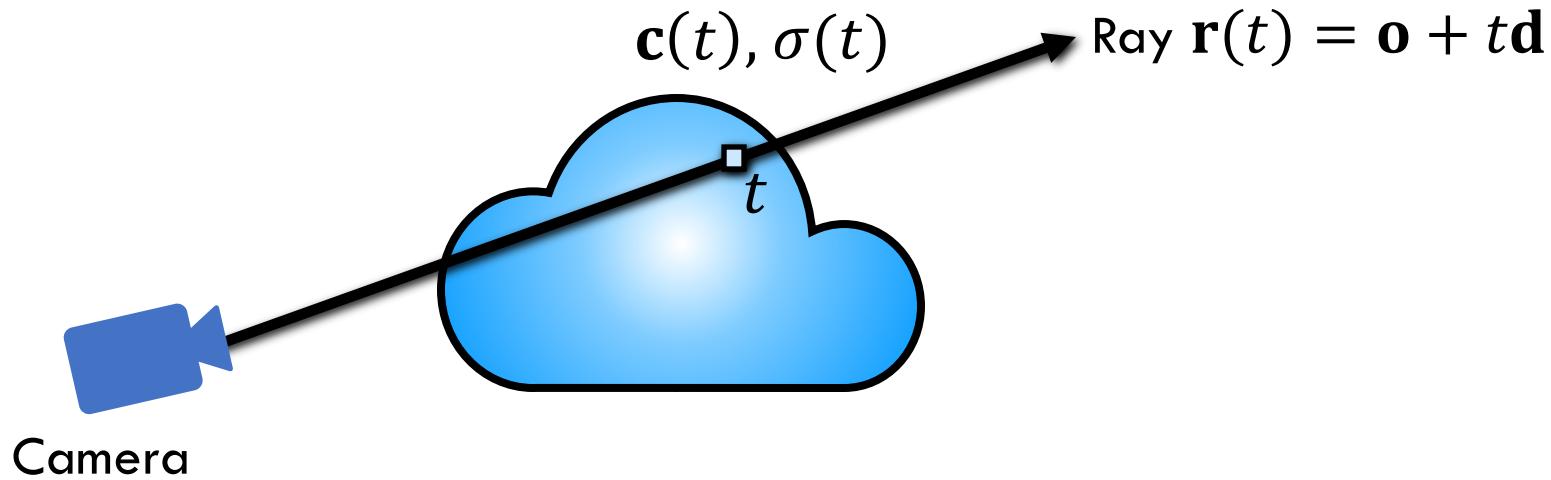


Volumetric formulation for NeRF



Scene is a cloud of tiny colored particles

Volumetric formulation for NeRF



at a point on the ray $\mathbf{r}(t)$, we can query color $\mathbf{c}(t)$ and density $\sigma(t)$

How to integrate all the info along the ray to get a color per ray?

Idea: Expected Color

- Pose probabilistically.
- Each point on the ray has a probability to be the first “hit” : $P[\text{first hit at } t]$
- Color per ray = Expected value of color with this probability of first “hit”

for a ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$:

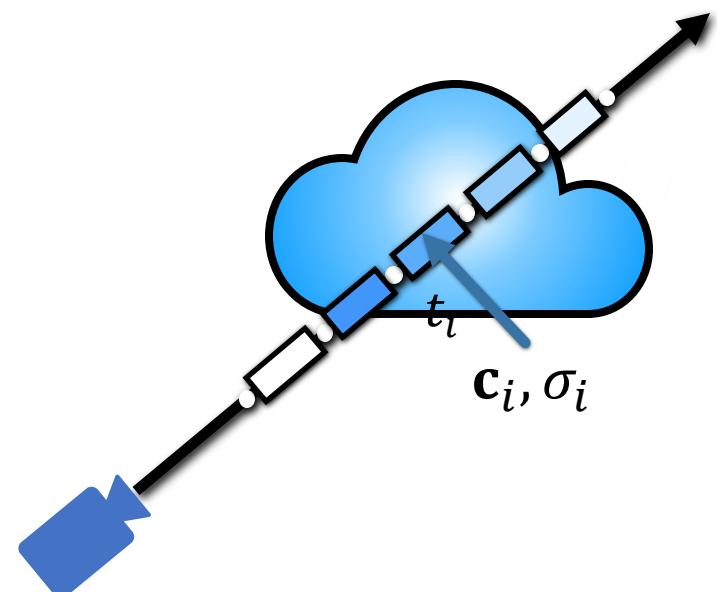
$$\mathbf{c}(\mathbf{r}) = \int_{t_0}^{t_1} P[\text{first hit at } t] \mathbf{c}(t) dt$$

$$\approx \sum_{t=0}^T P[\text{first hit at } t] \mathbf{c}(t)$$

$$\approx \sum_{t=0}^T w_t \mathbf{c}(t)$$

$$= \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$

$$\text{where } T_i = \prod_{j=1}^{i-1} (1 - \alpha_j) \quad \alpha_i = 1 - \exp(-\sigma_i \delta_i)$$



Differentiable Volumetric Rendering Formula

for a ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$:

$$\mathbf{c} \approx \sum_{i=1}^n w_i \mathbf{c}_i$$

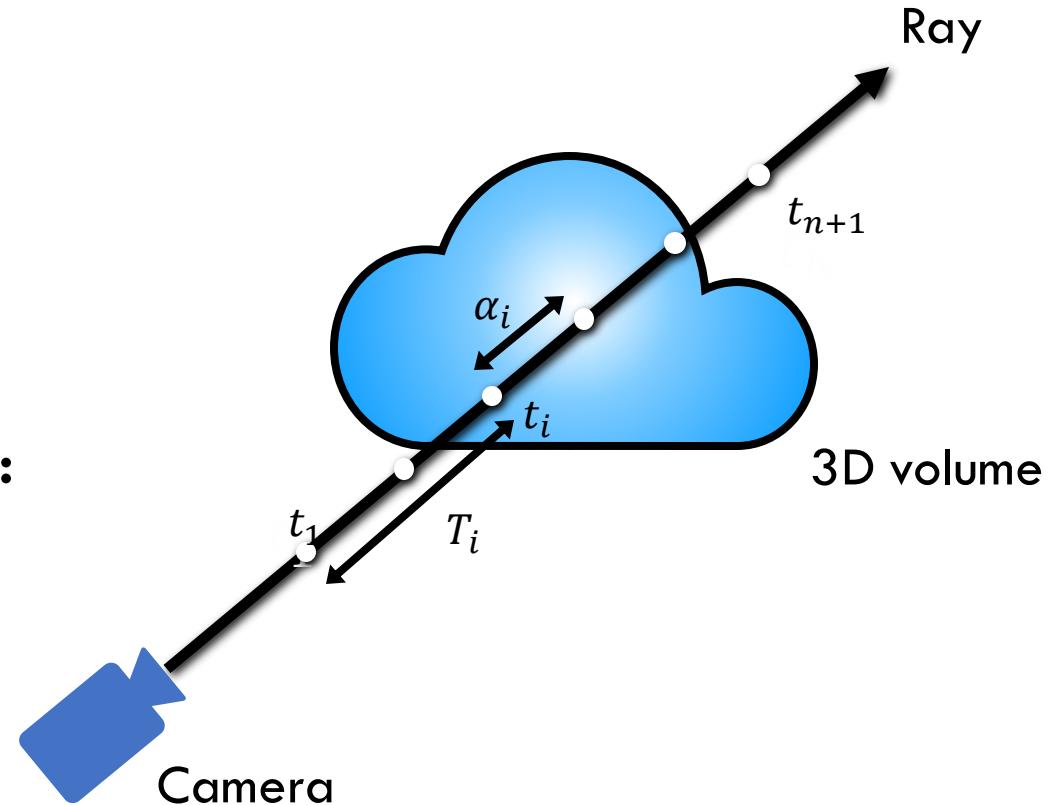
differentiable w.r.t. \mathbf{c}, σ

colors

weights

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$



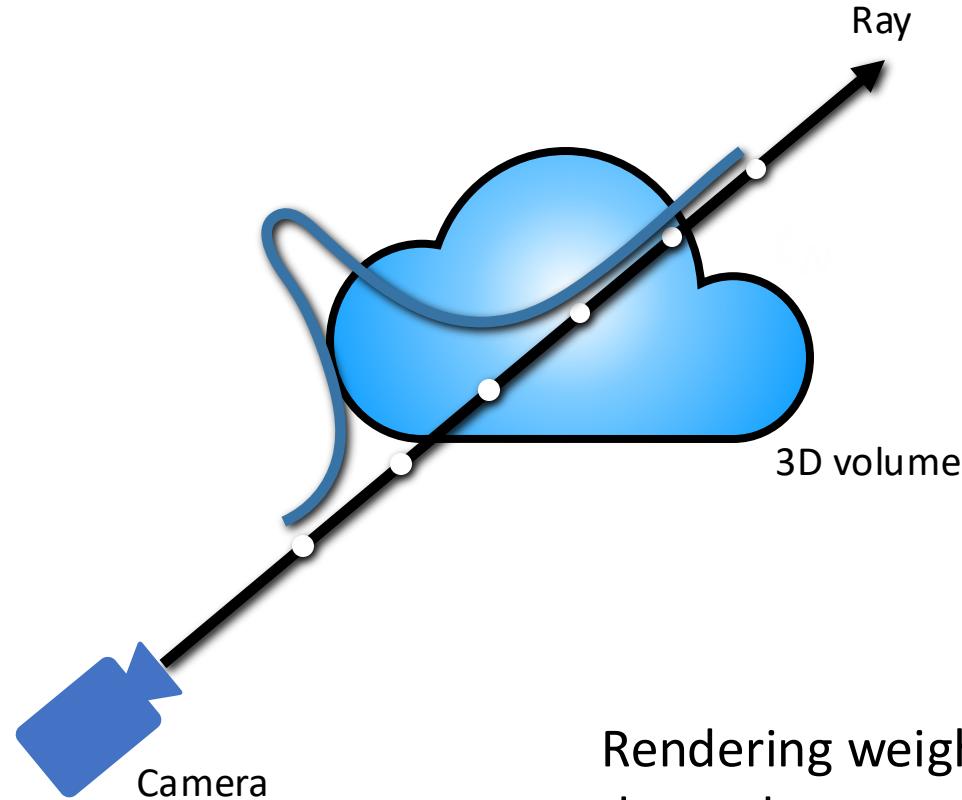
How much light is contributed by ray segment i :

$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$

Delta is the length of each segment: $\delta_i = t_{i+1} - t_i$

Visual intuition: rendering weights is specific to a ray

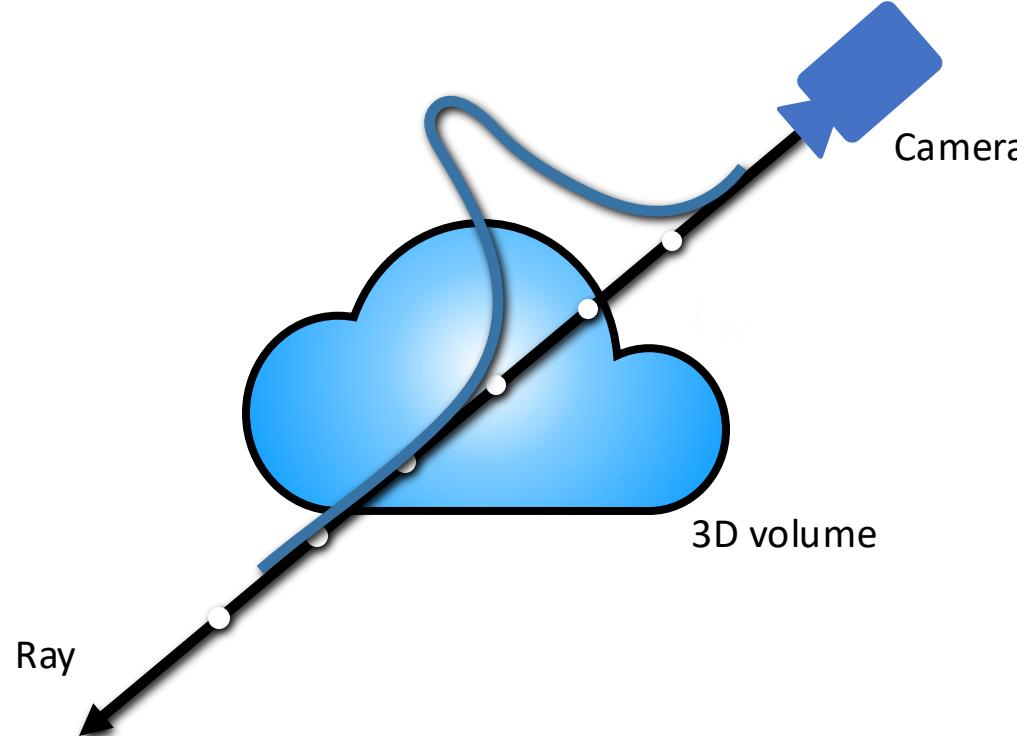
$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$



Rendering weights are not a 3D function — depends on ray, because of transmittance!

Visual intuition: rendering weights is specific to a ray

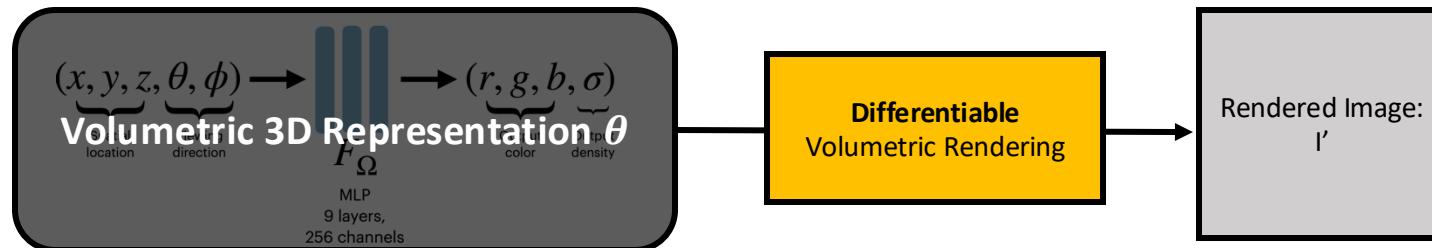
$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$



Rendering weights are not a 3D function — depends on ray, because of transmittance!

What's the point

- Remember, for each pixel or a ray we render a color with this formula based on the Volumetric 3D Representation
- We use this to supervise the 3D Representation (sigma, RGB volume)



“Training” Objective (aka Analysis-by-Synthesis):

$$\min_{\theta} \| \text{Rendered Image: } I' - \text{Observed Image: } I \|_2$$

Connection to alpha compositing

$$\text{Expected Color} = \sum_{i=1}^n T_i \mathbf{c}_i (1 - \exp(-\sigma_i \delta_i))$$

segment

opacity α_i

$$\text{Expected Color} = \sum_{i=1}^n T_i \mathbf{c}_i \alpha_i$$

where $T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$

$$= \prod_{j=1}^{i-1} (1 - \alpha_j)$$

$$\prod_i \exp(x_i) = \exp\left(\sum_i x_i\right)$$
$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$
$$1 - \alpha_i = -\exp(-\sigma_i \delta_i)$$

Summary

for a ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$:

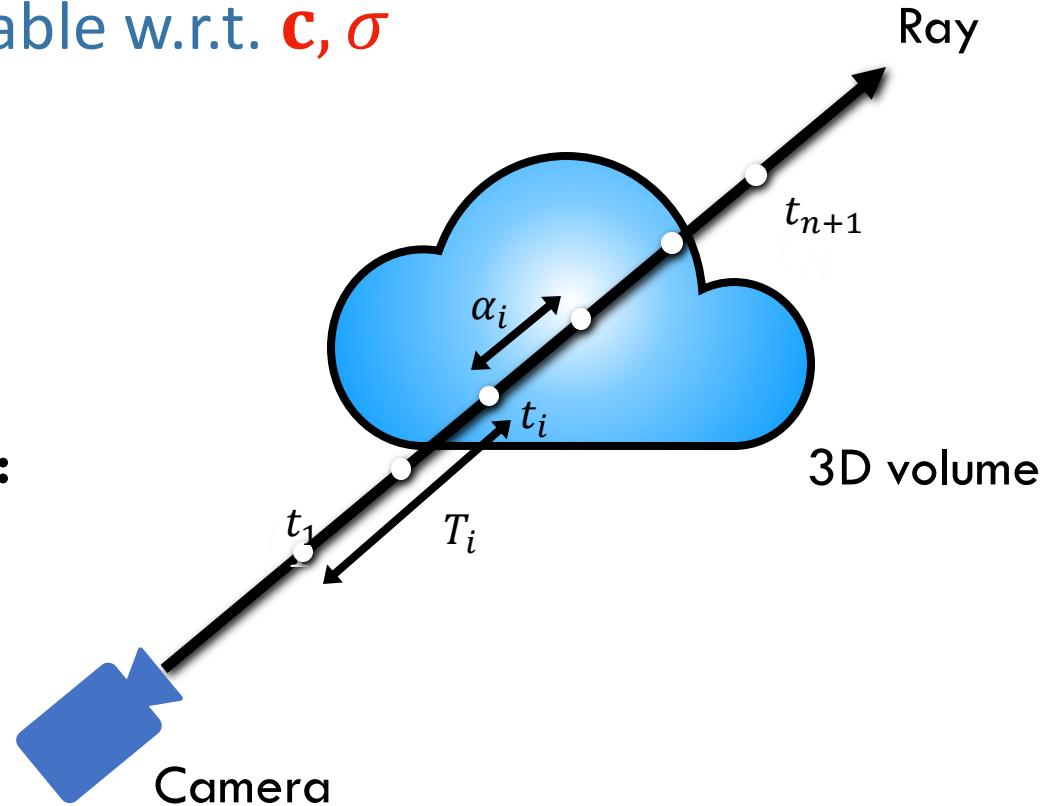
$$\mathbf{c} \approx \sum_{i=1}^n w_i \mathbf{c}_i = \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$

↑
weights colors

differentiable w.r.t. \mathbf{c}, σ

How much light is blocked earlier along ray:

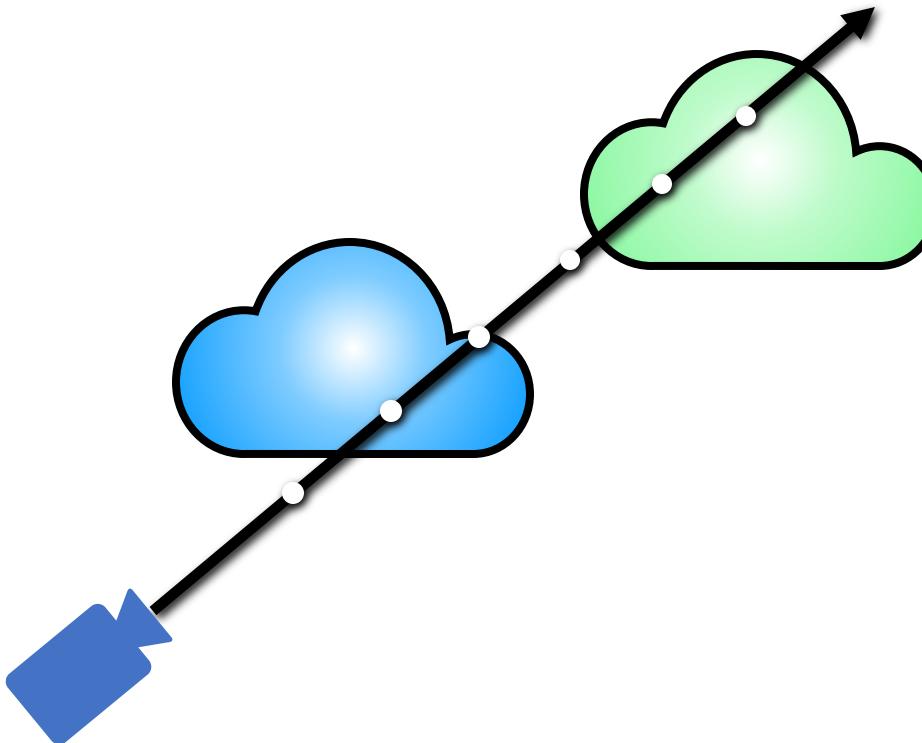
$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$



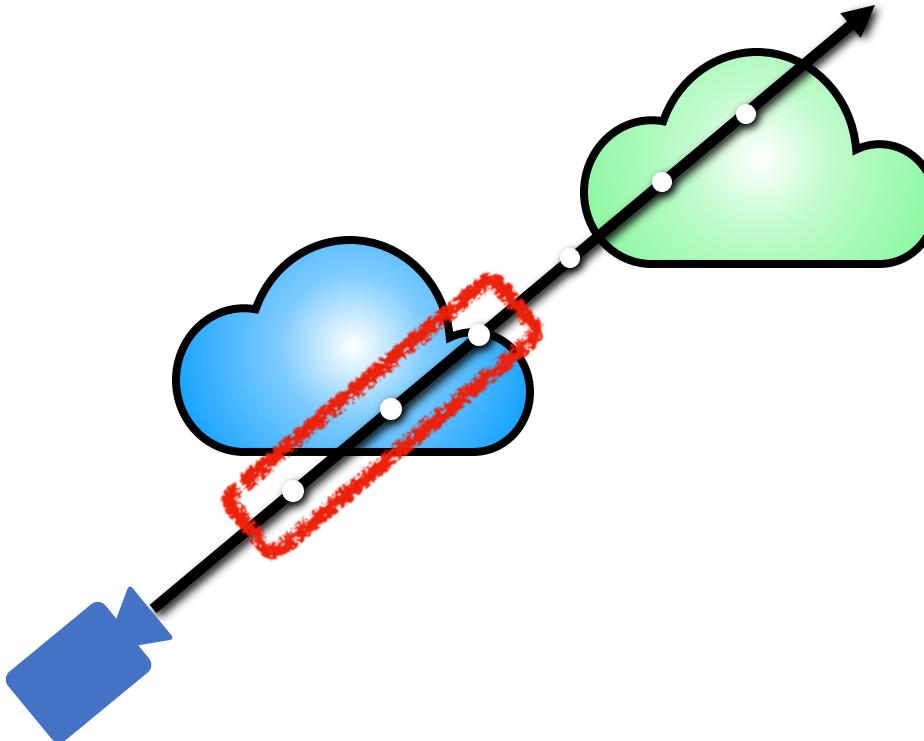
How much light is contributed by ray segment i :

$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$

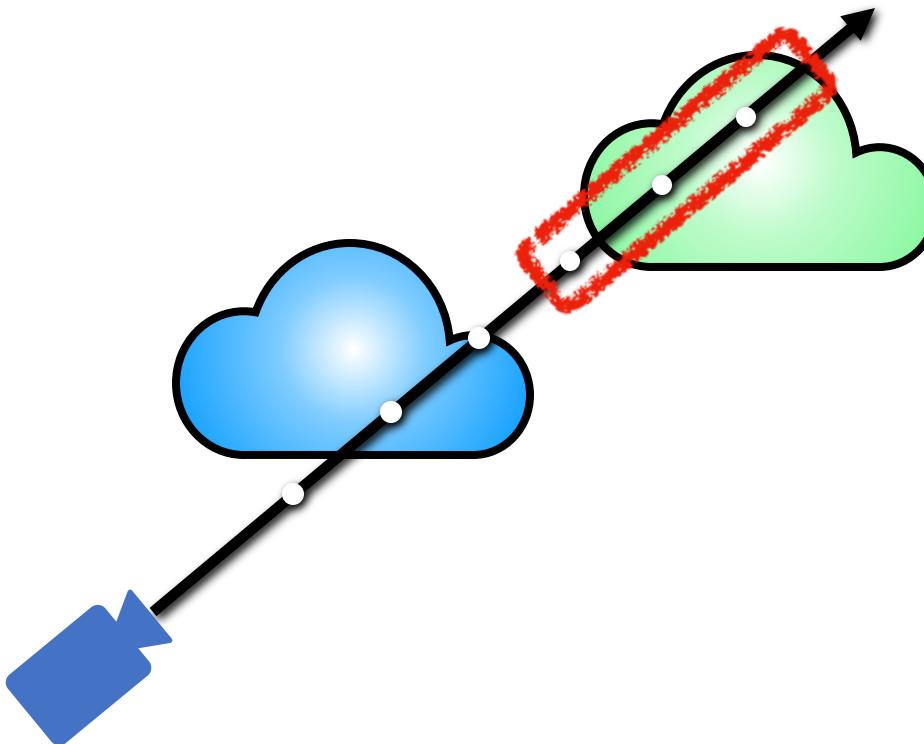
Alpha mattes and compositing



Alpha mattes and compositing



Alpha mattes and compositing



Alpha mattes and compositing



Mildenhall*, Srinivasan*, Tancik* et al 2020, NeRF

Poole et al 2022, DreamFusion

Tang et al 2022, Compressible-composable NeRF via Rank-residual Decomposition

Rendering weight PDF is important

Remember, expected color is equal to

$$\int T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_i T_i \alpha_i \mathbf{c}_i = \sum_i w_i \mathbf{c}_i$$

$T(t)\sigma(t)$ and $T_i \alpha_i$ are “rendering weights” — probability distribution along the ray (continuous and discrete, respectively)

You can also render entities other than color in 3D, for example it’s depth, or any other N-D vector \mathbf{v}_i

$$\text{Volume rendered "feature"} = \sum_i w_i \mathbf{v}_i$$

Rendering weight PDF is important — depth

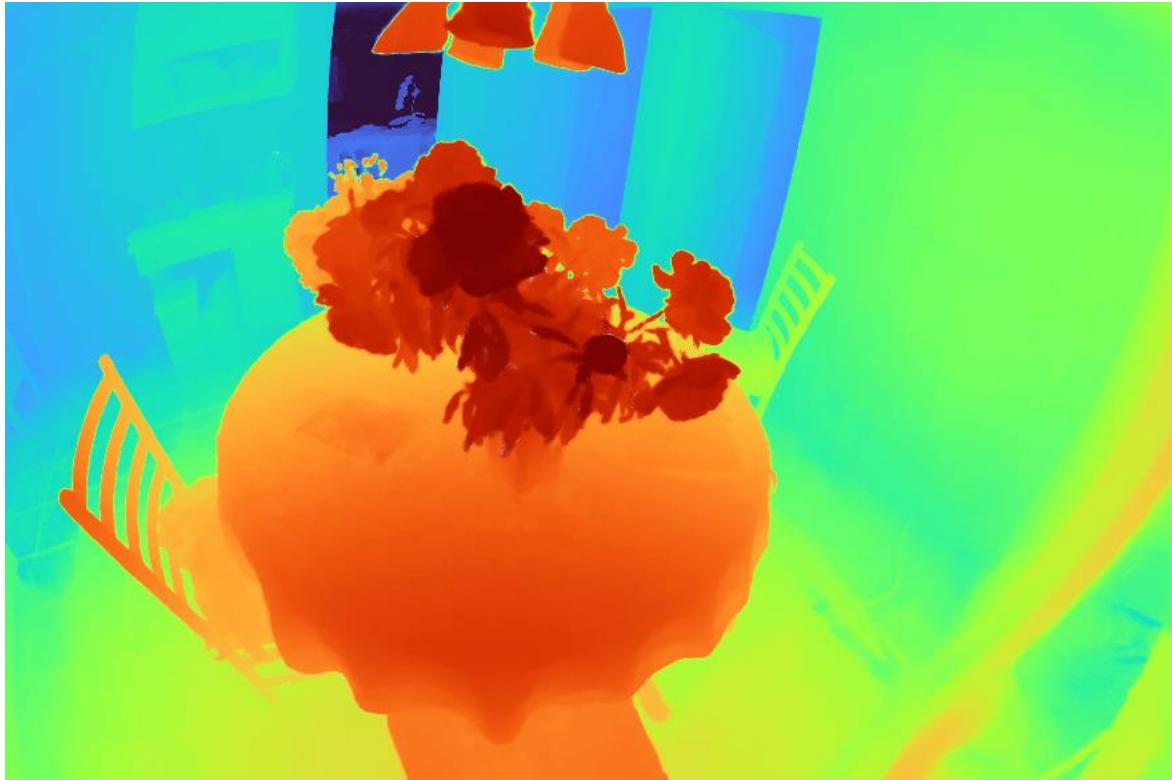
We can use this distribution to compute expectations for other quantities, e.g. “expected depth”:

$$\bar{t} = \sum_i T_i \alpha_i t_i$$

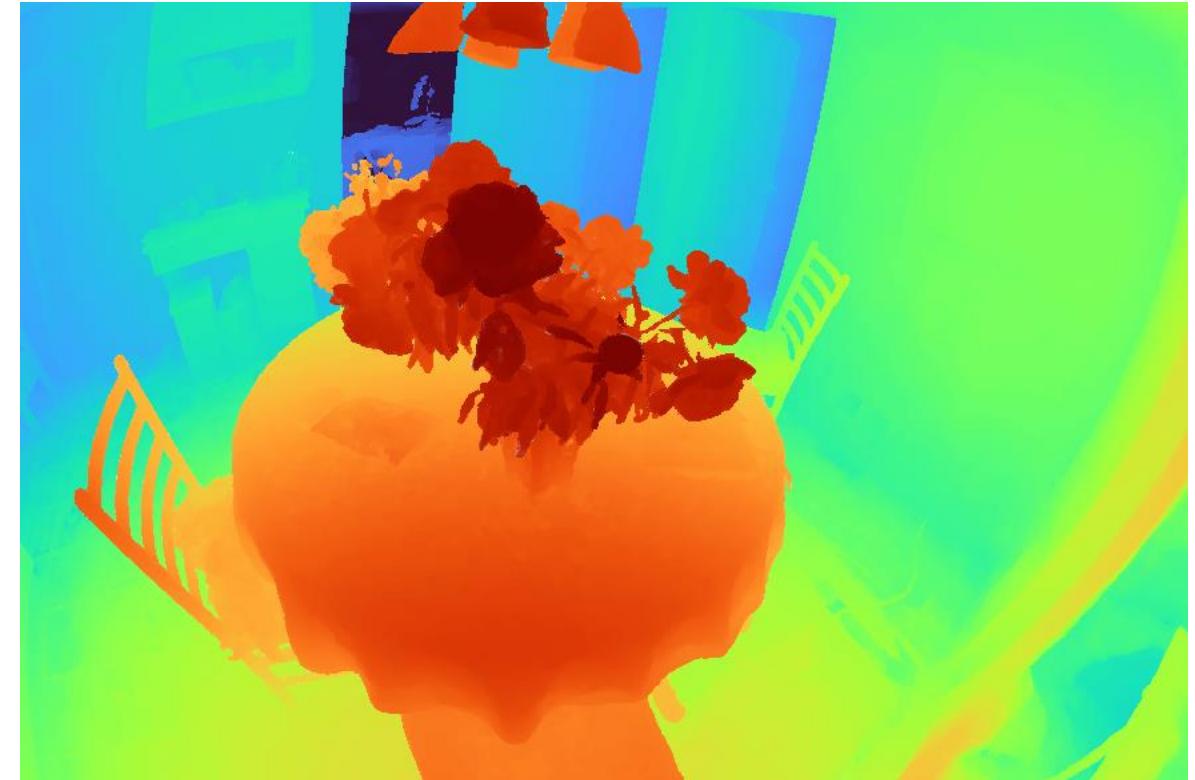
This is often how people visualise NeRF depth maps.

Alternatively, other statistics like mode or median can be used.

Rendering weight PDF is important – depth



Mean depth



Median depth

Volume rendering other quantities

This idea can be used for any quantity we want to “volume render” into a 2D image. If \mathbf{V} lives in 3D space (semantic features, normal vectors, etc.)

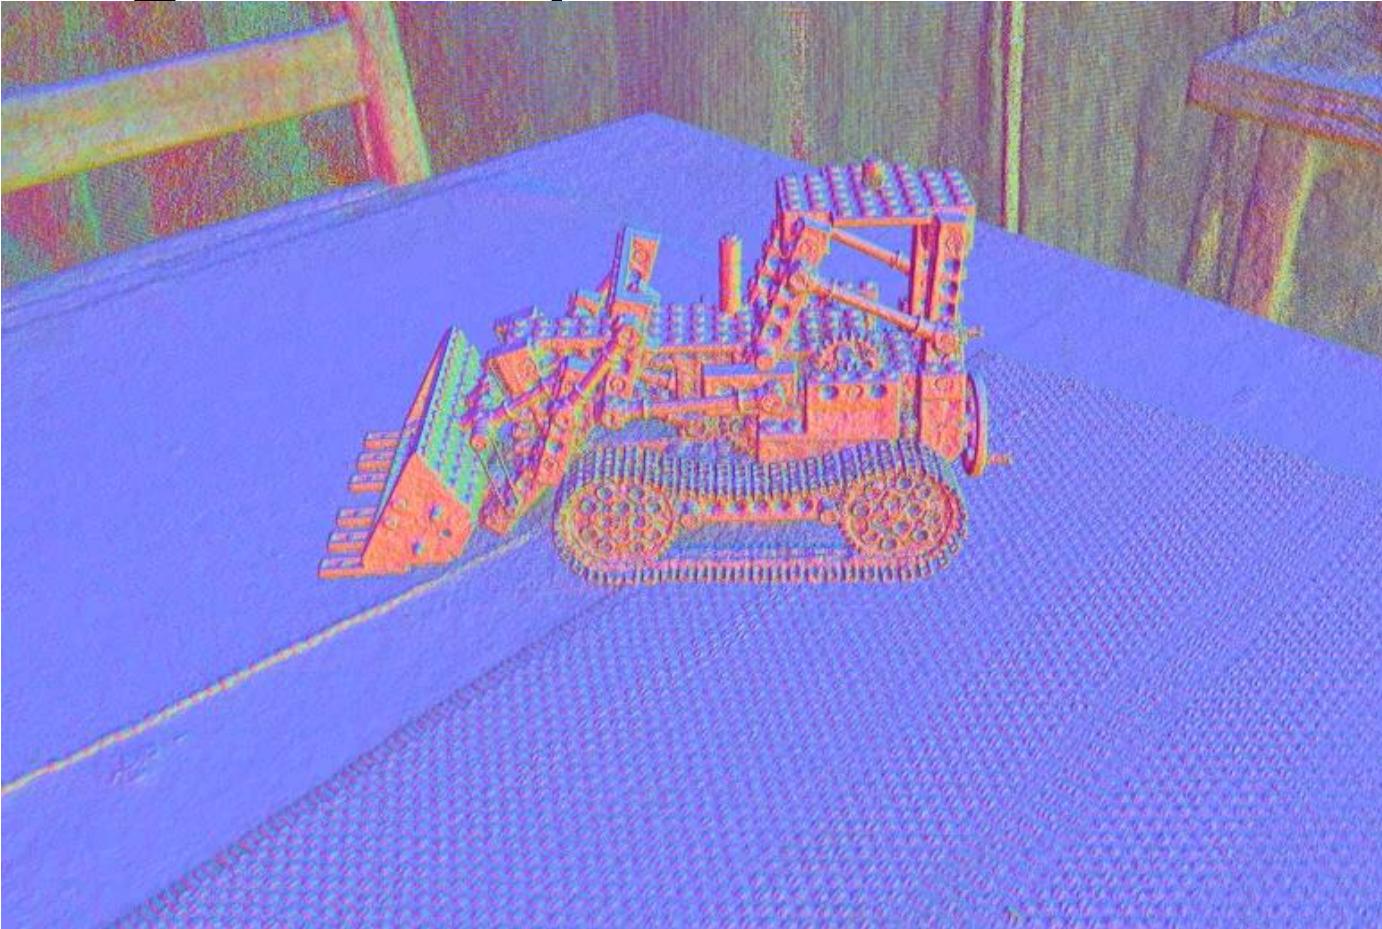
$$\sum_i T_i \alpha_i \mathbf{v}_i$$

can be taken per-ray to produce 2D output images.

Volume Rendering CLIP features

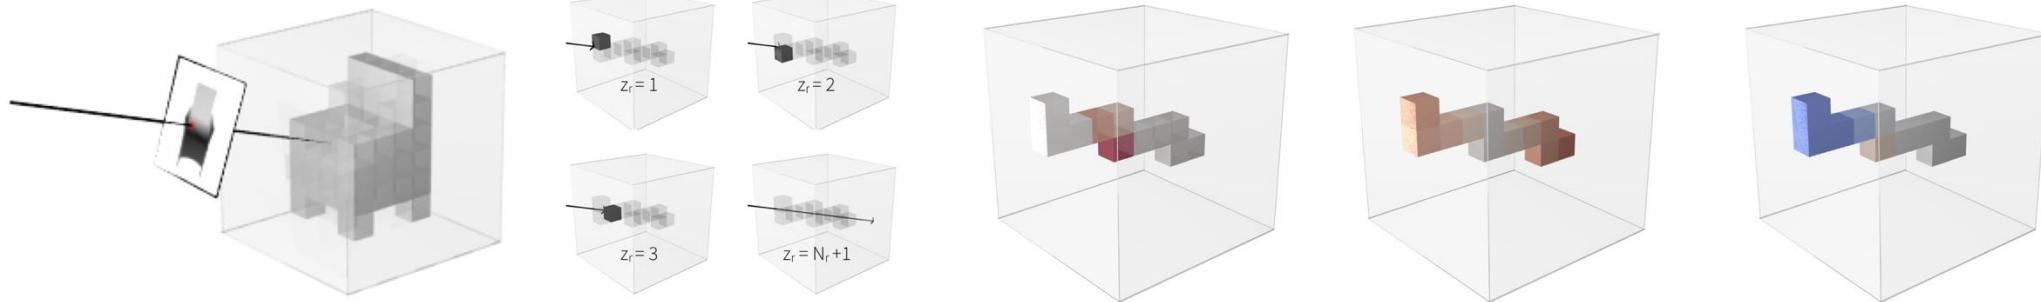


Density as geometry



Normal vectors (from analytic gradient of density)

Previous Papers



Differentiable ray consistency work used a forward model with “probabilistic occupancy” to supervise 3D-from-single-image prediction.
Same rendering model as alpha compositing!

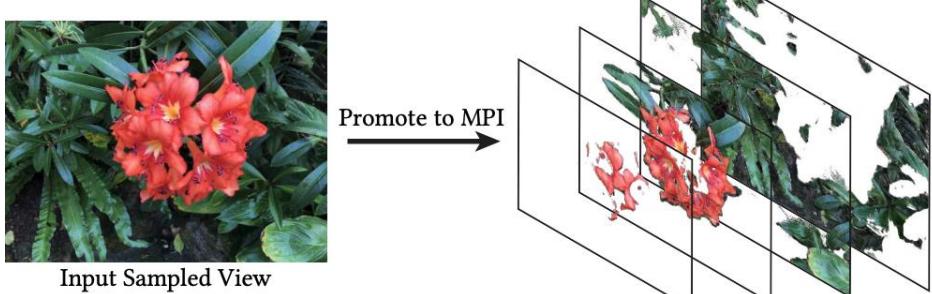
$$p(z_r = i) = \begin{cases} (1 - x_i^r) \prod_{j=1}^{i-1} x_j^r, & \text{if } i \leq N_r \\ \prod_{j=1}^{N_r} x_j^r, & \text{if } i = N_r + 1 \end{cases}$$

Similar Ideas before NeRF

Multiplane image methods

Stereo Magnification (Zhou et al. 2018)
Pushing the Boundaries... (Srinivasan et al. 2019)
Local Light Field Fusion (Mildenhall et al. 2019)
DeepView (Flynn et al. 2019)
Single-View... (Tucker & Snavely 2020)

Typical deep learning pipelines - images go into a 3D CNN, big RGBA 3D volume comes out



Neural Volumes

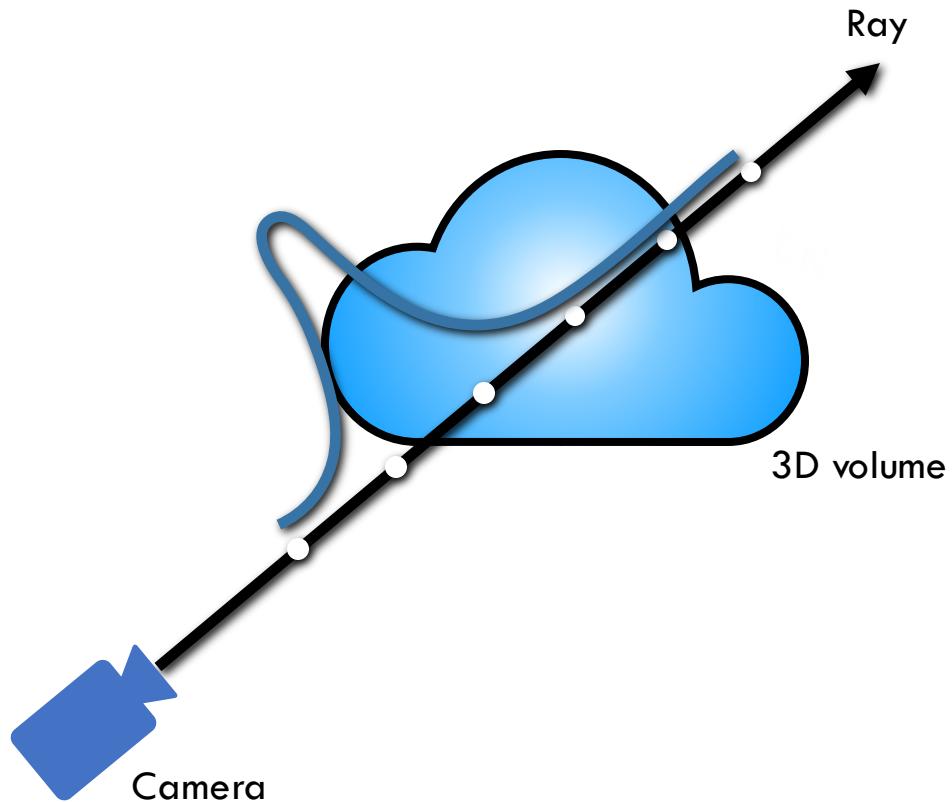
(Lombardi et al. 2019)

Direct gradient descent to optimize an RGBA volume, regularized by a 3D CNN



Signal Processing Consideration in NeRFs

What is this process?



What is happening here?

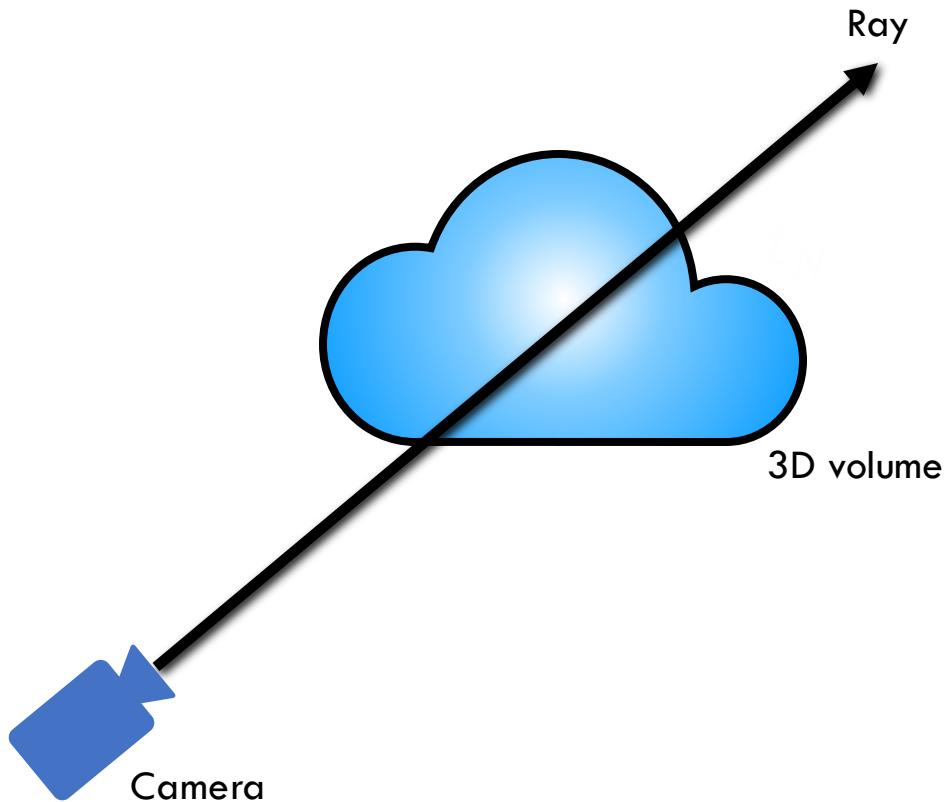
Aliasing!!



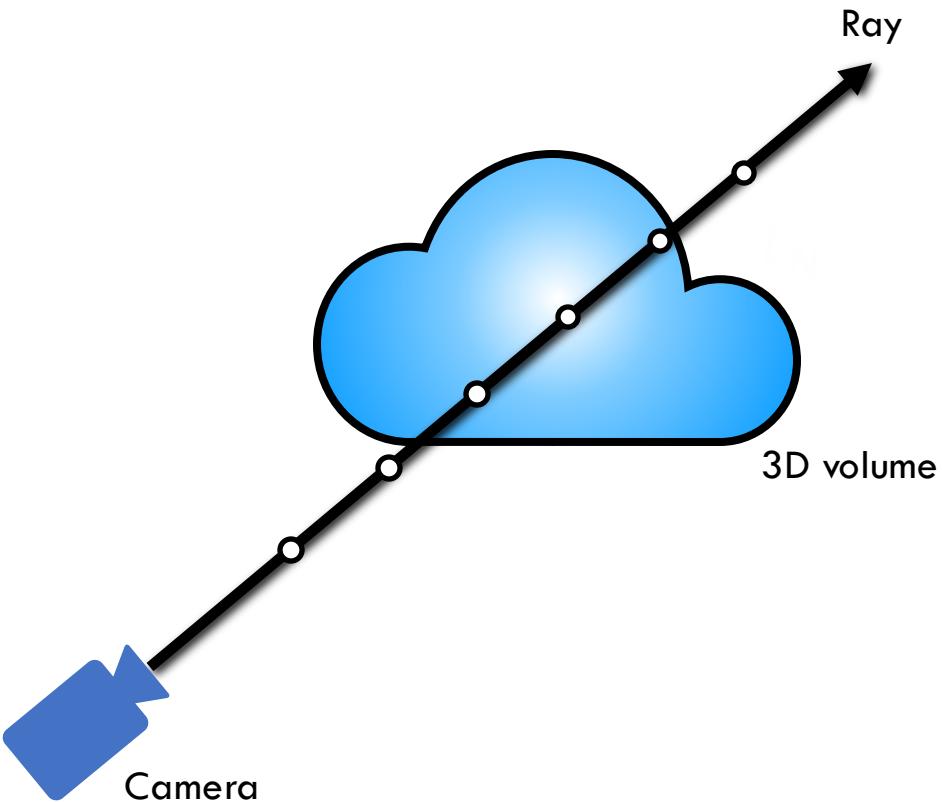
Naïve (original) NeRF



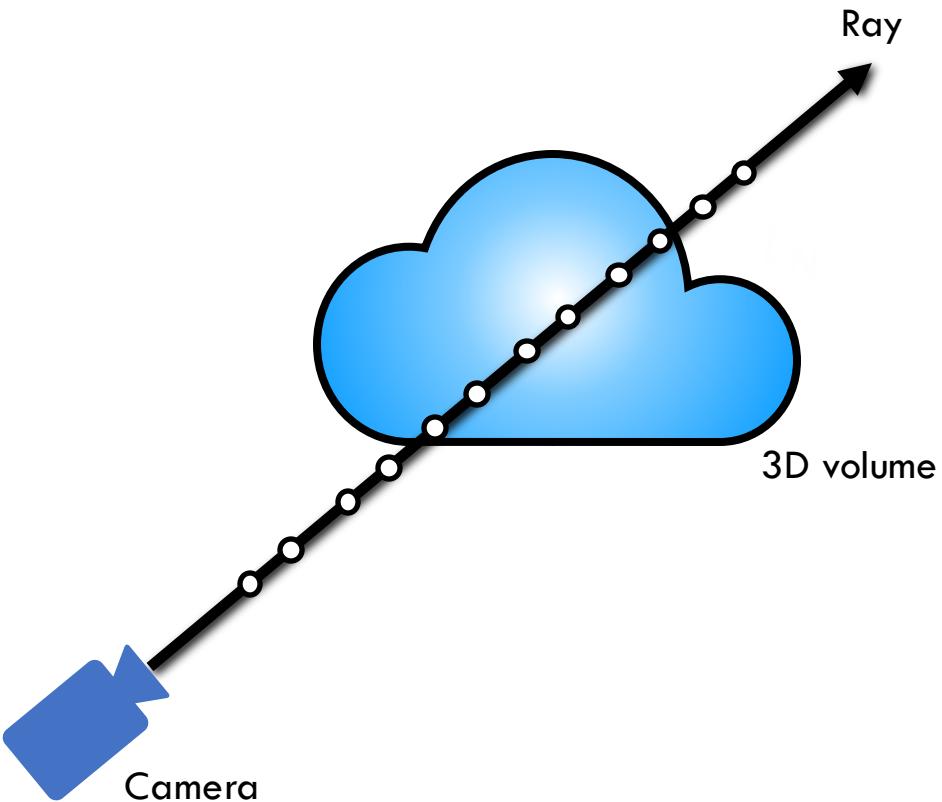
Where to place samples along rays?



How to be more efficient than dense sampling?



How to be more efficient than dense sampling?



Hierarchical Sampling vs. Acceleration Structures

Hierarchical Sampling vs. Acceleration Structures

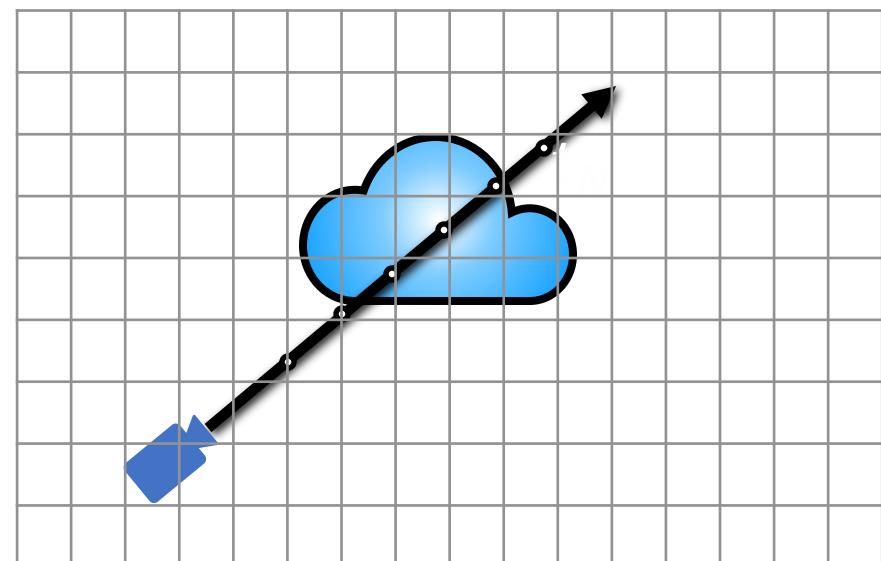
Hierarchical Sampling

Iteratively use samples from NeRF to more efficiently sample visible scene content

Acceleration Structures

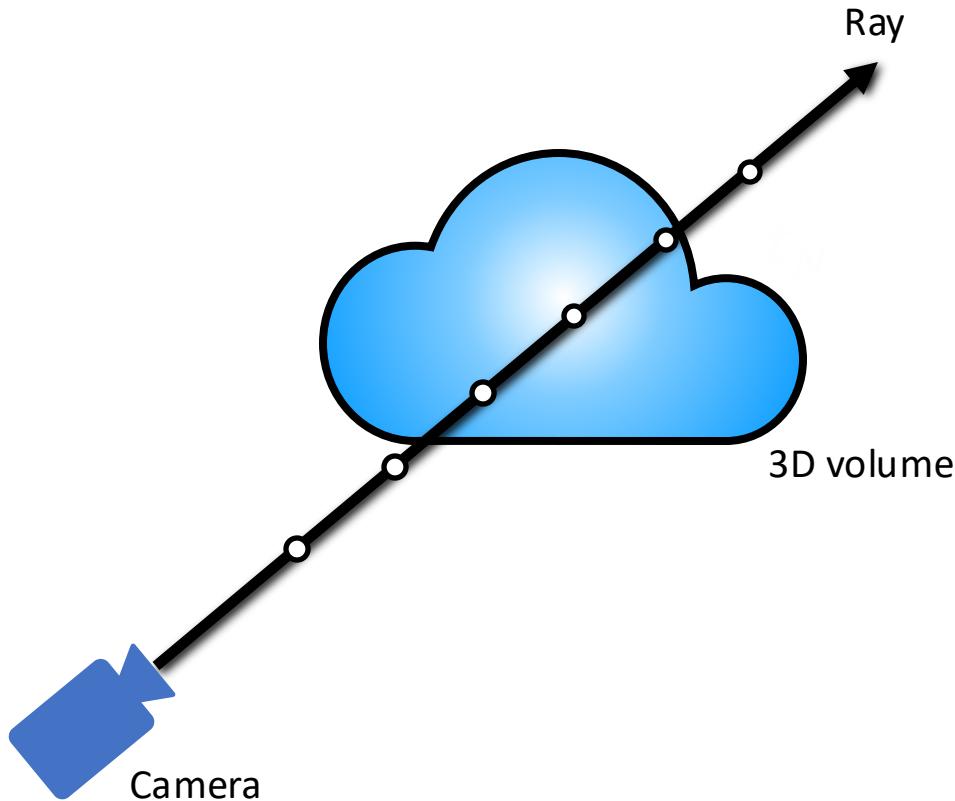
Distill/cache properties of NeRF into a structure that helps generate samples: e.g. Occupancy Grids

Straightforward compute —> storage tradeoff



Hierarchical ray sampling

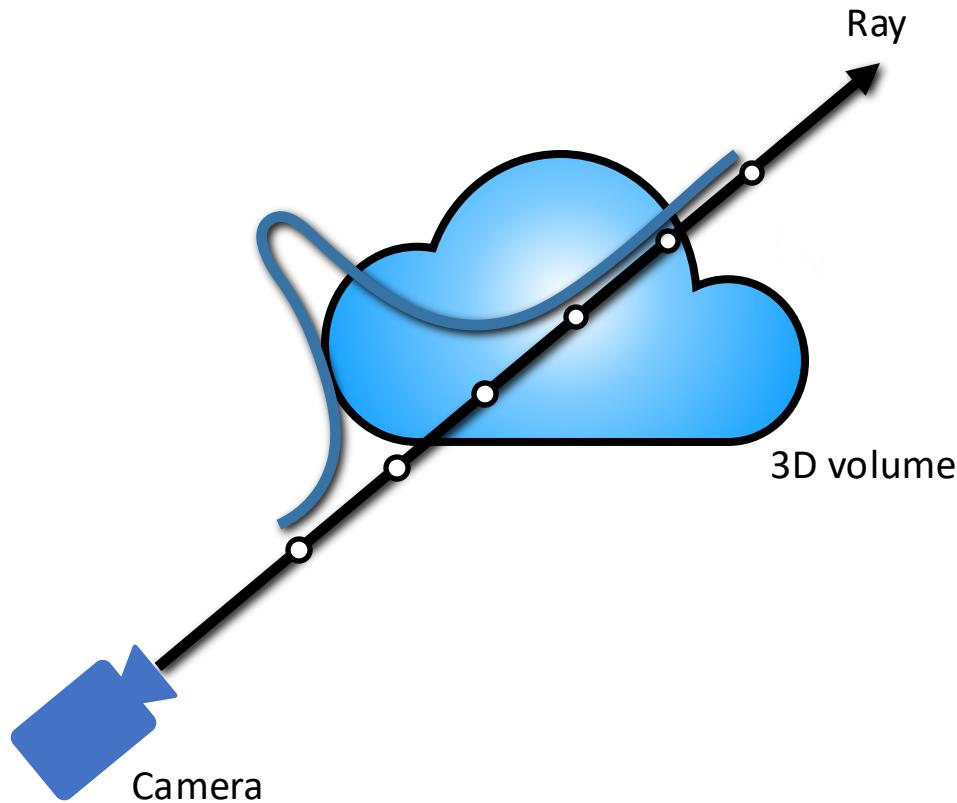
Key Idea: sample points proportionally to expected effect on final rendering



Key Idea: sample points proportionally to expected effect on final rendering

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$

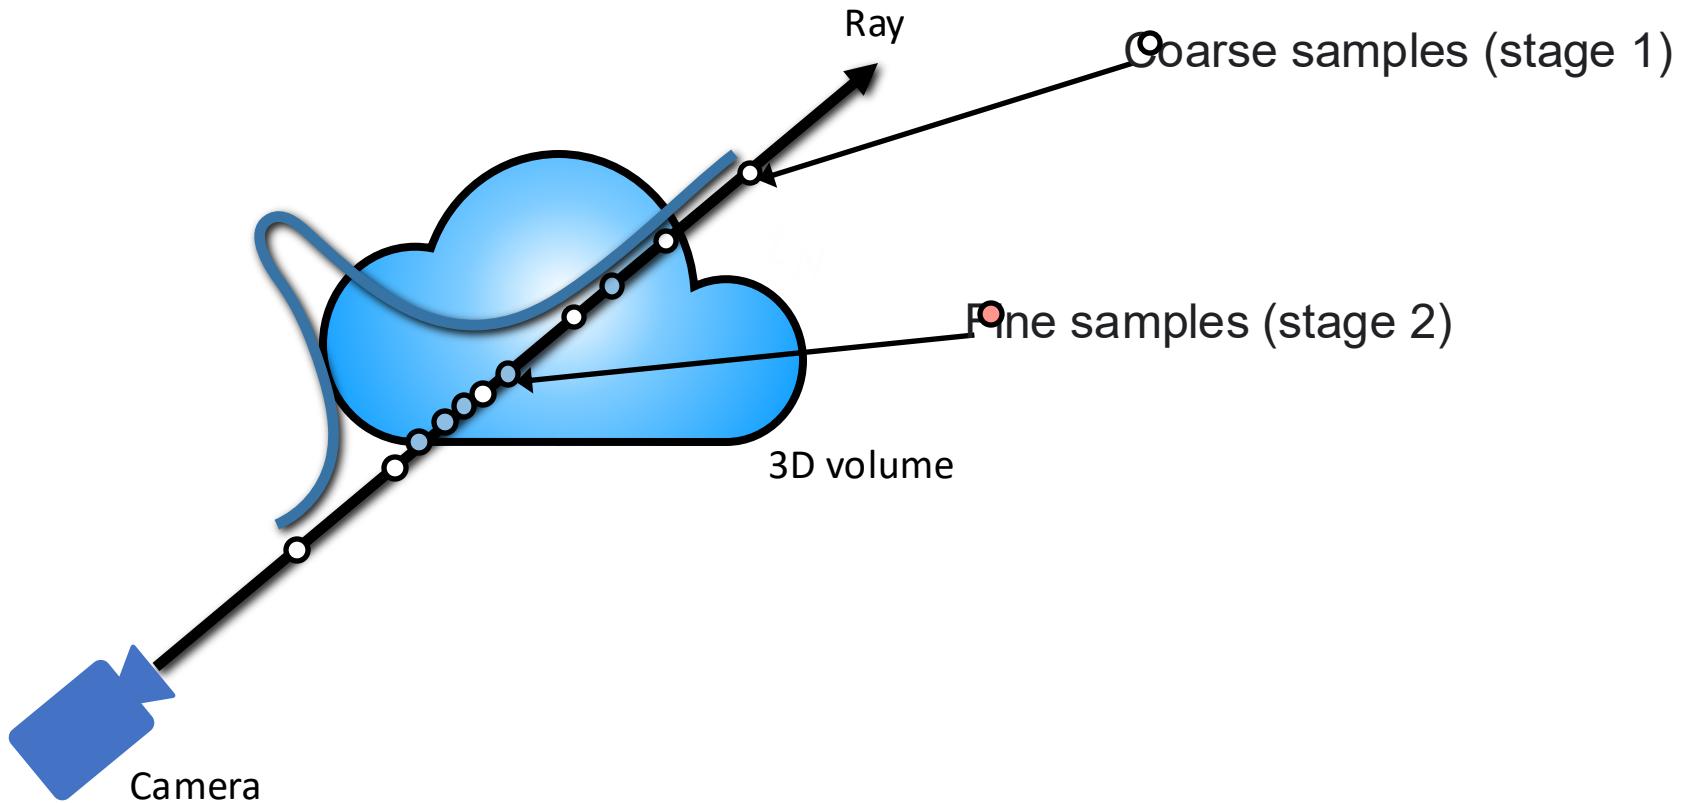
treat weights as probability distribution for new samples



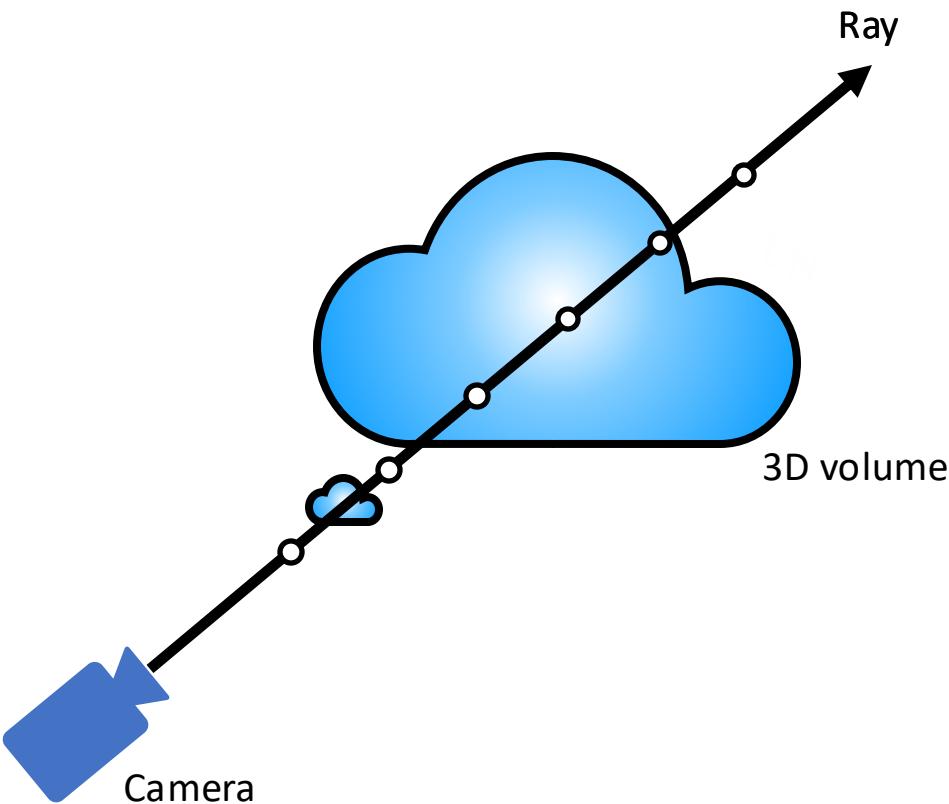
Key Idea: sample points proportionally to expected effect on final rendering

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$

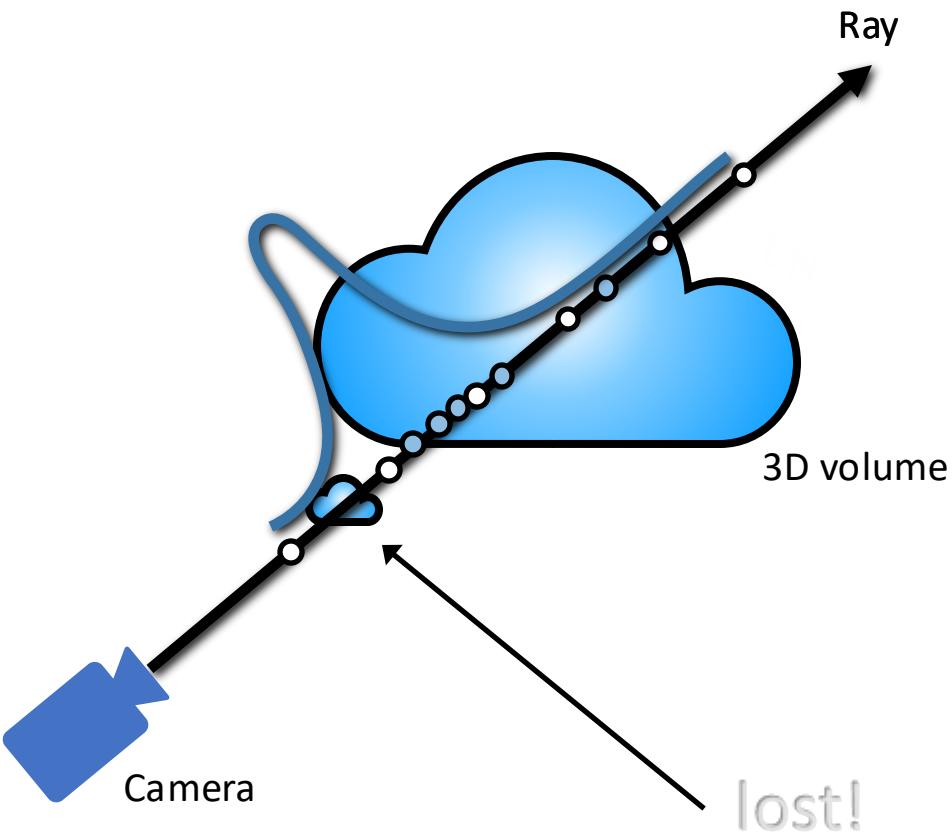
treat weights as probability distribution for new samples



What about aliasing during coarse sampling?

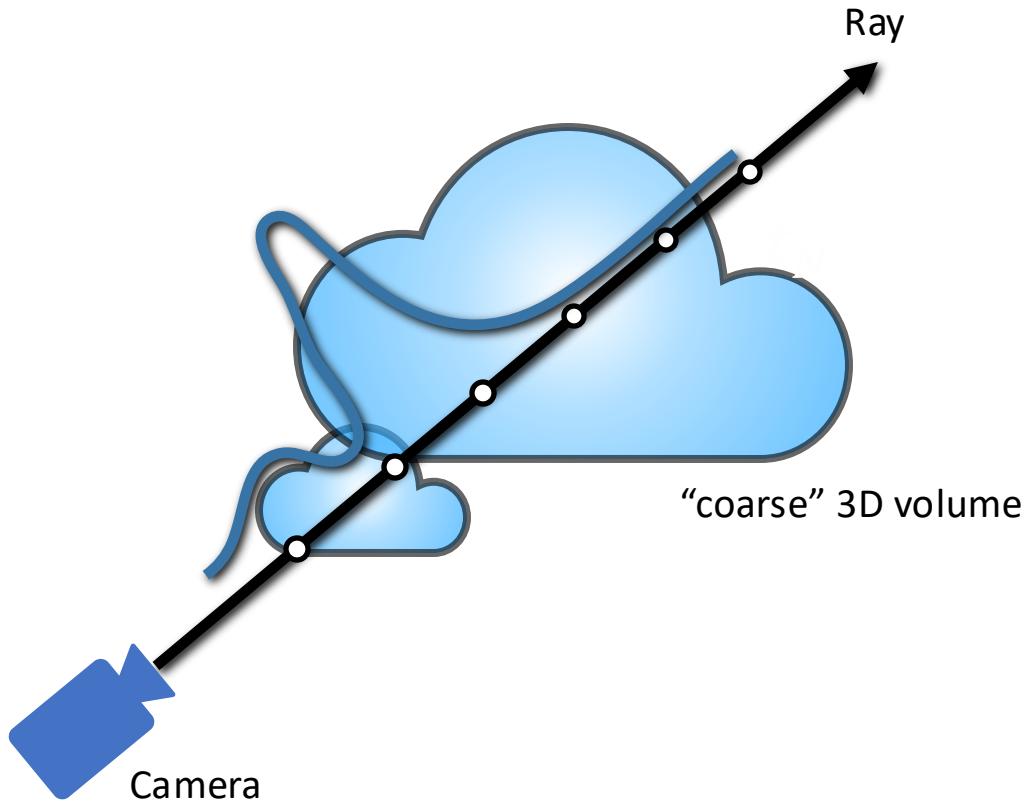


What about aliasing during coarse sampling?



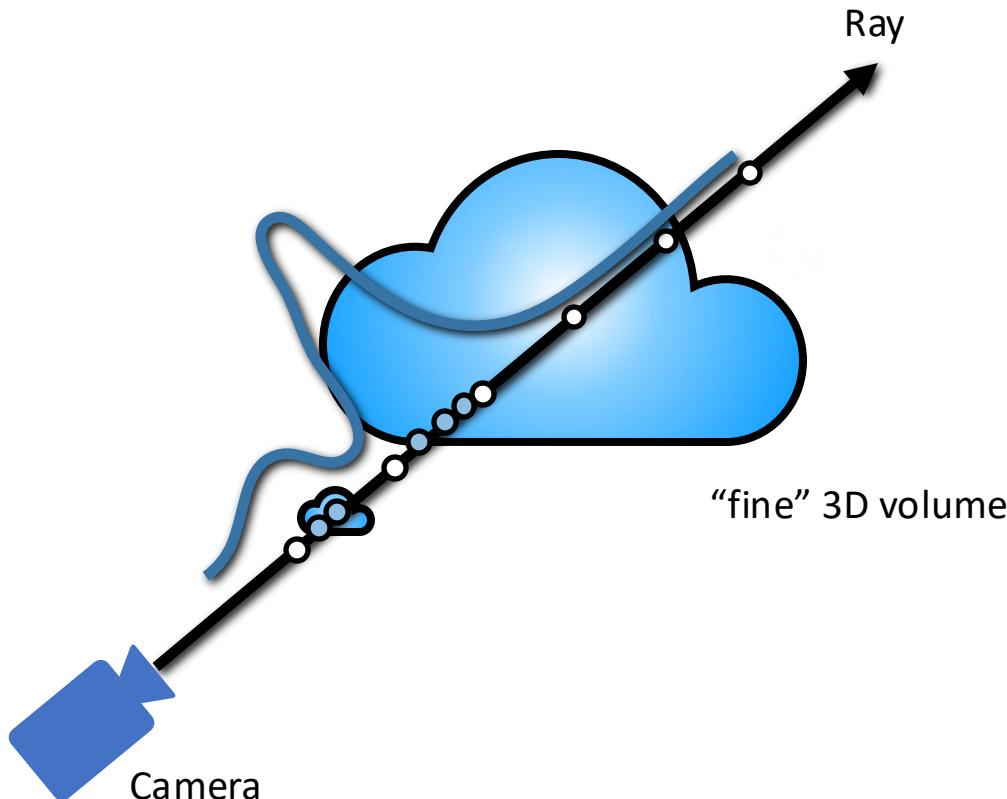
What about aliasing during coarse sampling?

Solution: train two NeRFs! —> lower resolution for first “coarse” level



What about aliasing during coarse sampling?

Solution: train two NeRFs! —> higher resolution for second “fine” level



This strategy is used in the original NeRF paper, you can implement it as an EC.

Further Reading on this Topic

- [MipNeRF](#)
 - Low-pass filter the positional encoding
- [MipNeRF360](#)
 - For 360 scenes
 - Train a “Proposal Network” instead of Coarse & Fine Networks that tells where to sample

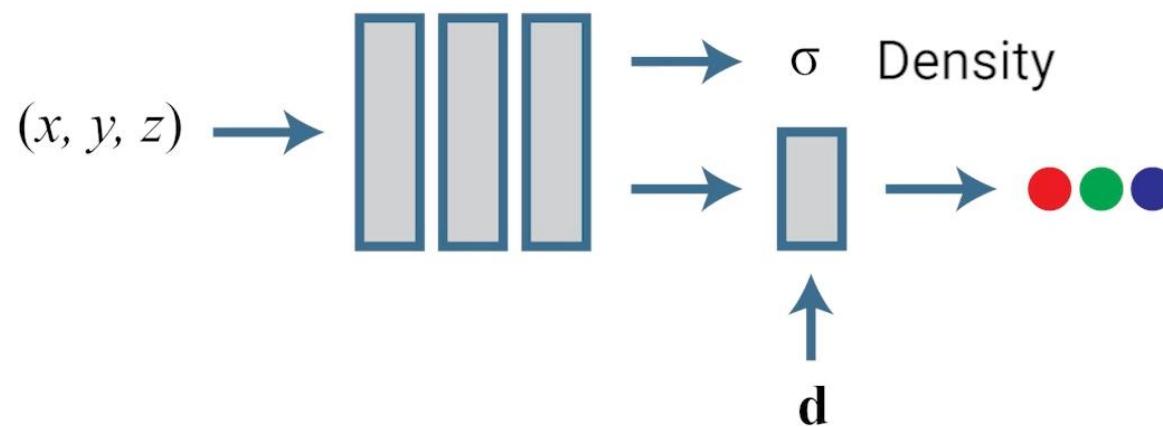
Other Advanced Topics

Compression capability of NeRF

- If Image is 256x256x3, why memorize this with a MLP?
 - 3 layer MLP with 256 neurons $\sim 3 \cdot 256^2$ learnable parameters
- In 3D:
 - if there are 100 input images: $100 \cdot 256 \times 256 \times 3 = 19M$
 - 8 layer MLP with 256 neurons $\sim 524K$
 - Just 3% of what it takes to hold 100 images
 - MLP size doesn't change even if we have 1000 input images
- Trade off?
 - SPEED!

Why is NeRF Slow?

NeRF



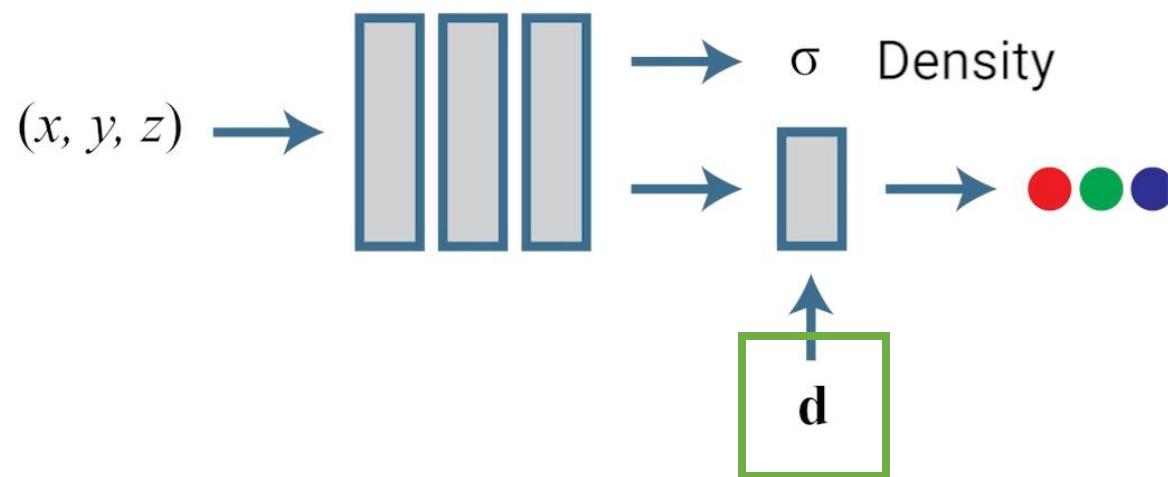
For each image
For each pixel (800, 800)
For each sample (256)
Eva1 NeRF Network

You have to sample densely in R^5

~ 163 million network calls

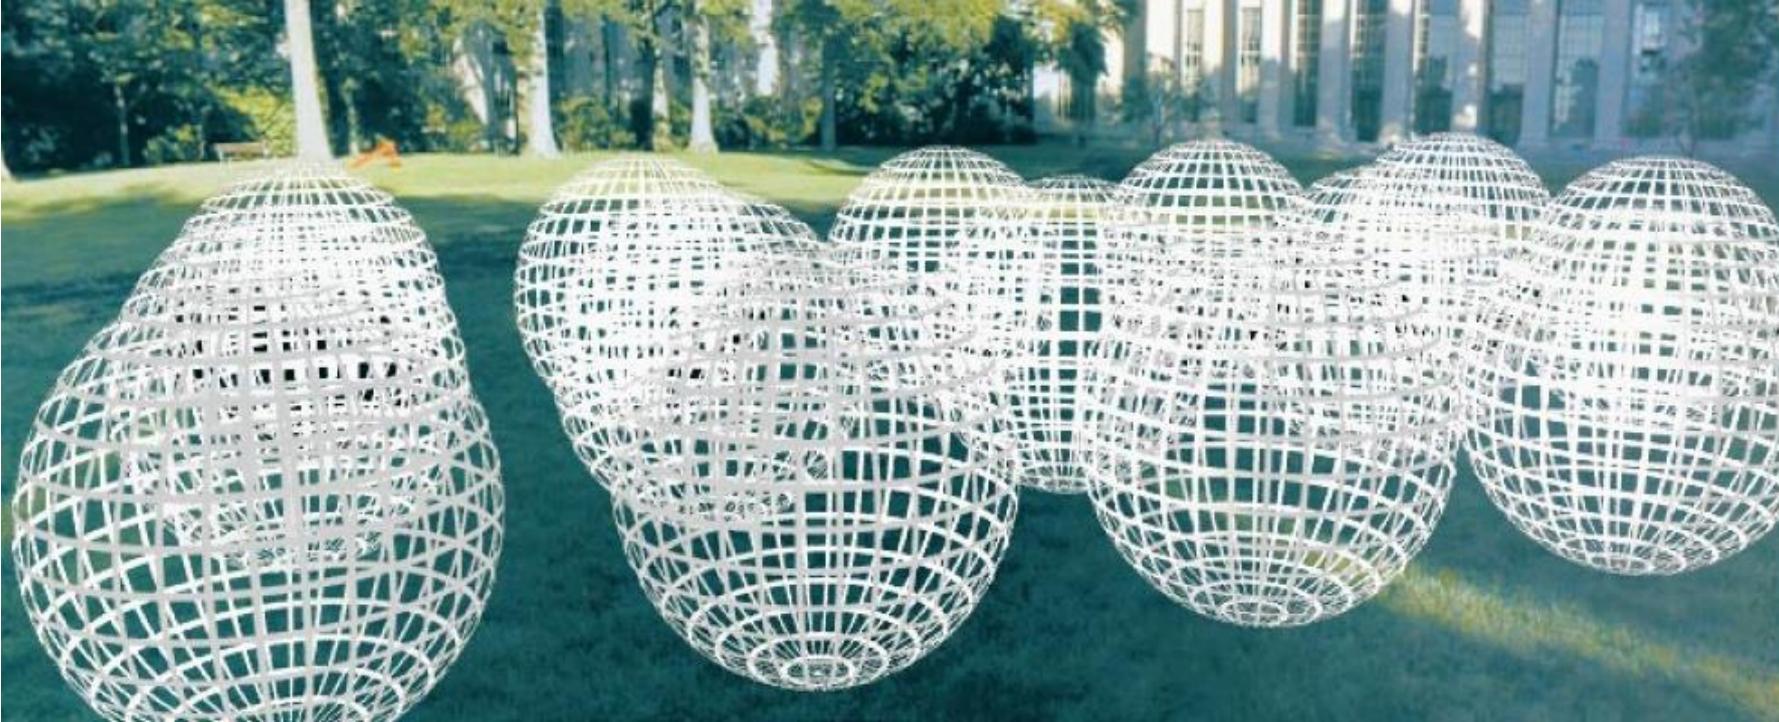
Cache everything?

NeRF



R^5 is a lot to cache

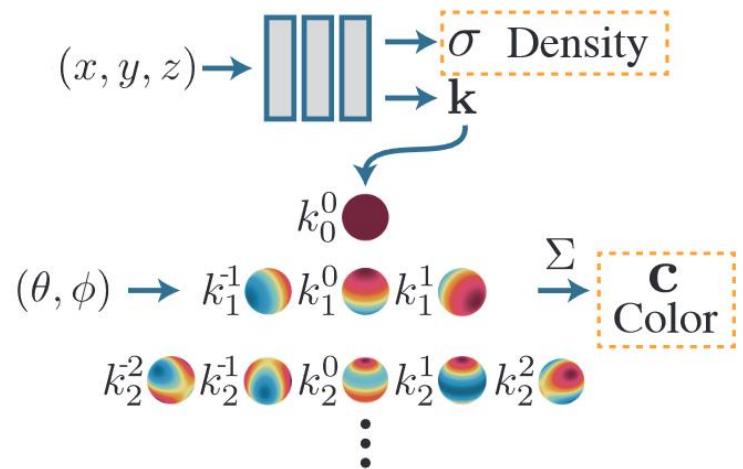
Recall:



We can factorize the location R^3 & the view direction R^2 !

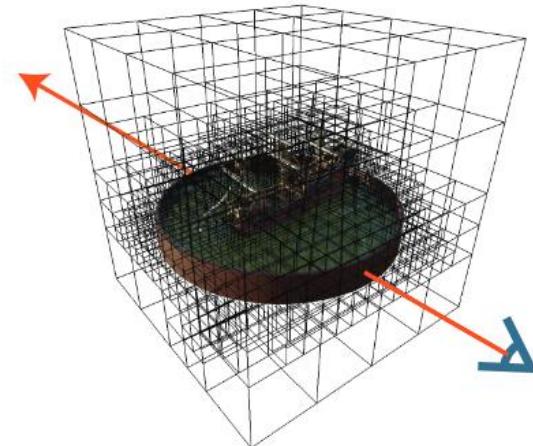
Idea

1. Factor out view-dependent effects



Spherical Harmonics

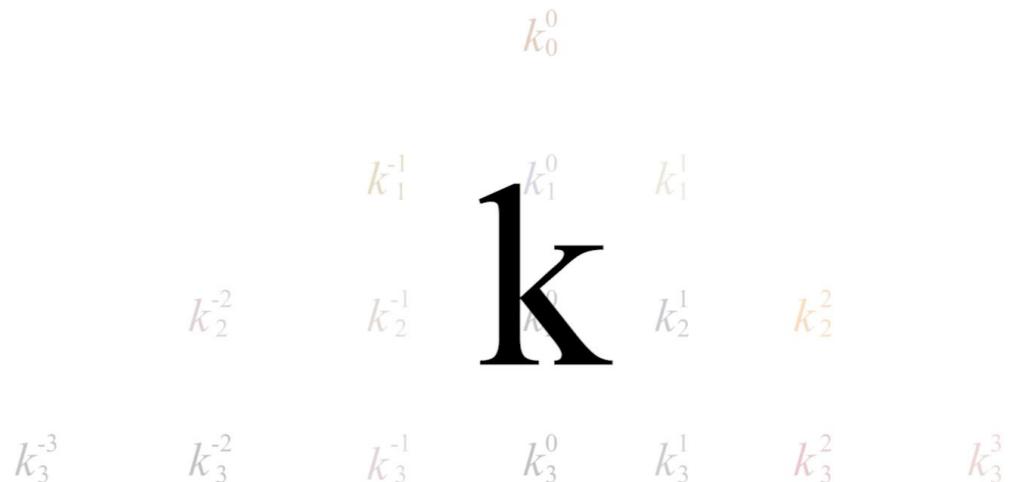
2. Store density + SH in efficient data structure (Octree)



PlenOctree =
View-dependent Octree

Spherical Harmonics to model view dependent Color

- Fourier basis on the sphere



SH models function on a sphere:

$$f(s) = \sum_{i=0}^{n^2} k_i y_i(s)$$

s – point on a unit sphere, $s = (\varphi, \theta)$

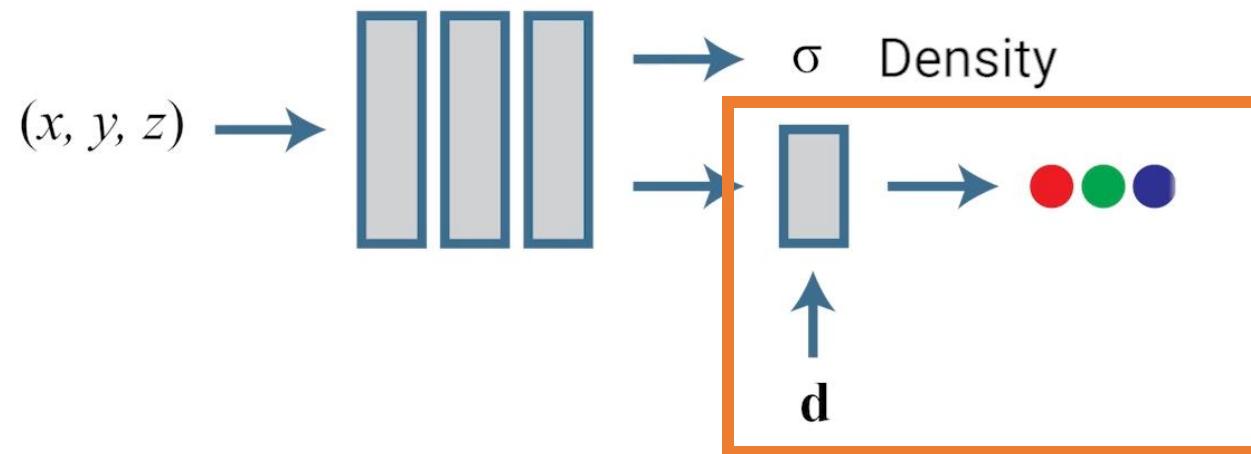
k_i - coefficient of i-th basis

$y_i(s)$ - SH basis, analytically computed, as shown on the video

Learn k_i

NeRF with Spherical Functions

NeRF

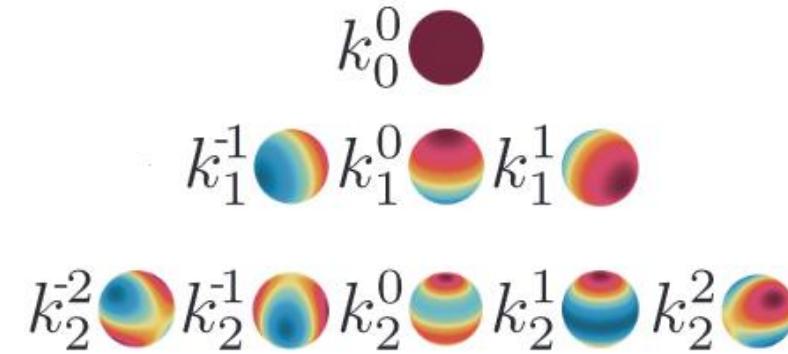
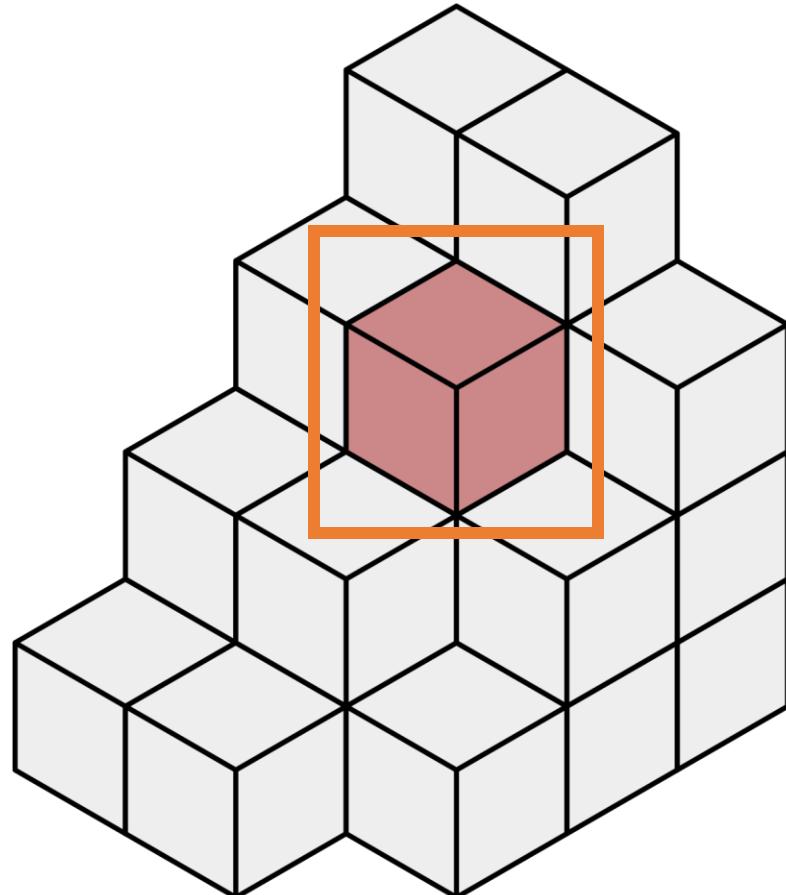


NeRF with Spherical Harmonics (NeRF-SH)



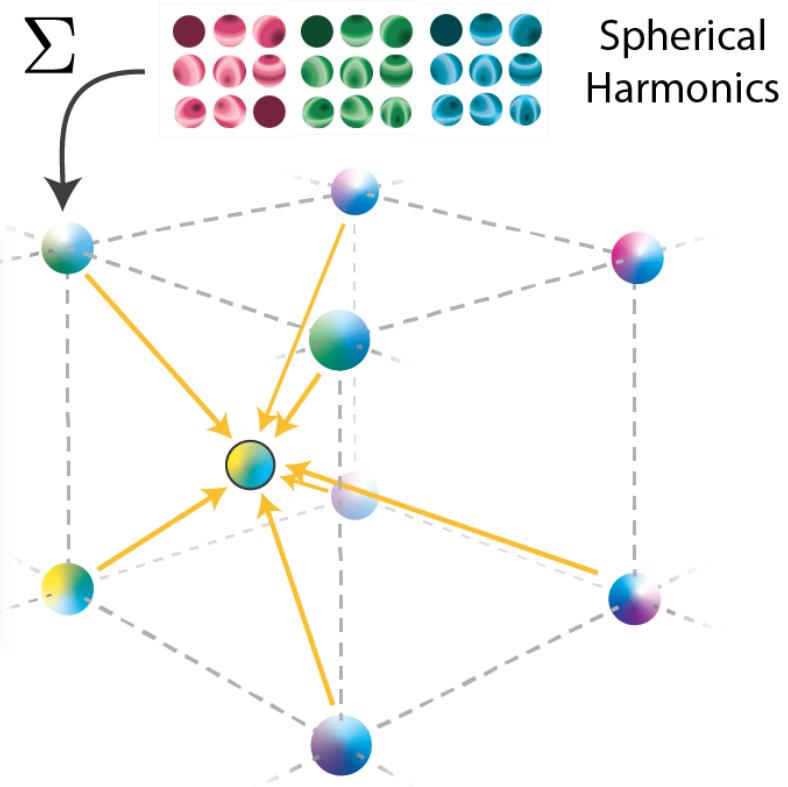
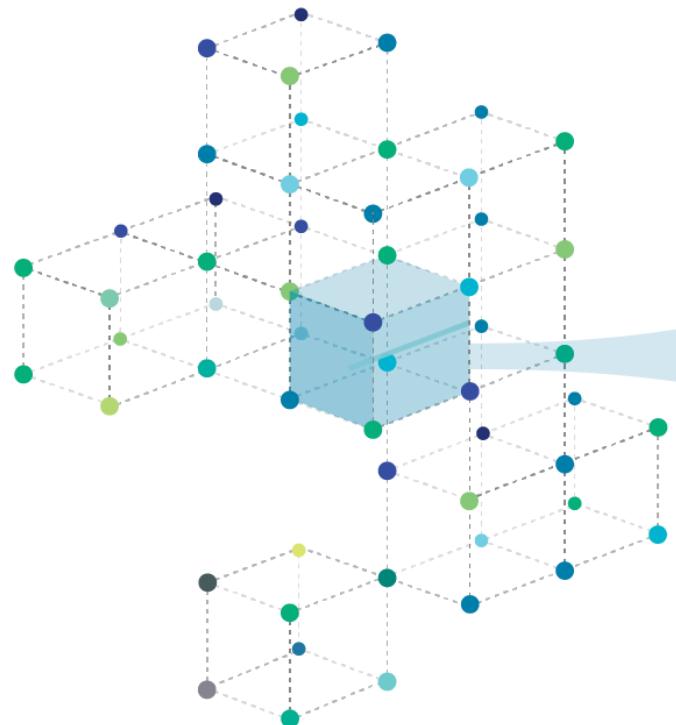
$$\mathbf{k} = \begin{matrix} k_0^0 \\ k_1^{-1} \\ k_1^0 \\ k_1^1 \\ k_2^{-2} \\ k_2^{-1} \\ k_2^0 \\ k_2^1 \\ k_2^2 \\ k_3^{-3} \\ k_3^{-2} \\ k_3^{-1} \\ k_3^0 \\ k_3^1 \\ k_3^2 \\ k_3^3 \end{matrix}$$

PlenOctree = Sparse Voxels + SH coefficients



Only stores nonempty voxels

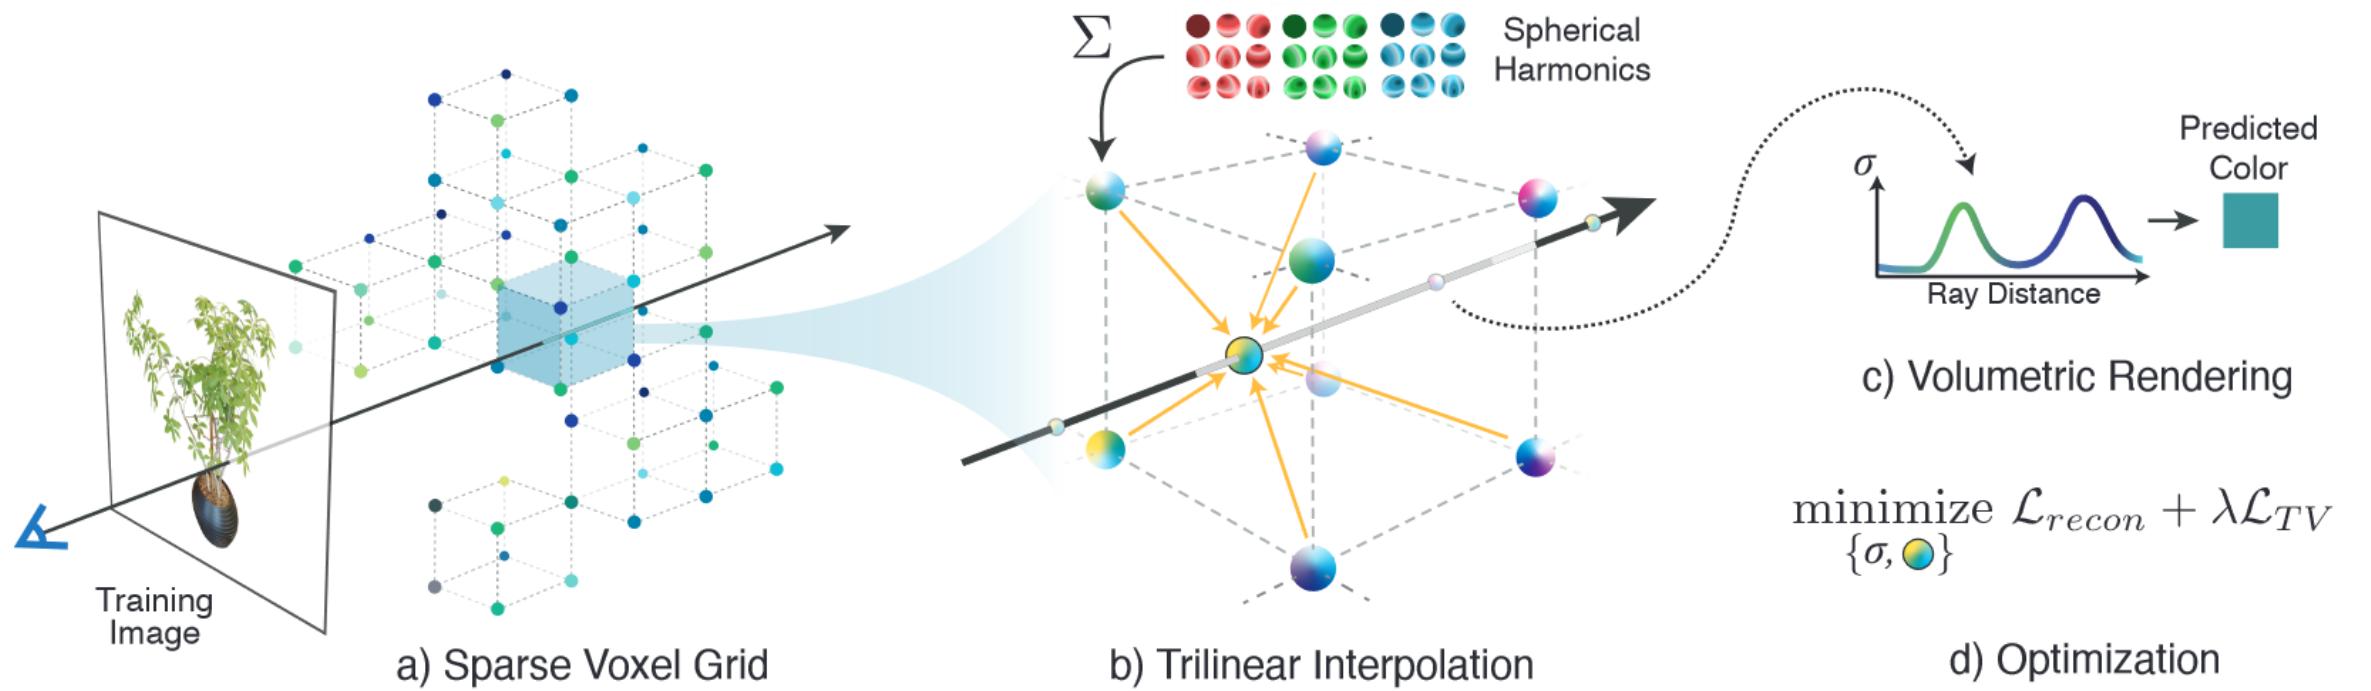
Plenoxel = “Plenoptic Volume Element”



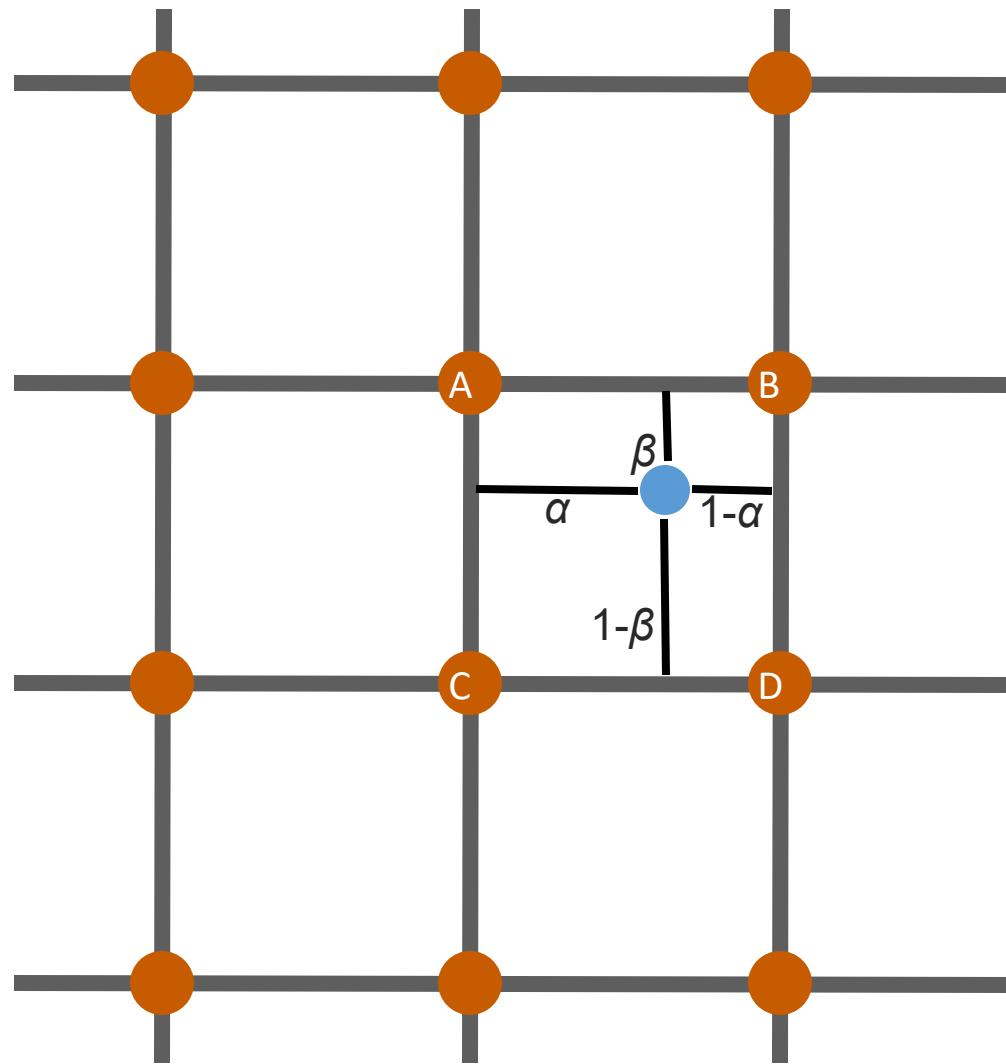
Generalization of PlenOctree – A data structure consisting of Plenoxels

Yu et al. CVPR 2022

No MLP required..



Make continuous via tri-linear interpolation



$$\bullet = \beta(\alpha A + (1 - \alpha)B) + (1 - \beta)(\alpha C + (1 - \alpha)D)$$

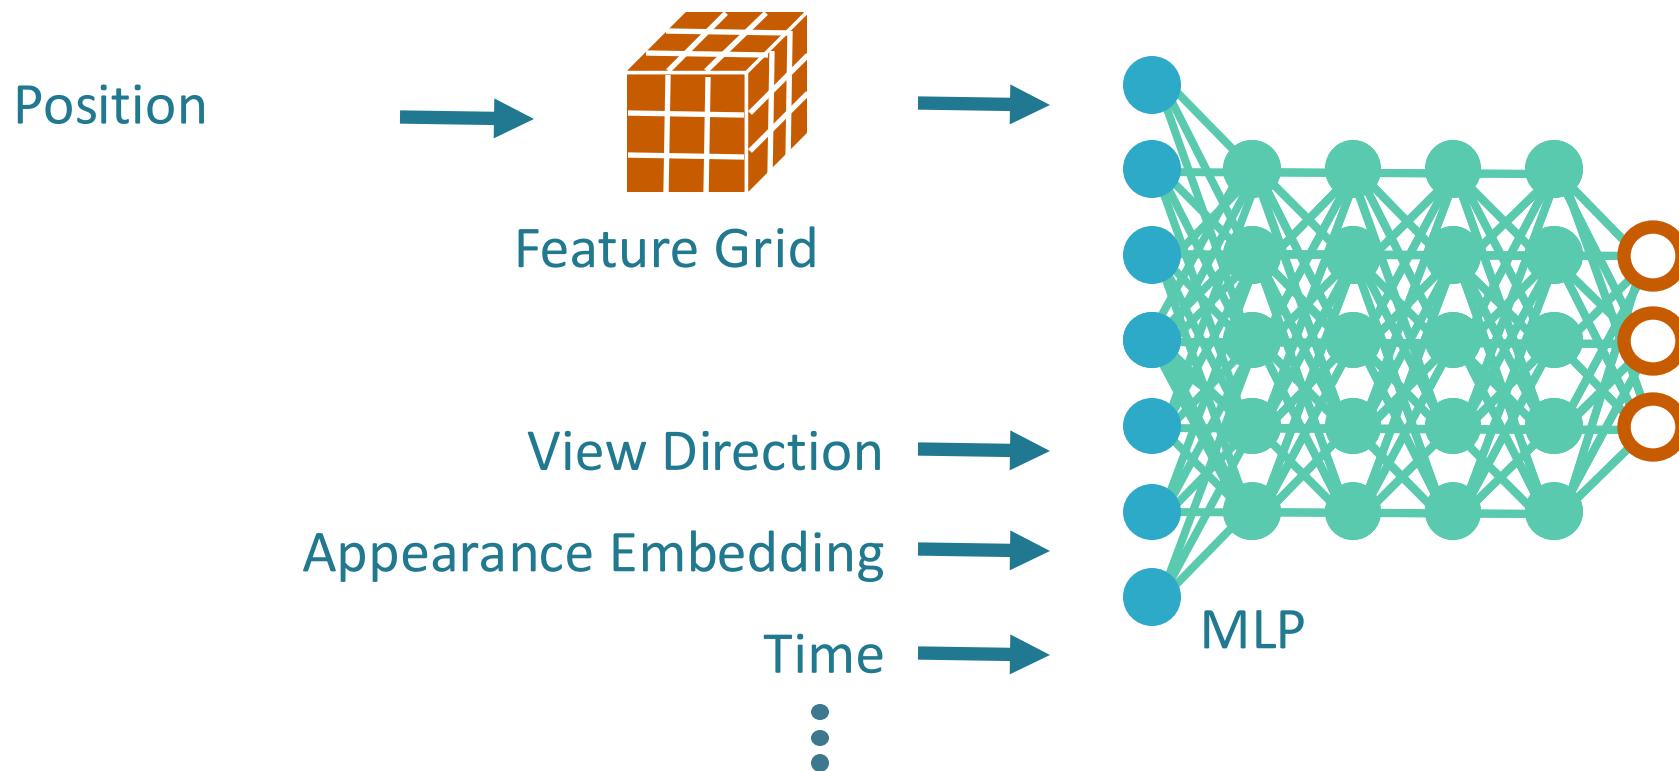
NeRF

[Mildenhall et al. ECCV 2020]

Plenoxels

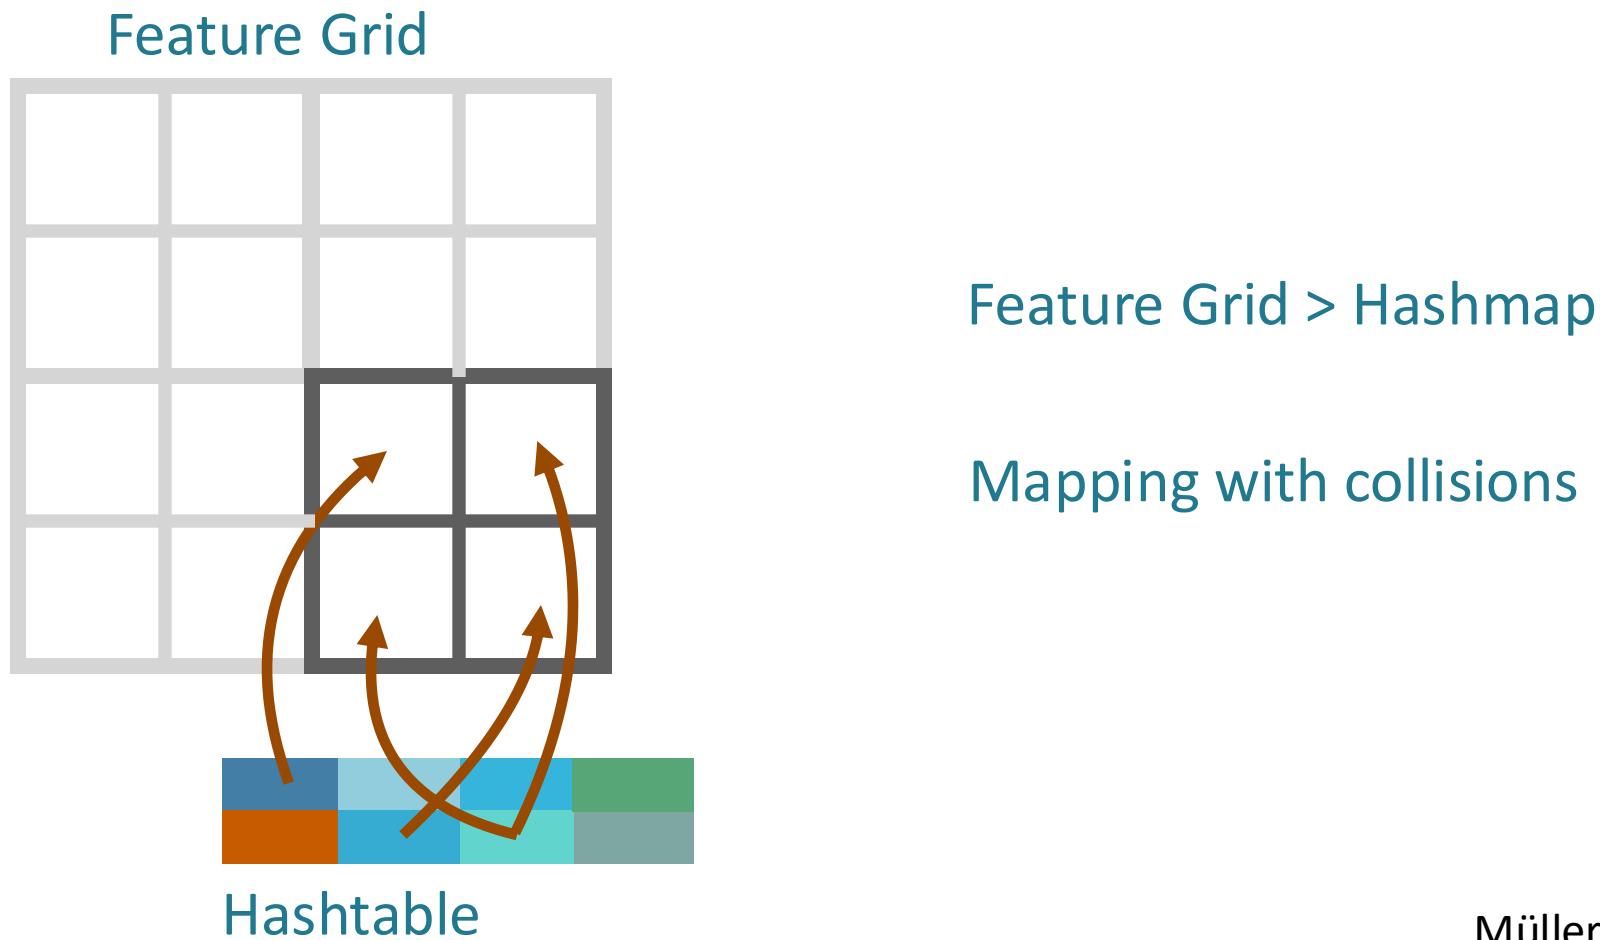
minutes seconds
00:00

But MLPs are convenient



InstantNGP

- Hybrid approach
- Grid → hashmap → Feature retrieval
- Pass to a small MLP
- Fast and convenient!

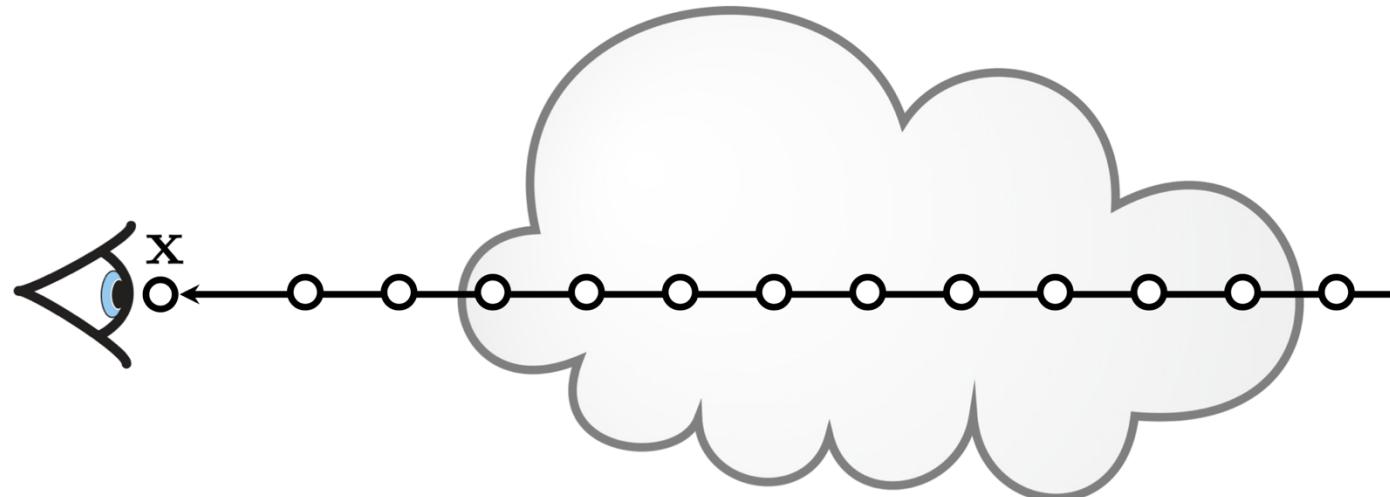


Latest: Gaussian Splatting

- Approximate with 3D Gaussian points
- Rasterize instead of volrend
- **Still alpha-blend**
- no MLP + SH like plenoxels



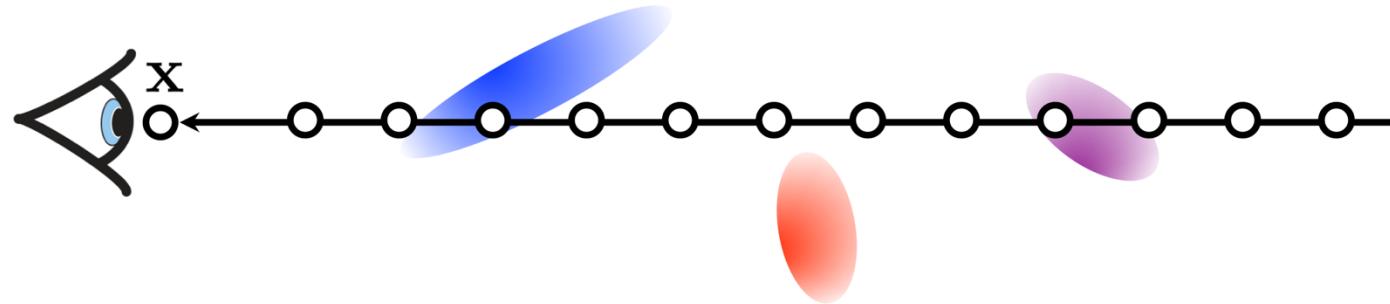
Recap: Rendering Volumes



1. Draw samples along the ray

2. Aggregate their contributions to render

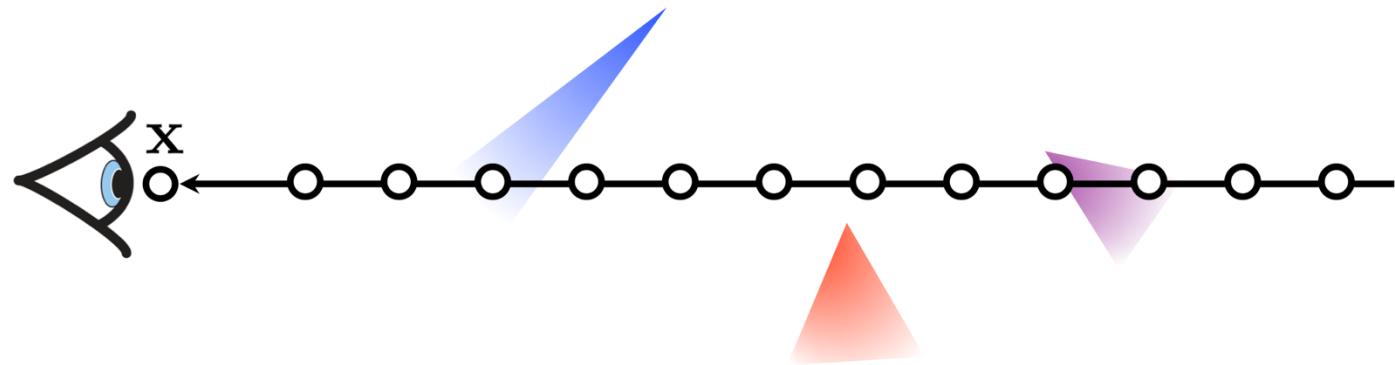
Rendering Primitives (e.g. Gaussians)



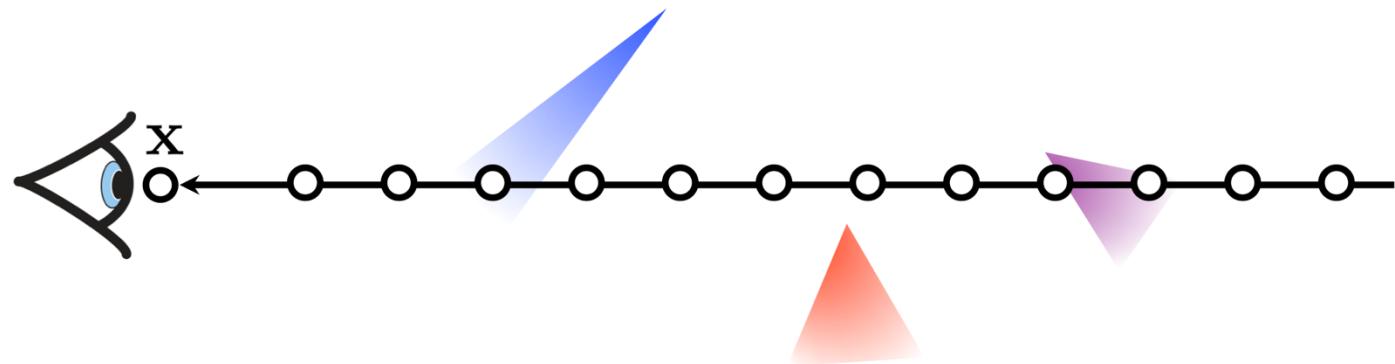
1. Draw samples along the ray

2. Aggregate their contributions to render

Rendering Primitives (e.g. Triangles/Meshes)



Rendering Primitives (e.g. Triangles/Meshes)



~~1. Draw samples along the ray~~

(wasteful — we know where
the primitives are!)

1. Find primitives that affect the pixel

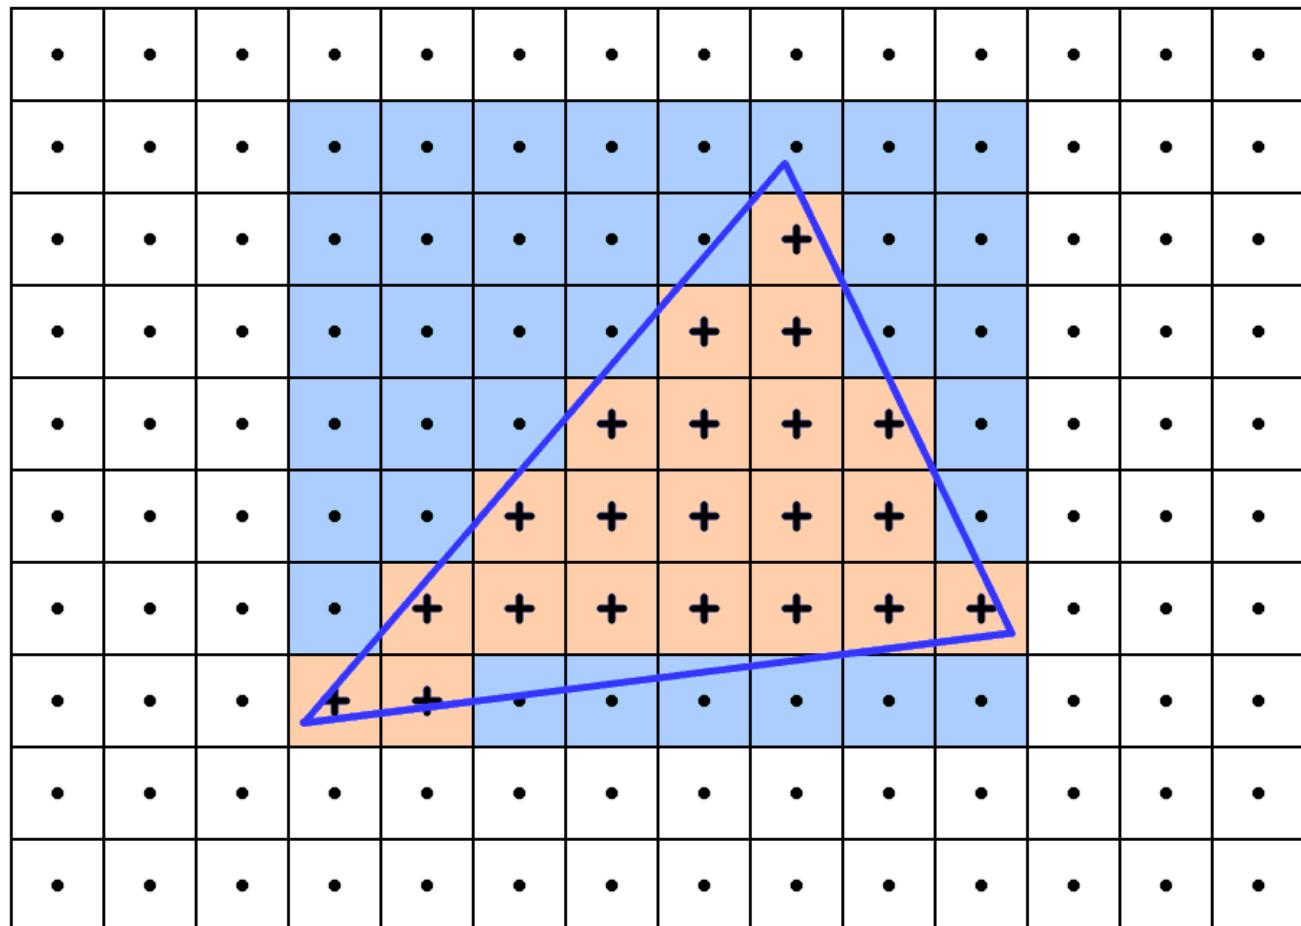
2. Aggregate their contributions to render

Rendering a mesh via Rasterization

Rasterization = conversion of primitives to pixels (details in CS184)

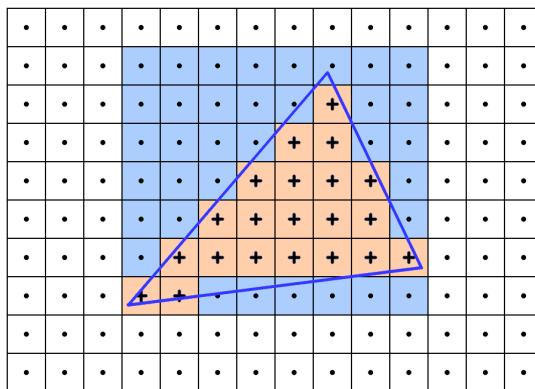
Rasterizat

Blending (+ S|



Rendering a mesh via Rasterization

Rasterization = conversion of primitives to pixels (details in CS184)



Rasterization

{

Blending (+ Shading)

```
def render(mesh, camera):
    ### some structure to store K triangles
    ### for each pixel
    sorted_k_closest_img = ...

    ### Iterate over all triangles
    for triangle in mesh:
        tri_2d = project(triangle, image)
        #####
        # Update closest_K_img for pixels within tri_2d
        #####
        #####
        # Iterate over all pixels
        for pixel in camera.grid:
            ### Iterate over triangles influencing this pixel
            for triangle in sorted_k_closest_img[pixel]:
                #####
                # Aggregate appearance
                #####|
```

Differentiable Gaussian Rendering

**What is the representation
of a 3D Gaussian?**

**How to project to
2D and rasterize?**

**How to model/aggregate
appearance?**

Rasterization



Blending



```
def render(gaussians, camera):
    """
    # Initialize a rasterization data structure
    # (records influencing primitives for each pixel)
    """
```

```
for gaussian in gaussians:
    gauss2d = project(gaussian, camera)
    """
    # Update rasterization data structure
    """
```

```
for pixel in camera.grid:
    """
    # Aggregate appearance from influencing gaussians
    """
```

Differentiable Gaussian Rendering

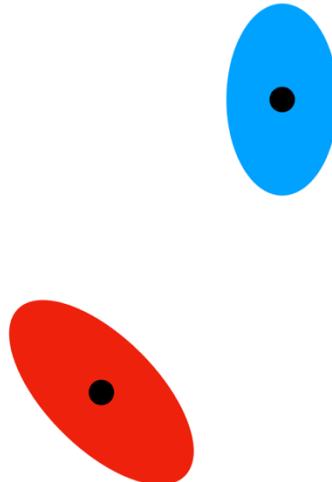
**What is the representation
of a 3D Gaussian?**

**How to project to
2D and rasterize?**

**How to model/aggregate
appearance?**

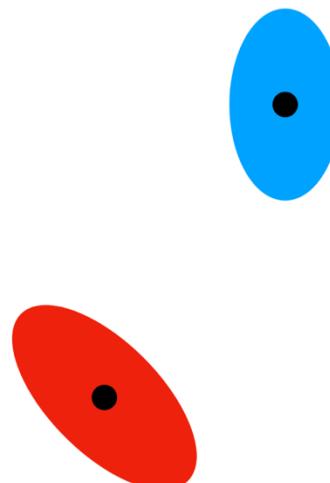
Position \mathbf{p}

$$\mathcal{G}_{\mathbf{V}}(\mathbf{x} - \mathbf{p}) = \frac{1}{2\pi|\mathbf{V}|^{\frac{1}{2}}} e^{-\frac{1}{2}(\mathbf{x}-\mathbf{p})^T \mathbf{V}^{-1} (\mathbf{x}-\mathbf{p})}$$



Differentiable Gaussian Rendering

**What is the representation
of a 3D Gaussian?**



Position \mathbf{p}

$$S = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & s_z \end{bmatrix}$$

$$R \in SO(3)$$

**How to project to
2D and rasterize?**

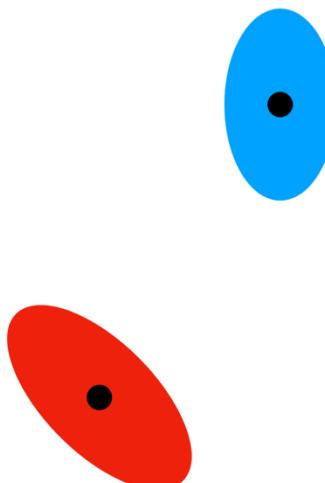
$$\mathcal{G}_{\mathbf{V}}(\mathbf{x} - \mathbf{p}) = \frac{1}{2\pi|\mathbf{V}|^{\frac{1}{2}}} e^{-\frac{1}{2}(\mathbf{x}-\mathbf{p})^T \mathbf{V}^{-1} (\mathbf{x}-\mathbf{p})}$$

Factorize as scale and rotation: $\mathbf{V} = RSS^T R^T$

**How to model/aggregate
appearance?**

Differentiable Gaussian Rendering

What is the representation
of a 3D Gaussian?



Position \mathbf{p}

$$S = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & s_z \end{bmatrix}$$

$$R \in SO(3)$$

How to project to
2D and rasterize?

$$\mathcal{G}_{\mathbf{V}}(\mathbf{x} - \mathbf{p}) = \frac{1}{2\pi|\mathbf{V}|^{\frac{1}{2}}} e^{-\frac{1}{2}(\mathbf{x}-\mathbf{p})^T \mathbf{V}^{-1} (\mathbf{x}-\mathbf{p})}$$

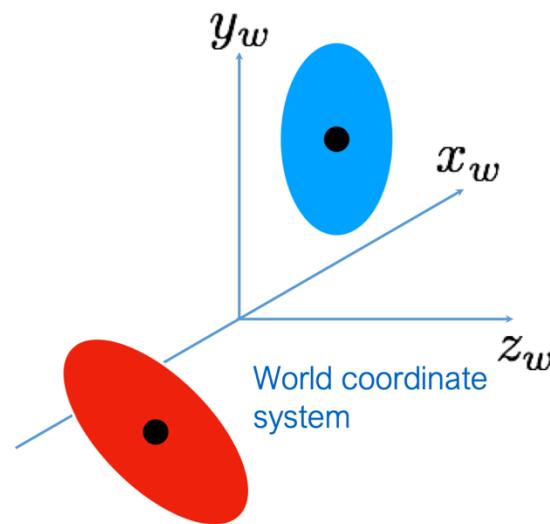
Factorize as scale and rotation: $\mathbf{V} = RSS^TR^T$

Each Gaussian also has an opacity and view-dependent color (via SH coefficients): α, \mathbf{c}

How to model/aggregate
appearance?

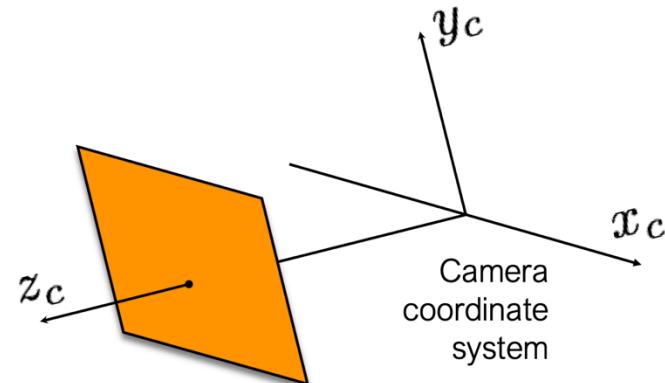
Differentiable Gaussian Rendering

**What is the representation
of a 3D Gaussian?**



\mathbf{p}, R, S

**How to project to
2D and rasterize?**



\mathbf{p}', R', S

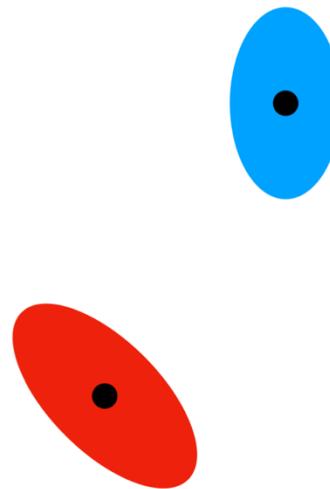
**How to model/aggregate
appearance?**

We can use the camera extrinsics to transform each 3D Gaussian to the camera frame

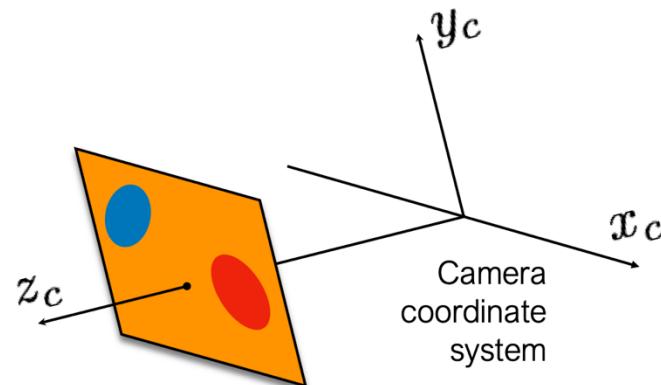
Differentiable Gaussian Rendering

$$J = \begin{bmatrix} \frac{f_x}{z} & 0 & -f_x \frac{x}{z^2} \\ 0 & \frac{f_y}{z} & -f_y \frac{y}{z^2} \end{bmatrix}$$

What is the representation of a 3D Gaussian?



\mathbf{p}', R', S



Q: What is the image-space projection of a 3D Gaussian?

A: Can approximate as a 2D Gaussian!

(EWA Volume Splatting. Zwicker et. al., 2001)

How to project to 2D and rasterize?

$$\pi(\mathbf{x}) = \mathbf{u}$$

π : Projection function for mapping 3D points to pixels

$$z \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = K \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

2D mean: $\mu_{2D} = \pi(\mu_{3D})$

2D covariance:

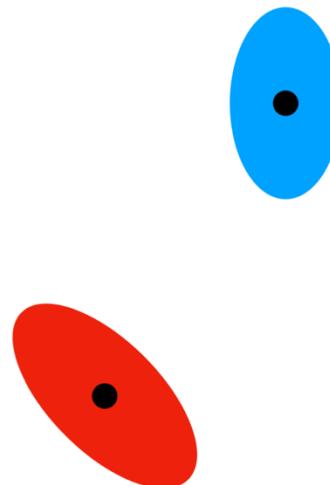
$$J = \frac{\partial \pi}{\partial \mathbf{x}}(\mu_{3D})$$

$$\Sigma_{2D} = J \Sigma_{3D} J^T$$

Slides thanks to Ioannis Gkioulekas & Shubham Tulsiani

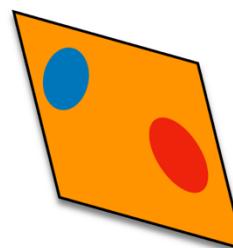
Differentiable Gaussian Rendering

**What is the representation
of a 3D Gaussian?**



$$\begin{aligned}\mu_{2D} &= \pi(\mu_{3D}) \\ \Sigma_{2D} &= J \Sigma_{3D} J^T\end{aligned}$$

**How to project to
2D and rasterize?**



**How to model/aggregate
appearance?**

1. Sort Gaussians from closest to furthest from the camera
2. For each pixel \mathbf{u} , compute opacity for each gaussian \mathcal{G}_k :

$$\bar{\alpha}_k = \alpha_k \frac{e^{-(\mathbf{u}-\mu_{2D}^k)^T (\Sigma_{2D}^k)^{-1} (\mathbf{u}-\mu_{2D}^k)}}{2\pi |\Sigma_{2D}^k|^{0.5}}$$

(In practice, can rasterize ‘blocks’ instead of entire image as
not all Gaussians influence all blocks)

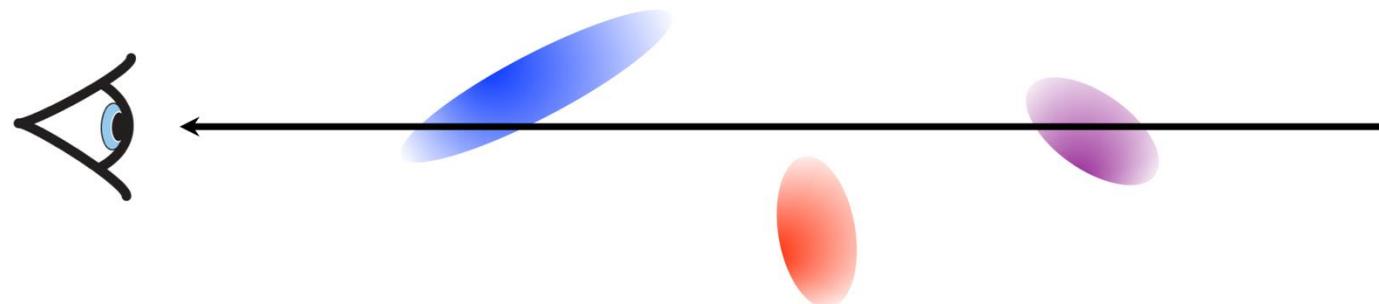
Slides thanks to Ioannis Gkioulekas & Shubham Tulsiani

Differentiable Gaussian Rendering

**What is the representation
of a 3D Gaussian?**

**How to project to
2D and rasterize?**

**How to model/aggregate
appearance?**



Compute per-Gaussian weights based on opacities of current and previous Gaussians:

$$w_k = \bar{\alpha}_k \prod_{j=1}^{k-1} (1 - \bar{\alpha}_j)$$

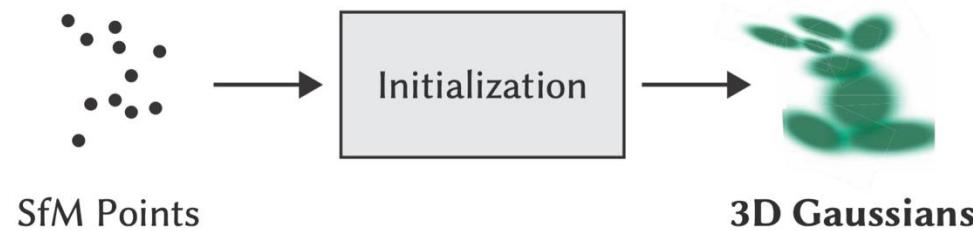
Use per-Gaussian SH coefficients and ray direction to get view-dependent color \mathbf{c}_k

Aggregate to obtain pixel color:

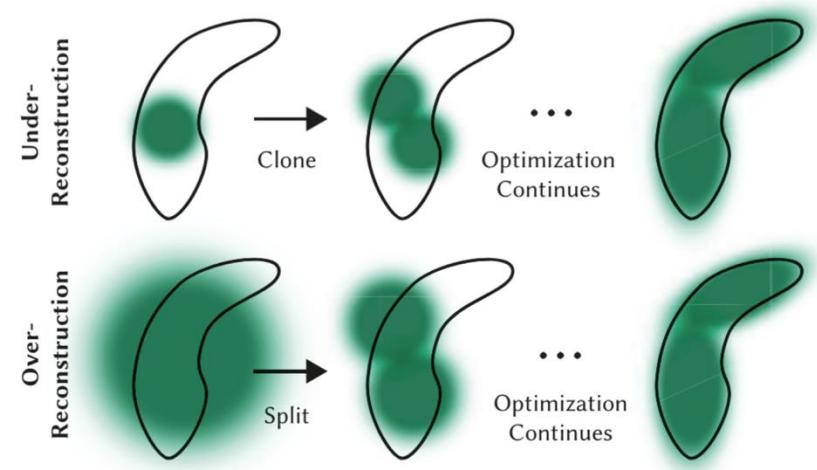
$$\mathbf{c} = \sum_k w_k \mathbf{c}_k$$

Gaussian Splatting: Bells and Whistles

+ Lots of efficient GPU optimization strategies



Initialize with sparse point cloud from SfM



Split/clone
Gaussians based on
heuristics

Bottom line

- Gaussian Splats still estimate volumetric rendering (same alpha-compositing, just with gaussians)
- MUCH faster because no sampling, no neural nets, rasterization
- From 30 sec / frame for 800x800 image original NeRF
 - Plenoxels/InstantNGP: 10-30FPS 1920x1080 image
 - GS: ~100-200+ FPS at 1920×1080
- Best balance of quality and speed – current status quo



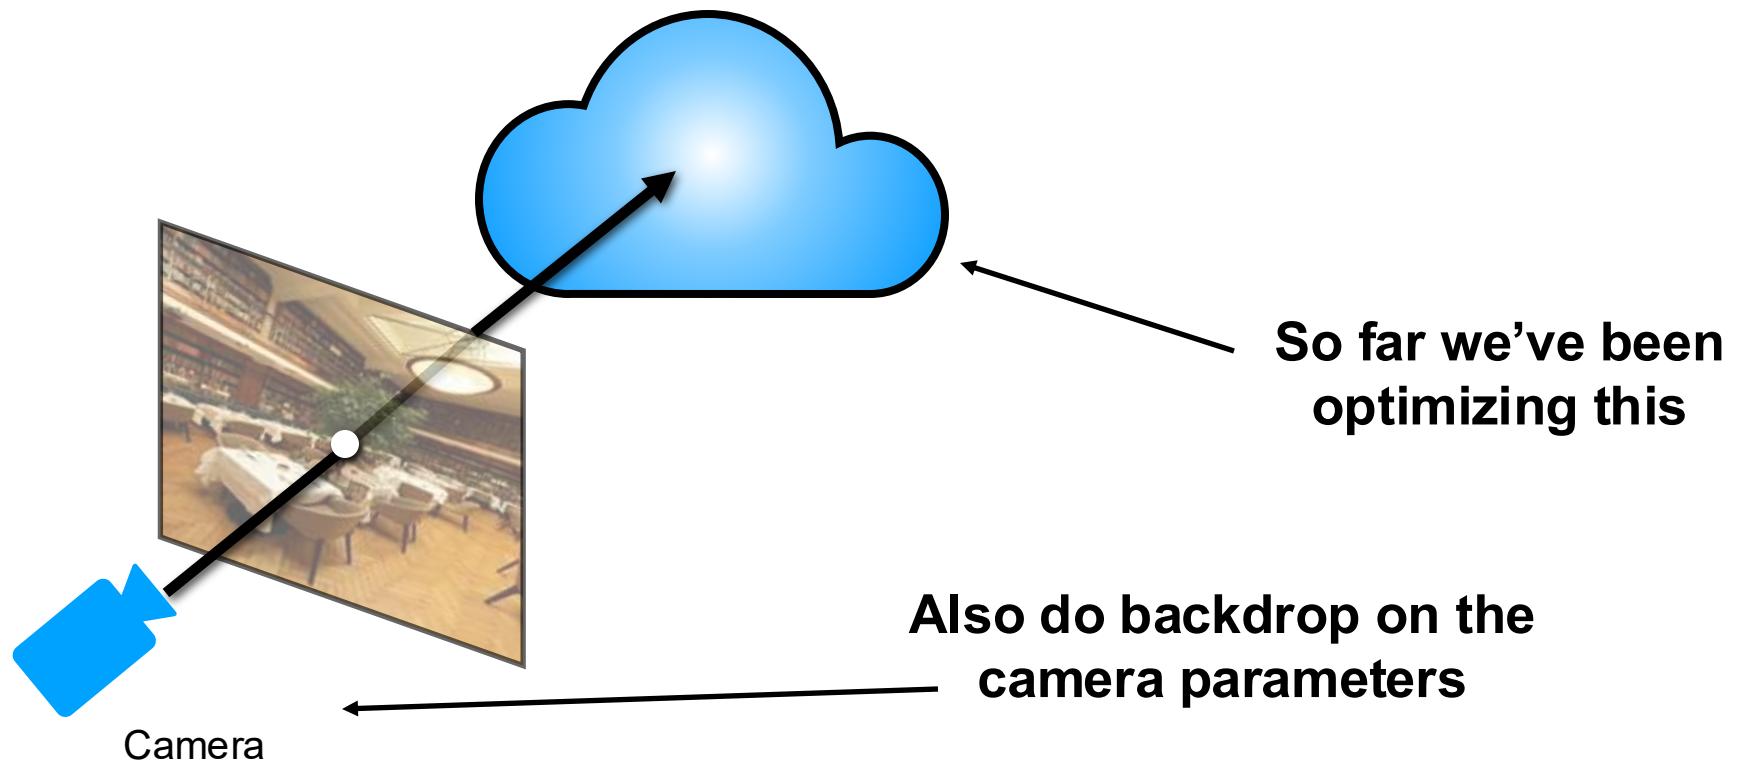
Ours



MipNeRF360

Camera Quality

Small noise in the camera can be made robust by also optimizing the camera



No Pose Optimization



Block-NeRF



Camera Optimization

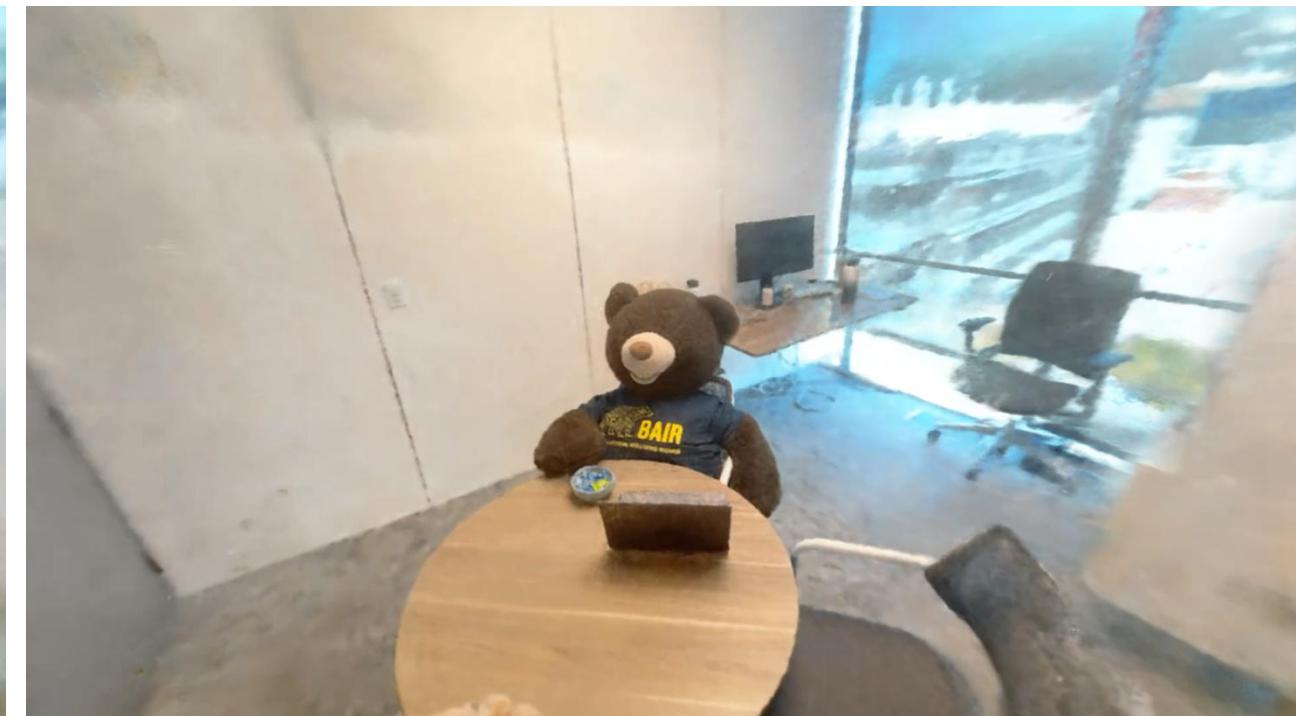
Small noise in the results can be improved

Starting from scratch is still an active area of research [Barf Lin et al. 2021, NeRF— ...]

Noisy Camera from IMU/Lidar



Result with Camera Optimization



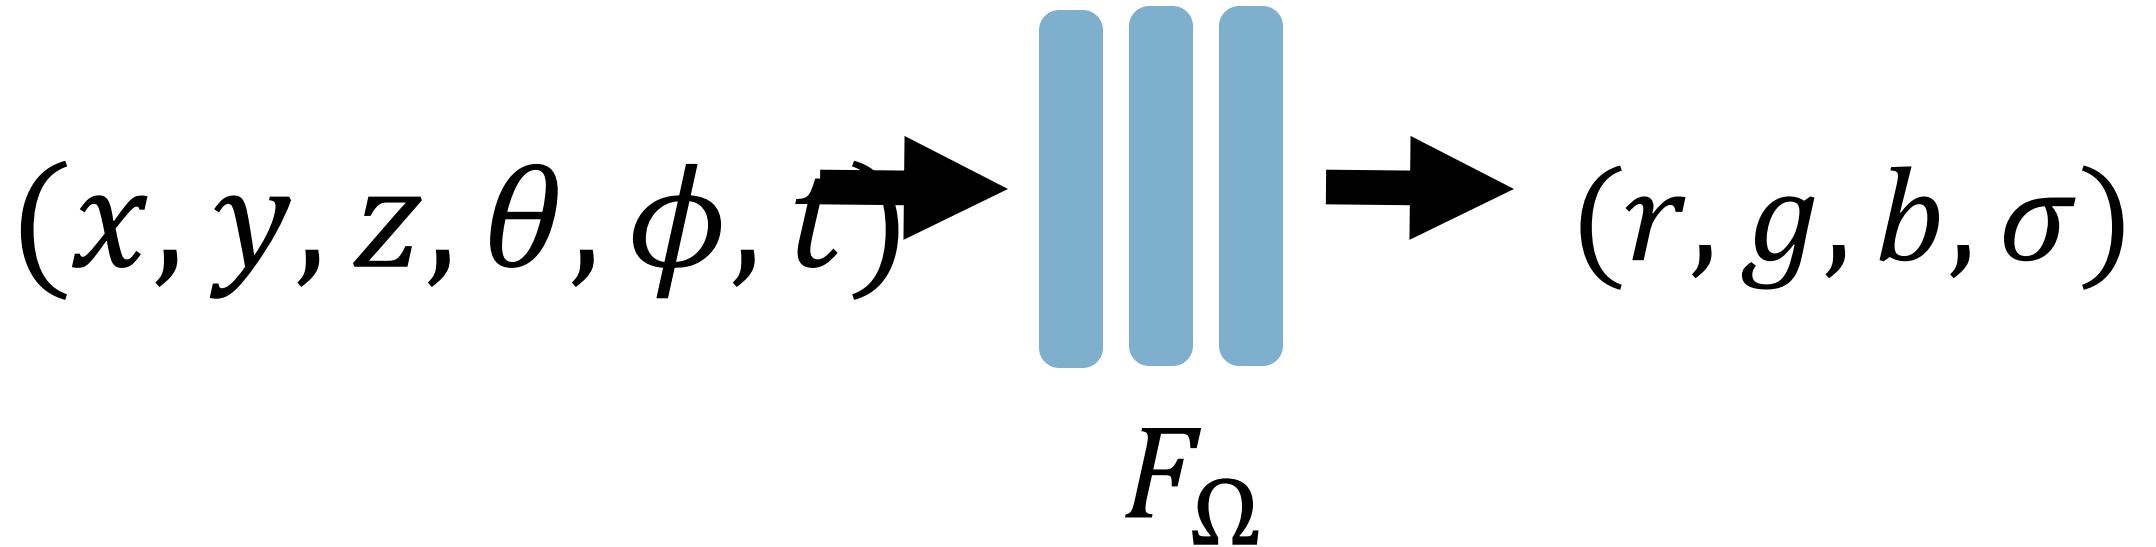
The Dynamic World



Holy grail

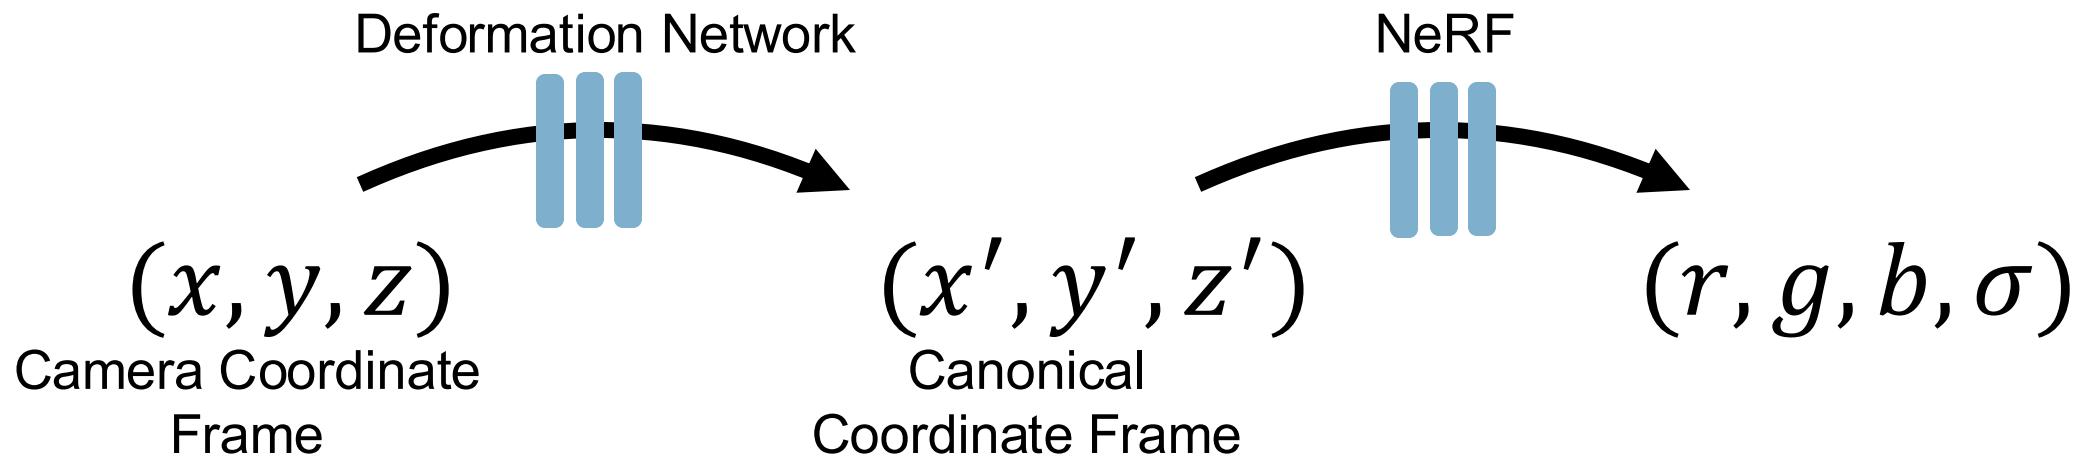
- Dynamic Novel View Synthesis from Monocular Camera
- Very difficult! Extremely under constrained problem

Simple baseline for adding time



Hard without simultaneous multiple view!

Through a deformation network



Still very under constrained

Dynamic View Synthesis: Monocular is hard



HyperNeRF (Ours)

D-NeRF [Pumarola et al. CVPR 2021], NSFF [Li et al., CVPR 2021], HyperNeRF [Park et al. SIGASia 2021]....

- But performance on in-the-wild monocular capture still far [Gao et al. NeurIPS 2022]



train view

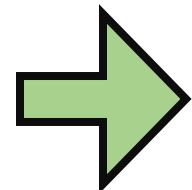
Nerfies



What if we knew how they deform?



HMMR, Kanazawa et al.
CVPR 2019



Other kinds of dynamic changes

Appearance Changes

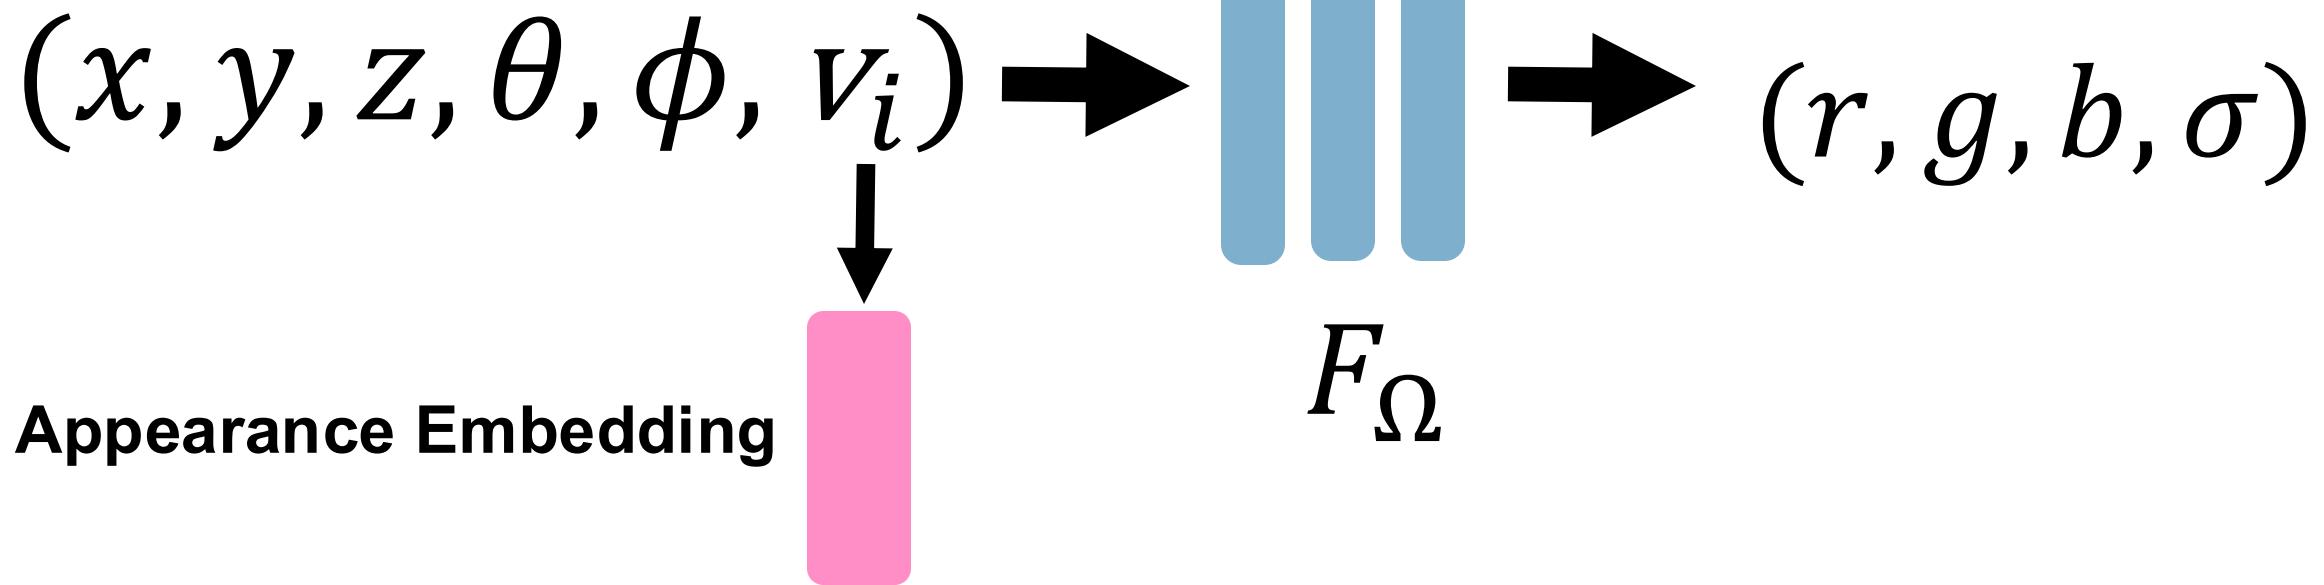
Exposure differences

Lighting changes (day, night)..

Clouds passing by..



Appearance Embedding: Pretty Robust Solution



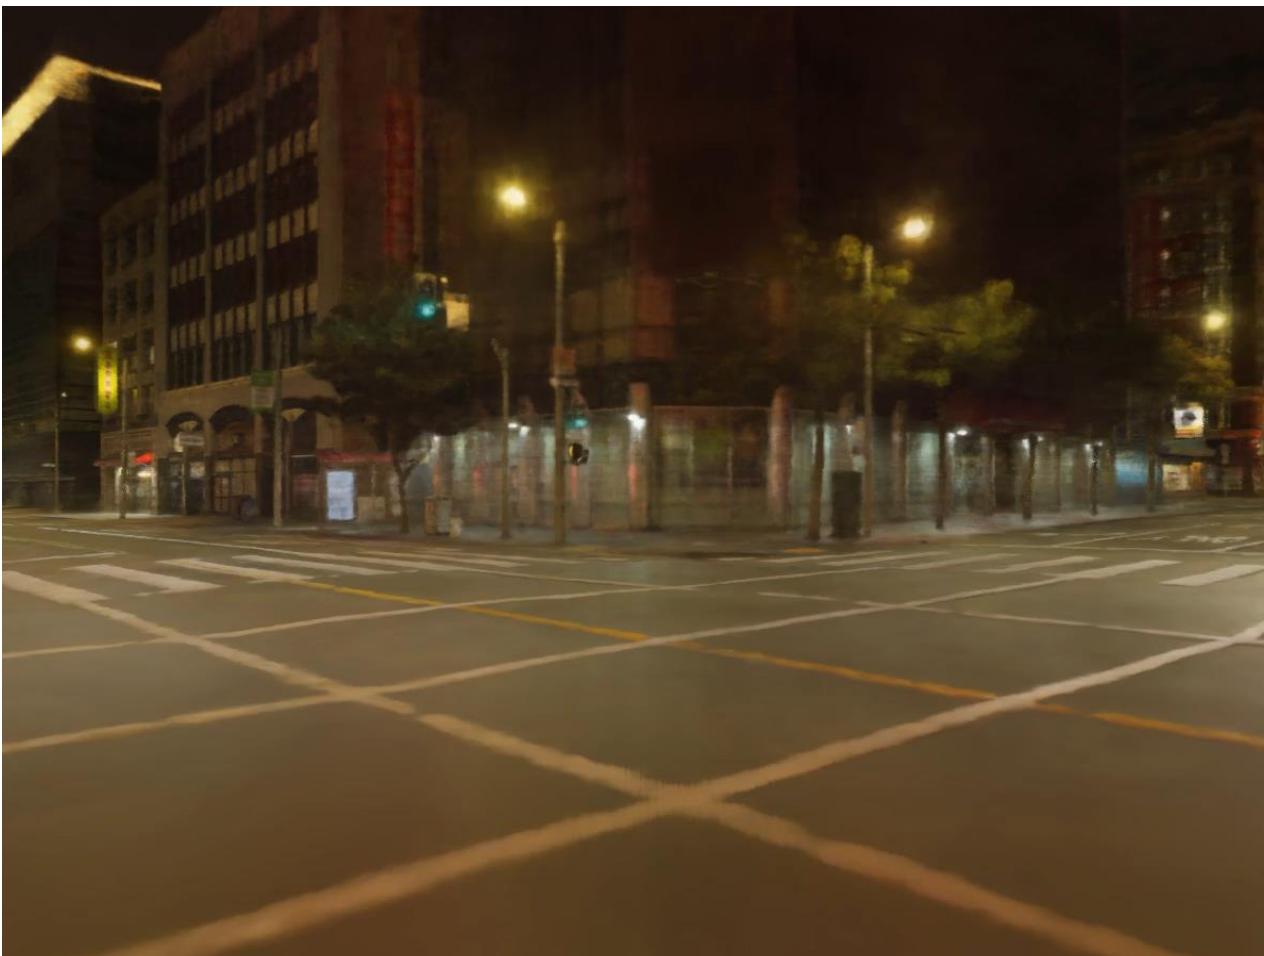
N-dim vector

Optimized *per* image: “Auto-Decoding”

ie GLO: Generative Latent Optimization [[Bojanowski et al.](#)] ICML 2018]

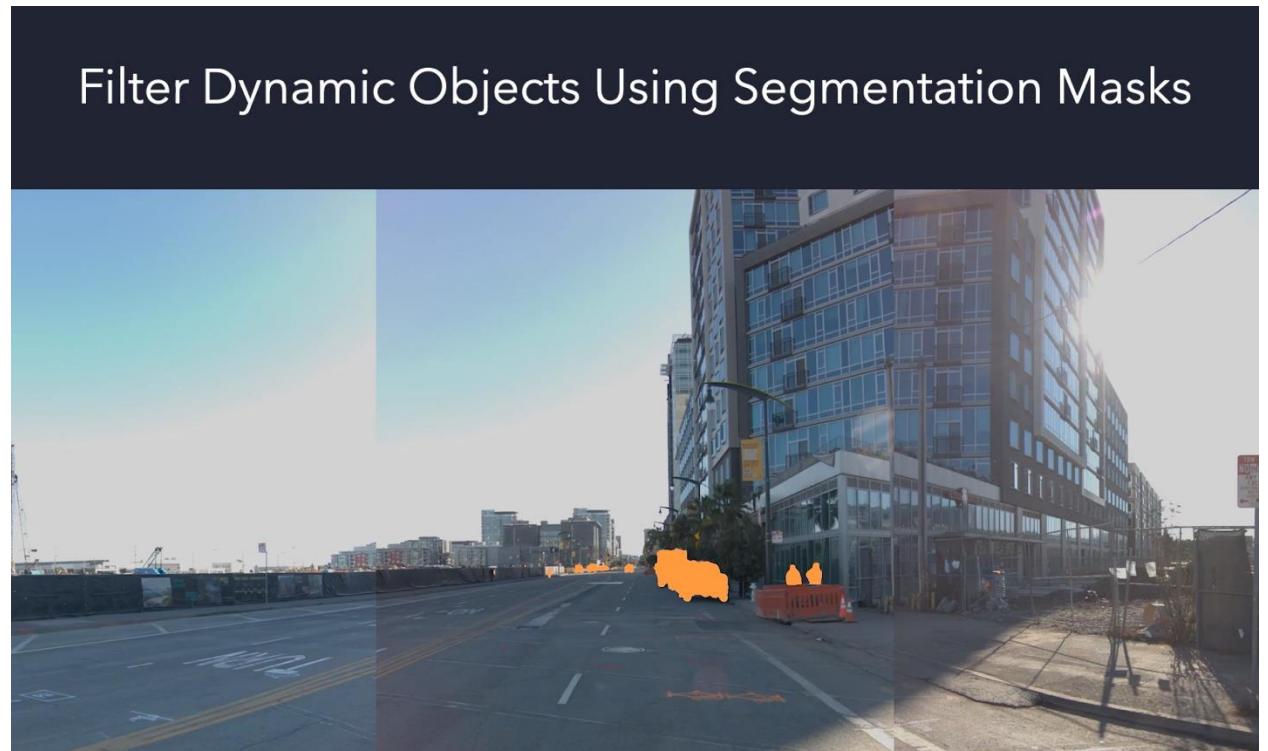
Appearance Changes

Appearance Encoding is Effective



Transient objects

- Happens all the time! People moving around, interacting with the world
- Difficult! Problem of Grouping
 - how do you know which part is connected or
 - Can use two NeRFs, one global, one per-image, but this often leads to degenerate solutions
- Current solution: Ignore (mask out)



Why is dynamic scenes hard?

- Unless you have a light dome
- Essentially you only have a single-view

Building & Reusing Prior Knowledge

Machine Learning

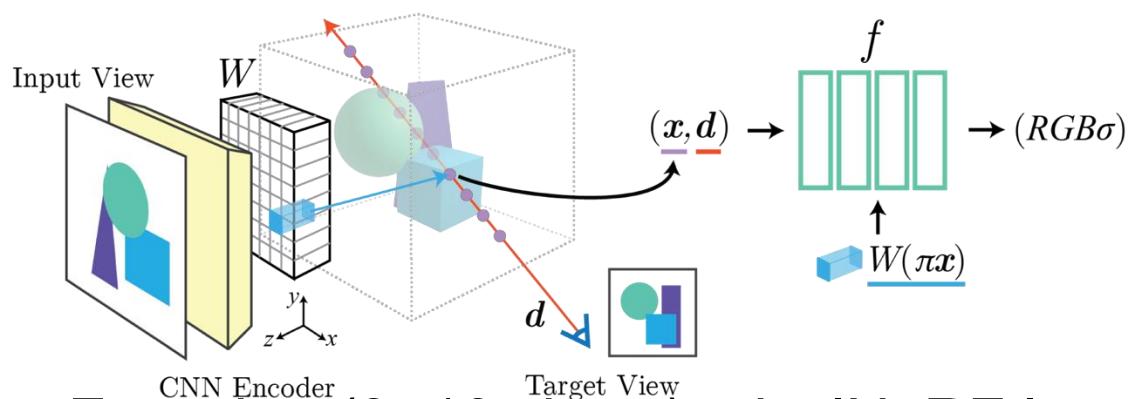
NeRF is per-scene optimization

- We need lots of images to get good view synthesis!!
- Also there's no knowledge reused from prior scene reconstructions
- How to bring learning in the picture?



Few-shot NeRF

- One-shot (single-view): pixelNeRF [Yu et al. CVPR'19]



- Few-shot (3~10 views): pixelNeRF, IBRNet [Wang et al. CVPR'21], MVSNet [Chen et al. ICCV'21], etc...
- Challenging for predicting completely unseen real scenes

- How to deal with the multi-modal nature of the problem??



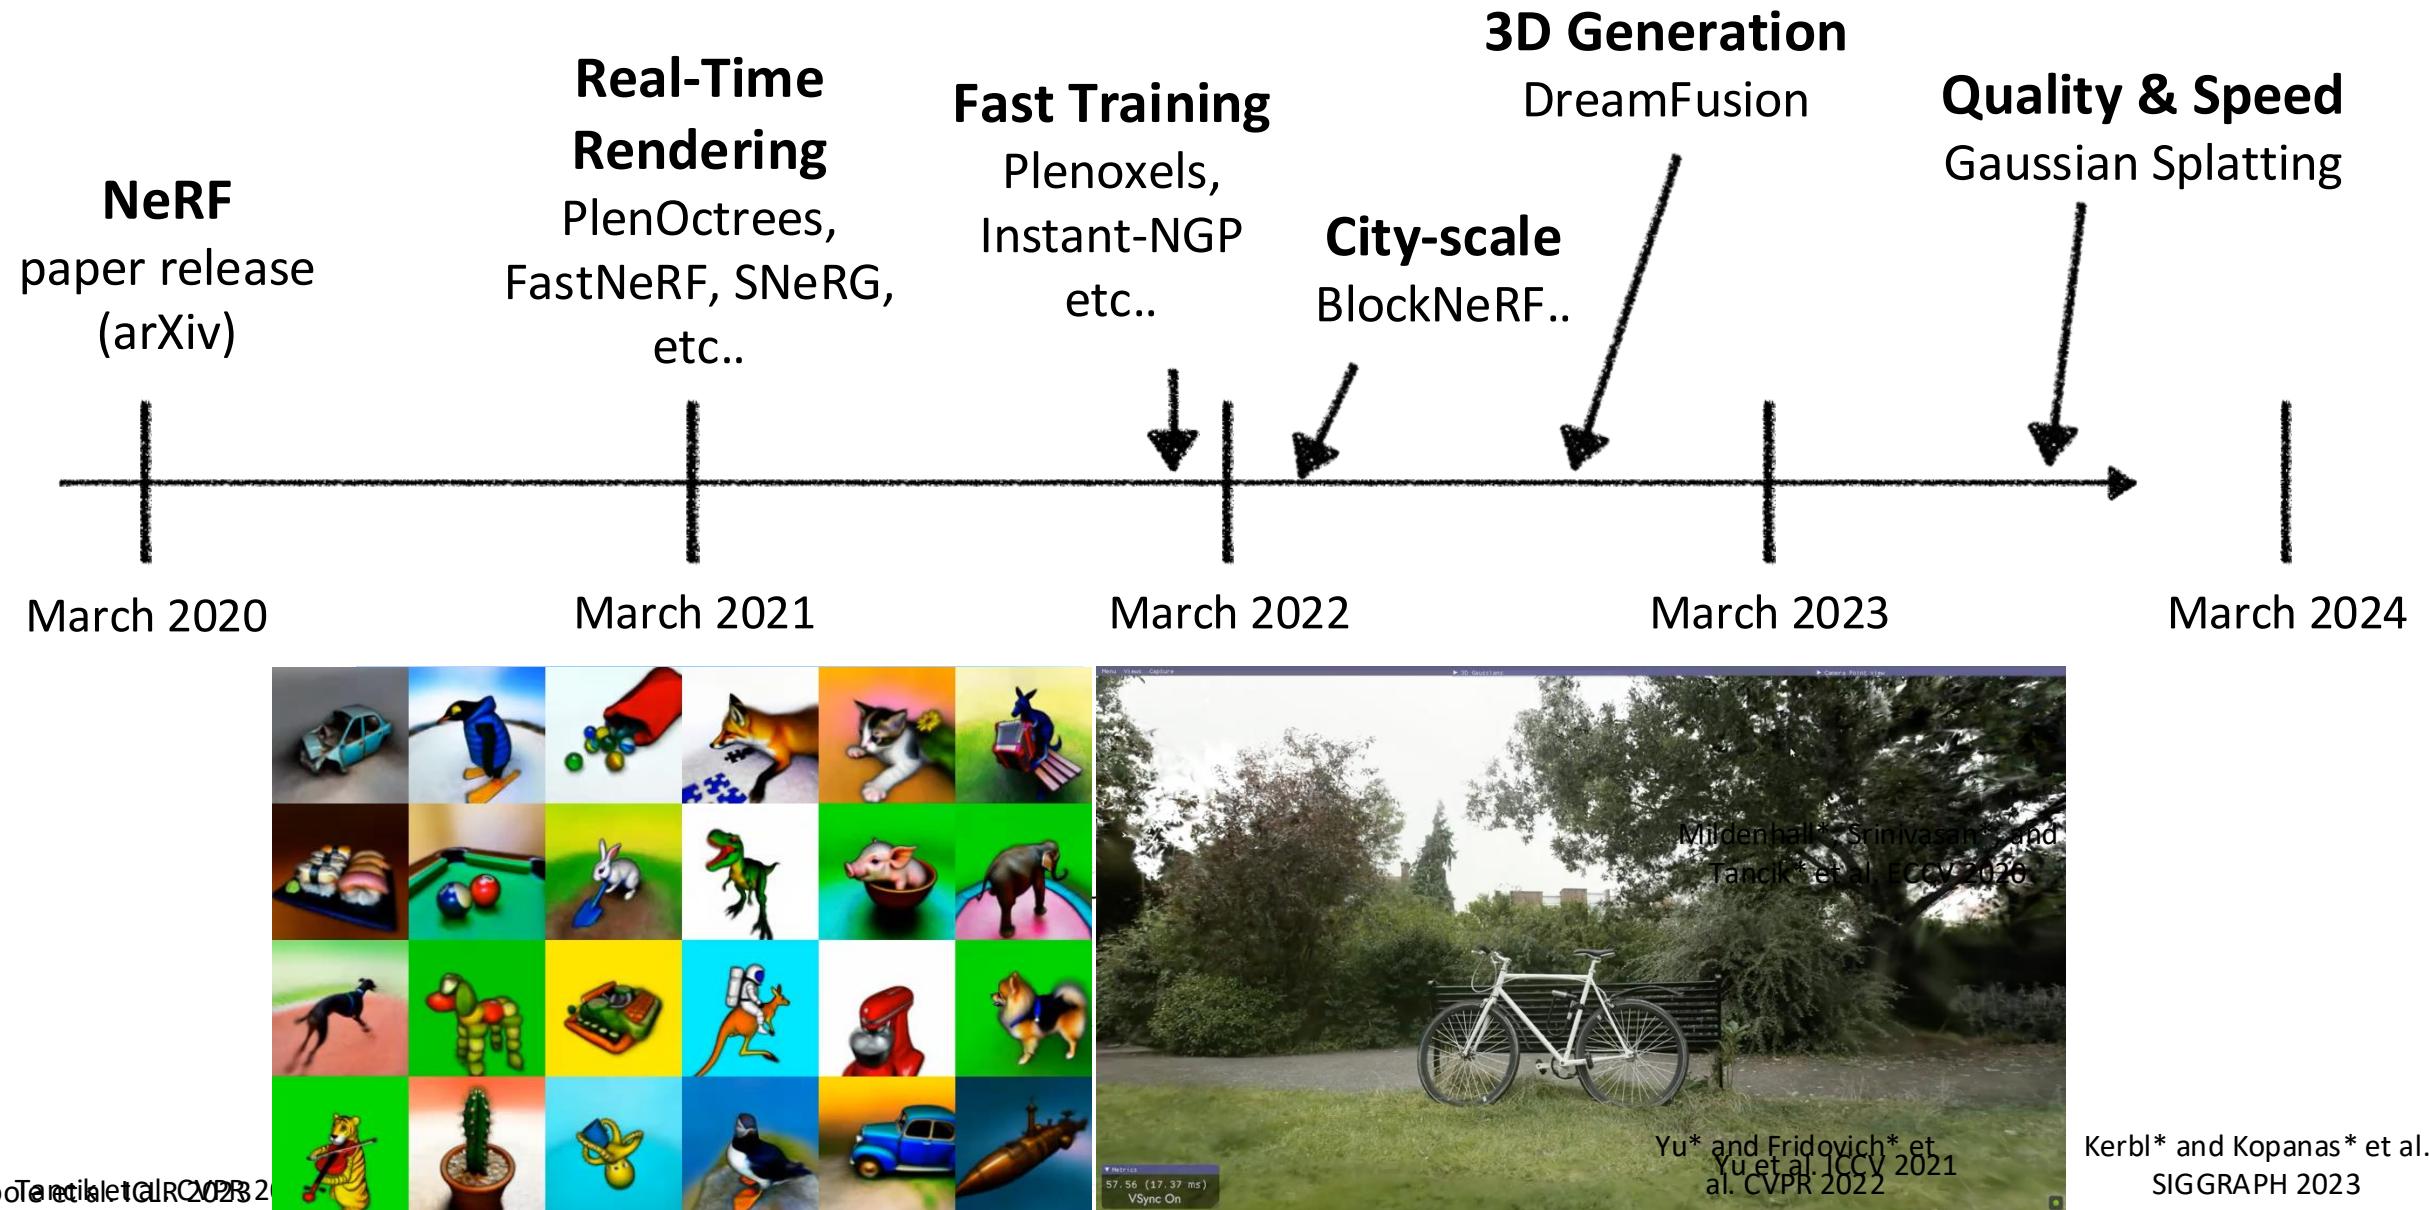
IBRNet

Data is the bottleneck

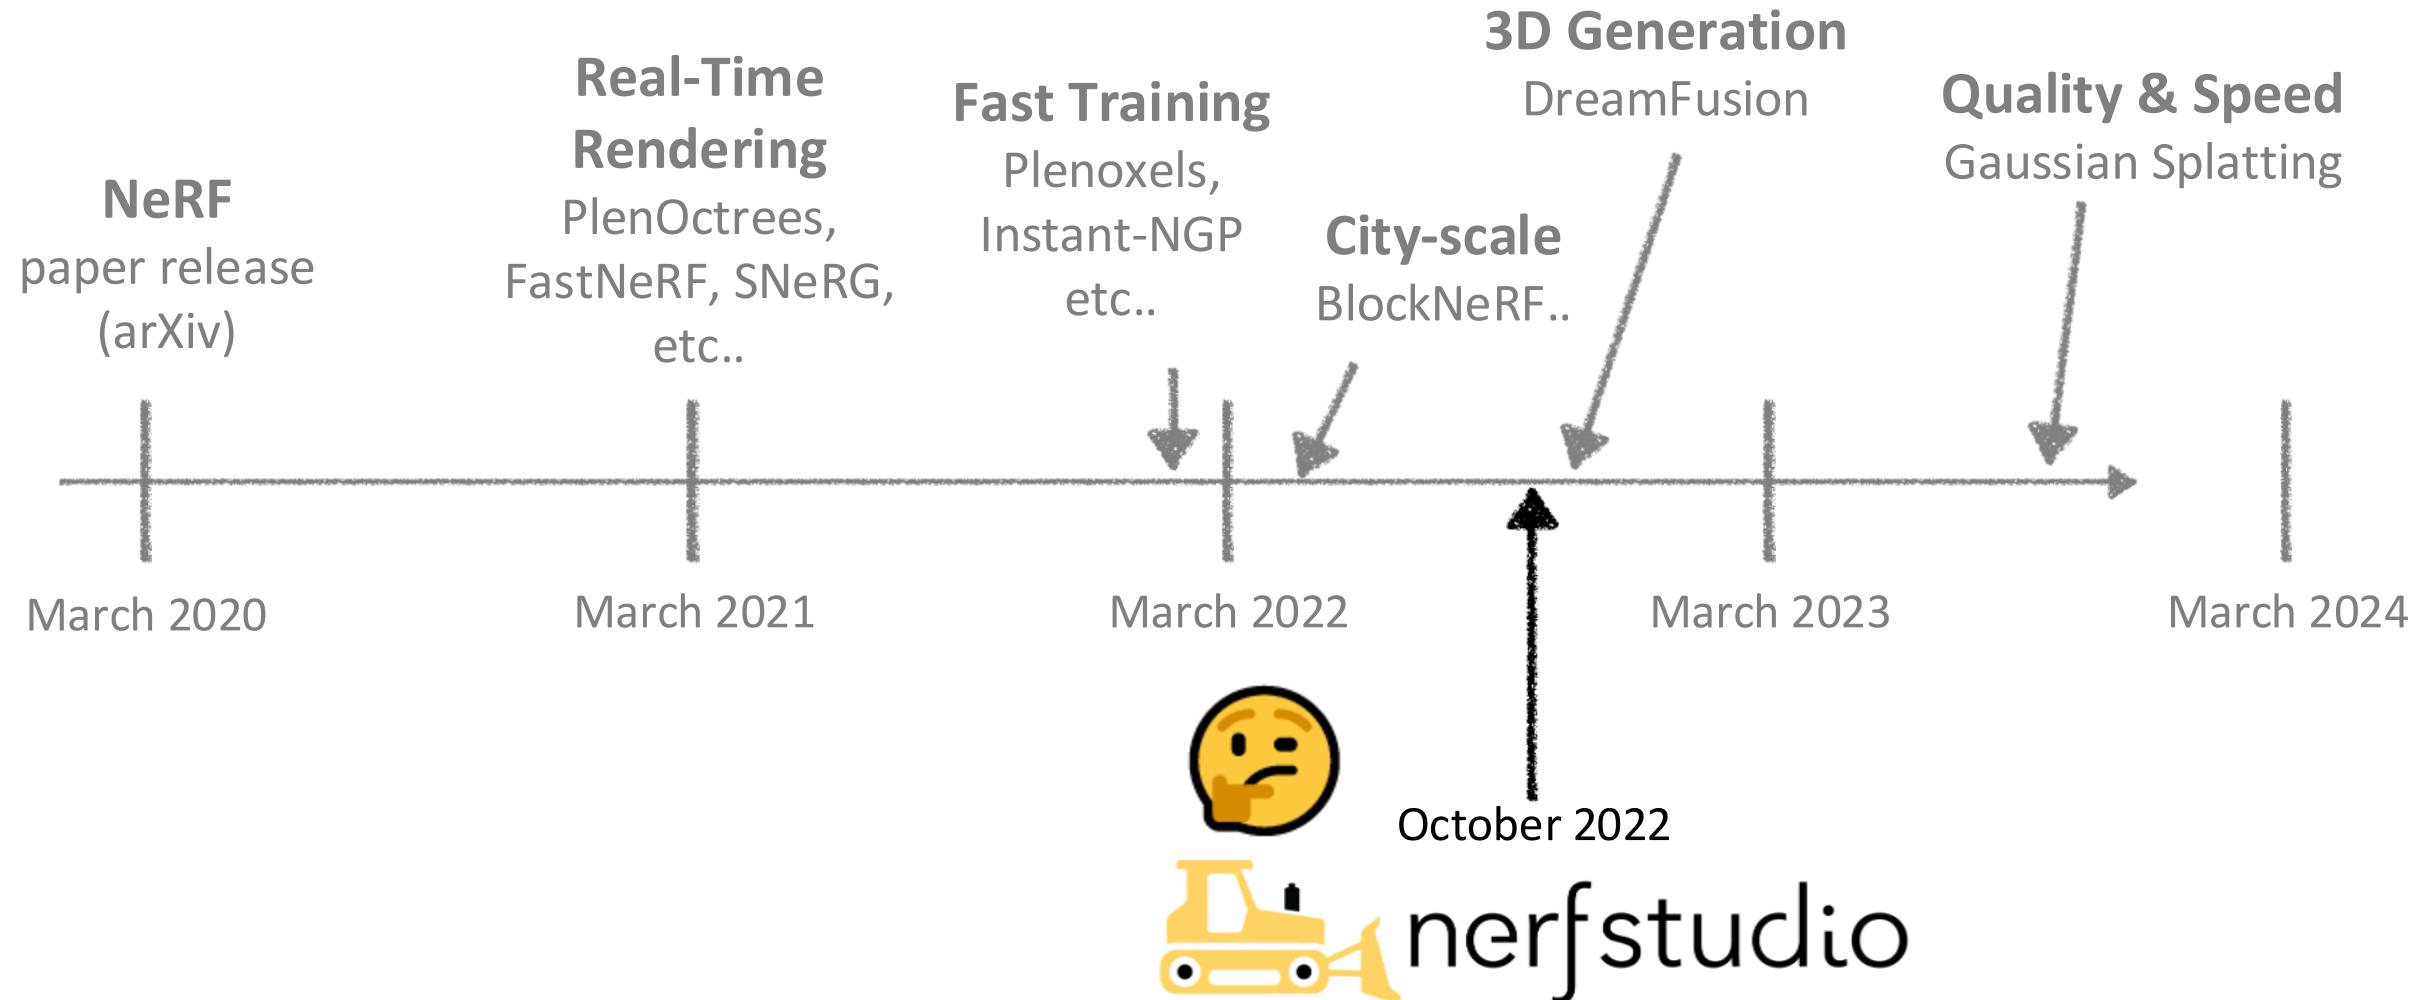
- Large-scale Real-World Multi-view Data is hard to collect:
CO3D [Reizenstein ICCV 2021]
- A lot to learn from other single-view 3D prediction models:



Time Line



Time Line



A Modular Framework for NeRF Development

Matthew Tancik*, Ethan Weber*, Evonne Ng*, Ruilong Li, Brent Yi, Justin Kerr,
Terrance Wang, Alexander Kristoffersen, Jake Austin, Kamyar Salahi, Abhik Ahuja,
David McAllister, Angjoo Kanazawa

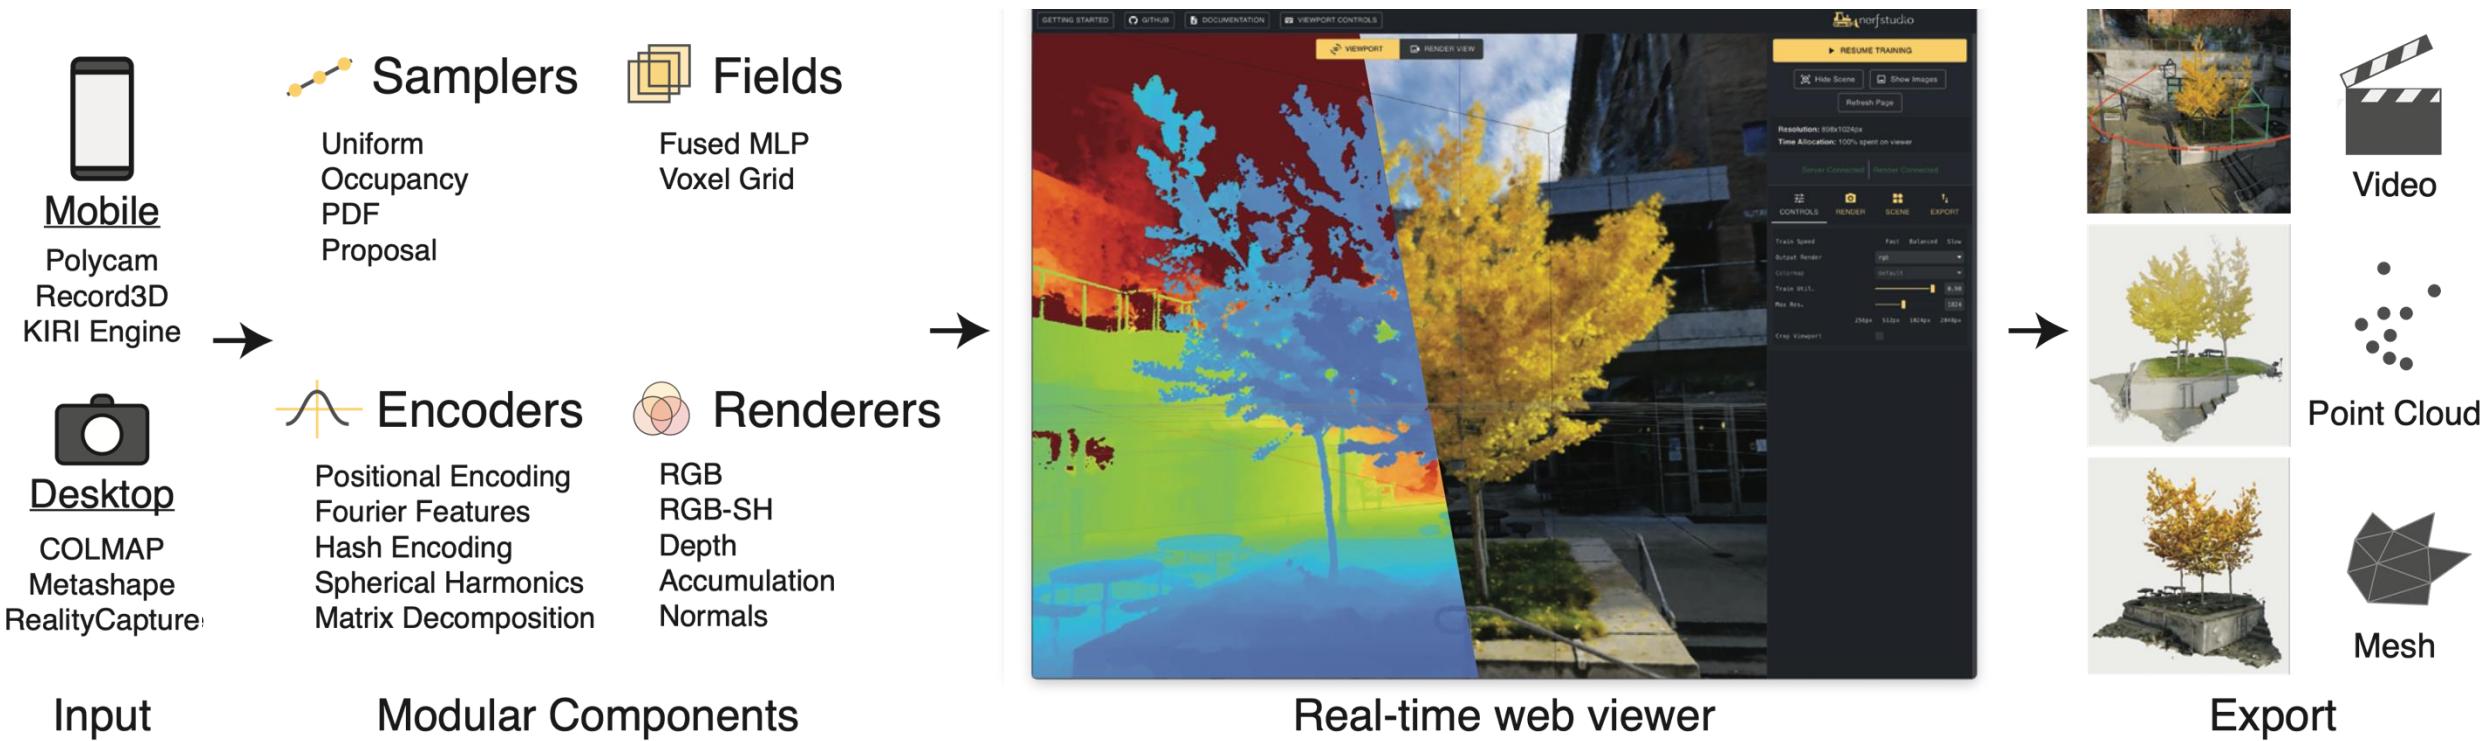
+143 additional Github collaborators
SIGGRAPH 2023

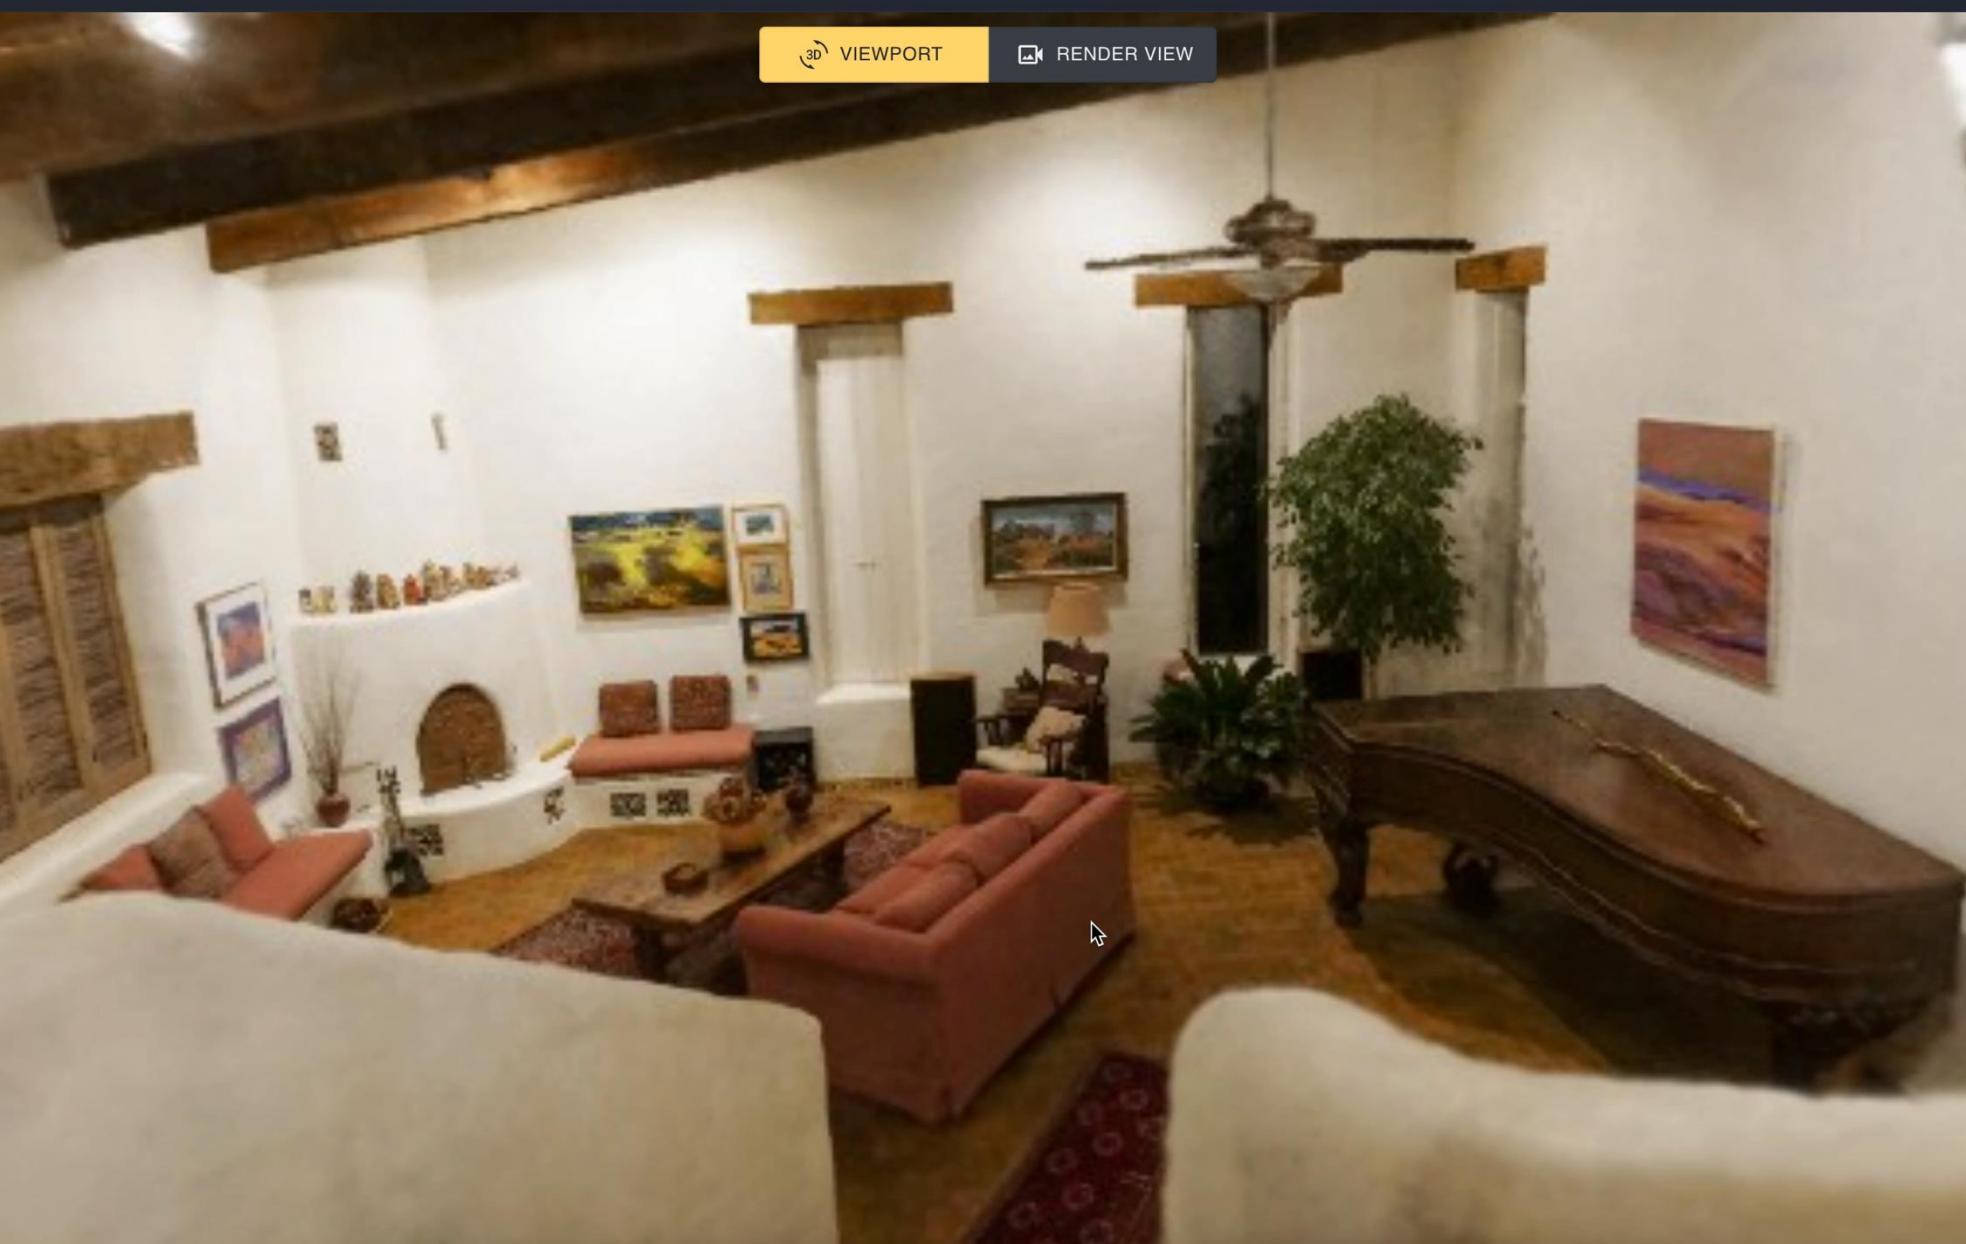
Design Goals

Easy to:

Use
Develop
Learn

An End-to-End Framework



[GETTING STARTED](#)[GITHUB](#)[DOCUMENTATION](#)[VIEWPORT CONTROLS](#)

3D VIEWPORT

RENDER VIEW

TRAINING COMPLETE

[Show Scene](#)[Hide Images](#)[Refresh Page](#)

Iteration: 29999

Resolution: 320x512px



CONTROLS



RENDER



SCENE



EXPORT

Train Speed

Balanced

Output Render

rgb

Colormap

default

Train Util

0.85

Max Res

512

Crop Viewport



Data Pipelines



Onboarding Pipelines

- COLMAP
- Polycam
- Record3D
- MetaShape
- RealityCapture
- Kiri Engine

Easy to Develop

Sampling

Fields & Encoders

Volumetric
Rendering

Pythonic and Modular

Easy to Develop

Sampling

- Uniform
- Occupancy
- PDF
- Proposal
- Spacing Fn

Fields & Encoders

- Positional Encoding
- Fourier Features
- Hash Encoding
- Spherical Harmonics
- Matrix Decomposition
- Fused MLP
- Voxel Grid

Volumetric Rendering

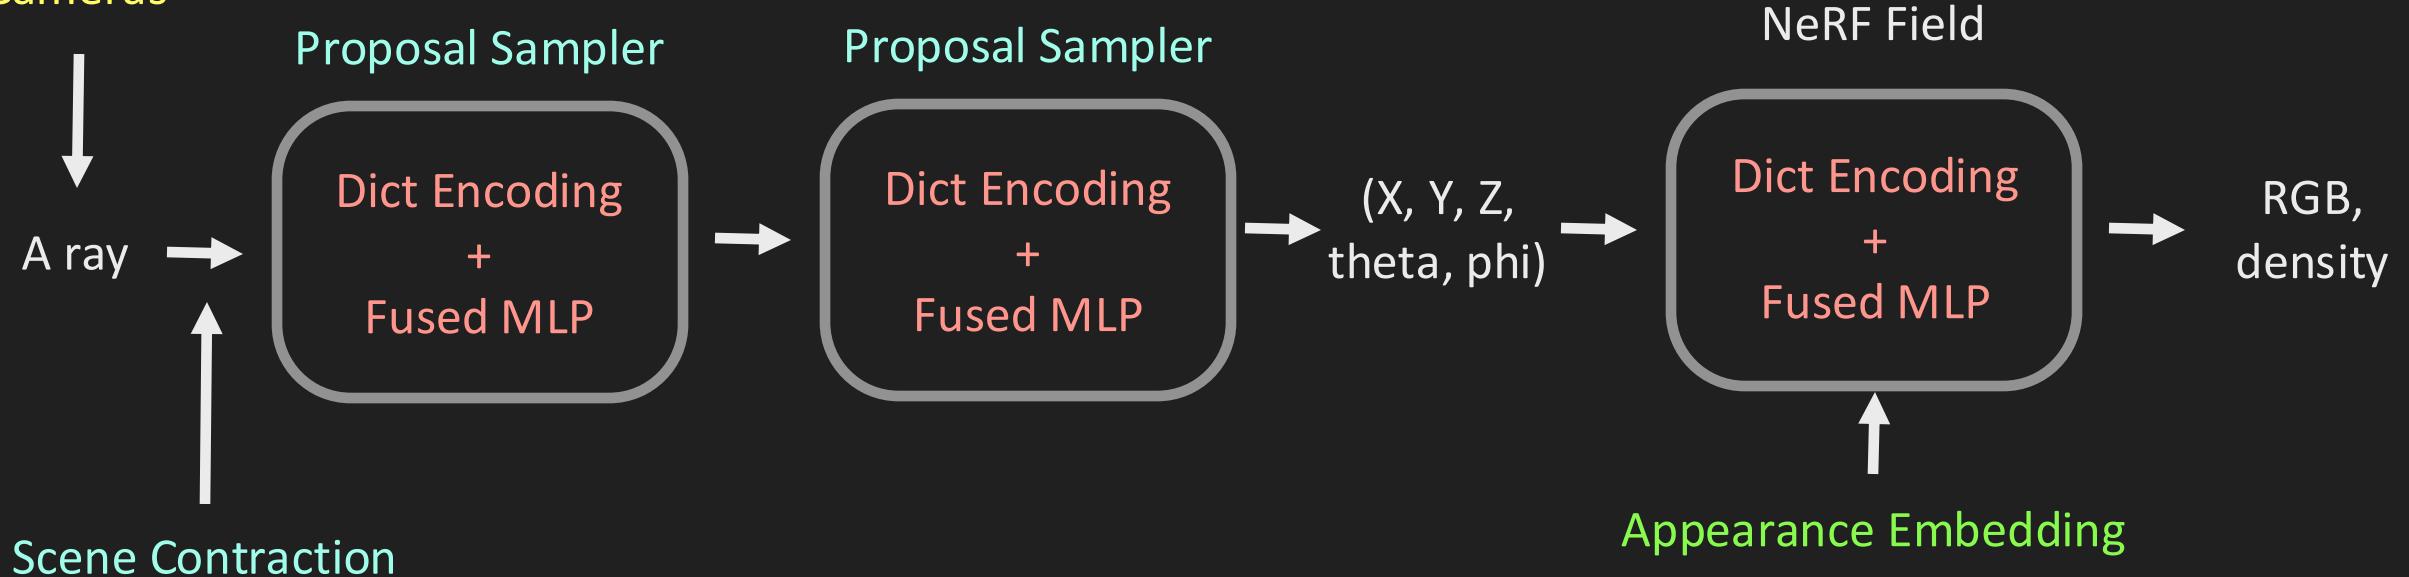
- RGB
- RGB-SH
- Depth
- Accumulation
- Normals

Pythonic and Modular

Nerfacto

Striking the balance between performance & easy development

Optimized
Cameras



Current Model

mip-NeRF 360, NeRF-W,
NeRF--/BaRF, InstantNGP

An Active Discord Community

nerfstudio

Events

rules

moderator-only

TEXT CHANNELS +

announcements

landing-pad

general

renders

troubleshooting

capture-tips

random

feature-requests

research

questions

VOICE CHANNELS +

General

renders

Why did you downscale? Because of memory or was the result better?

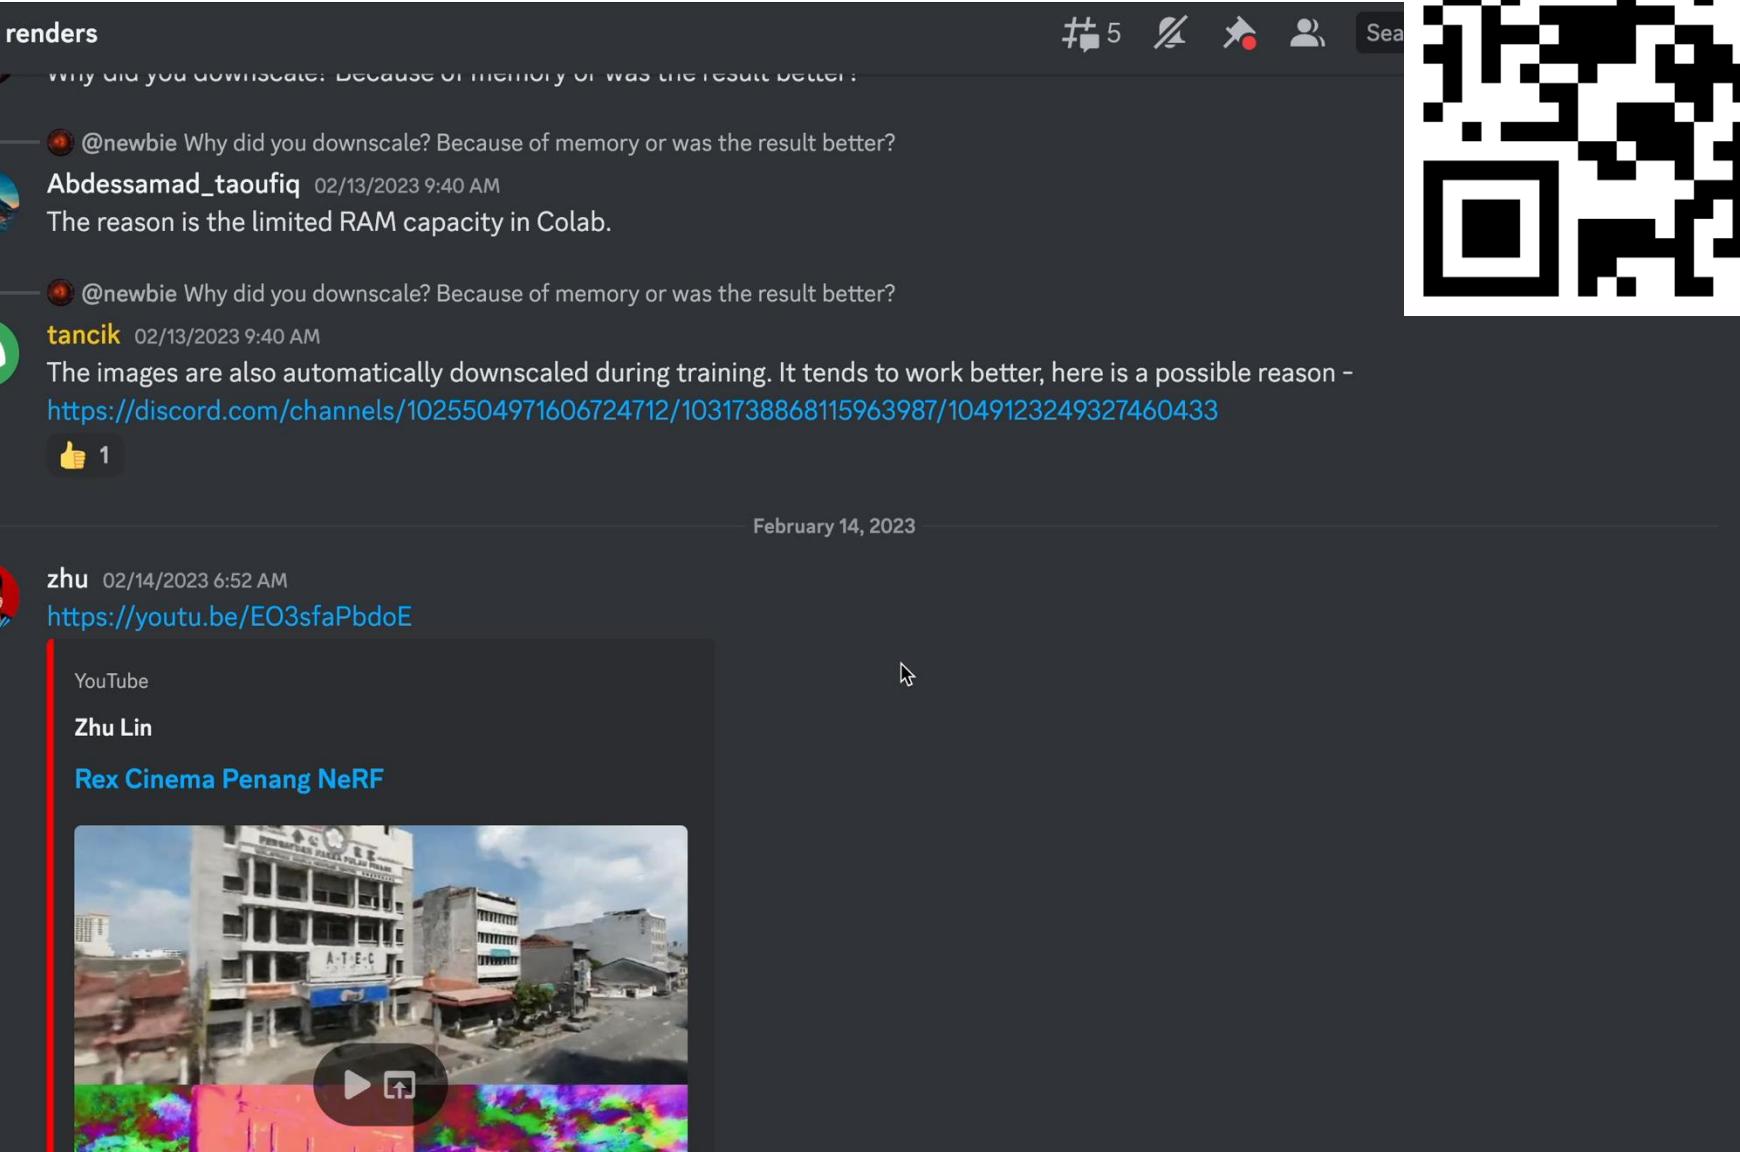
@newbie Why did you downscale? Because of memory or was the result better?
Abdessamad_taoufiq 02/13/2023 9:40 AM
The reason is the limited RAM capacity in Colab.

@newbie Why did you downscale? Because of memory or was the result better?
tancik 02/13/2023 9:40 AM
The images are also automatically downscaled during training. It tends to work better, here is a possible reason -
<https://discord.com/channels/1025504971606724712/1031738868115963987/1049123249327460433>

February 14, 2023

zhu 02/14/2023 6:52 AM
<https://youtu.be/EO3sfaPbdoE>

YouTube
Zhu Lin
Rex Cinema Penang NeRF









Nerfacto



Nerfacto-huge





Viewer

[GETTING STARTED](#)[GITHUB](#)[DOCUMENTATION](#)[VIEWPORT CONTROLS](#)

3D VIEWPORT

RENDER VIEW



▶ RESUME TRAINING

 Hide Scene Show Images

Refresh Page

Iteration: 29999

Resolution: 200x319px

CONTROLS

RENDER

SCENE

EXPORT

Train Speed

Balanced

Output Render

rgb

Colormap

default

Train Util

0.85

Max Res

512

Crop Viewport

Background color

#17234a

Crop Min

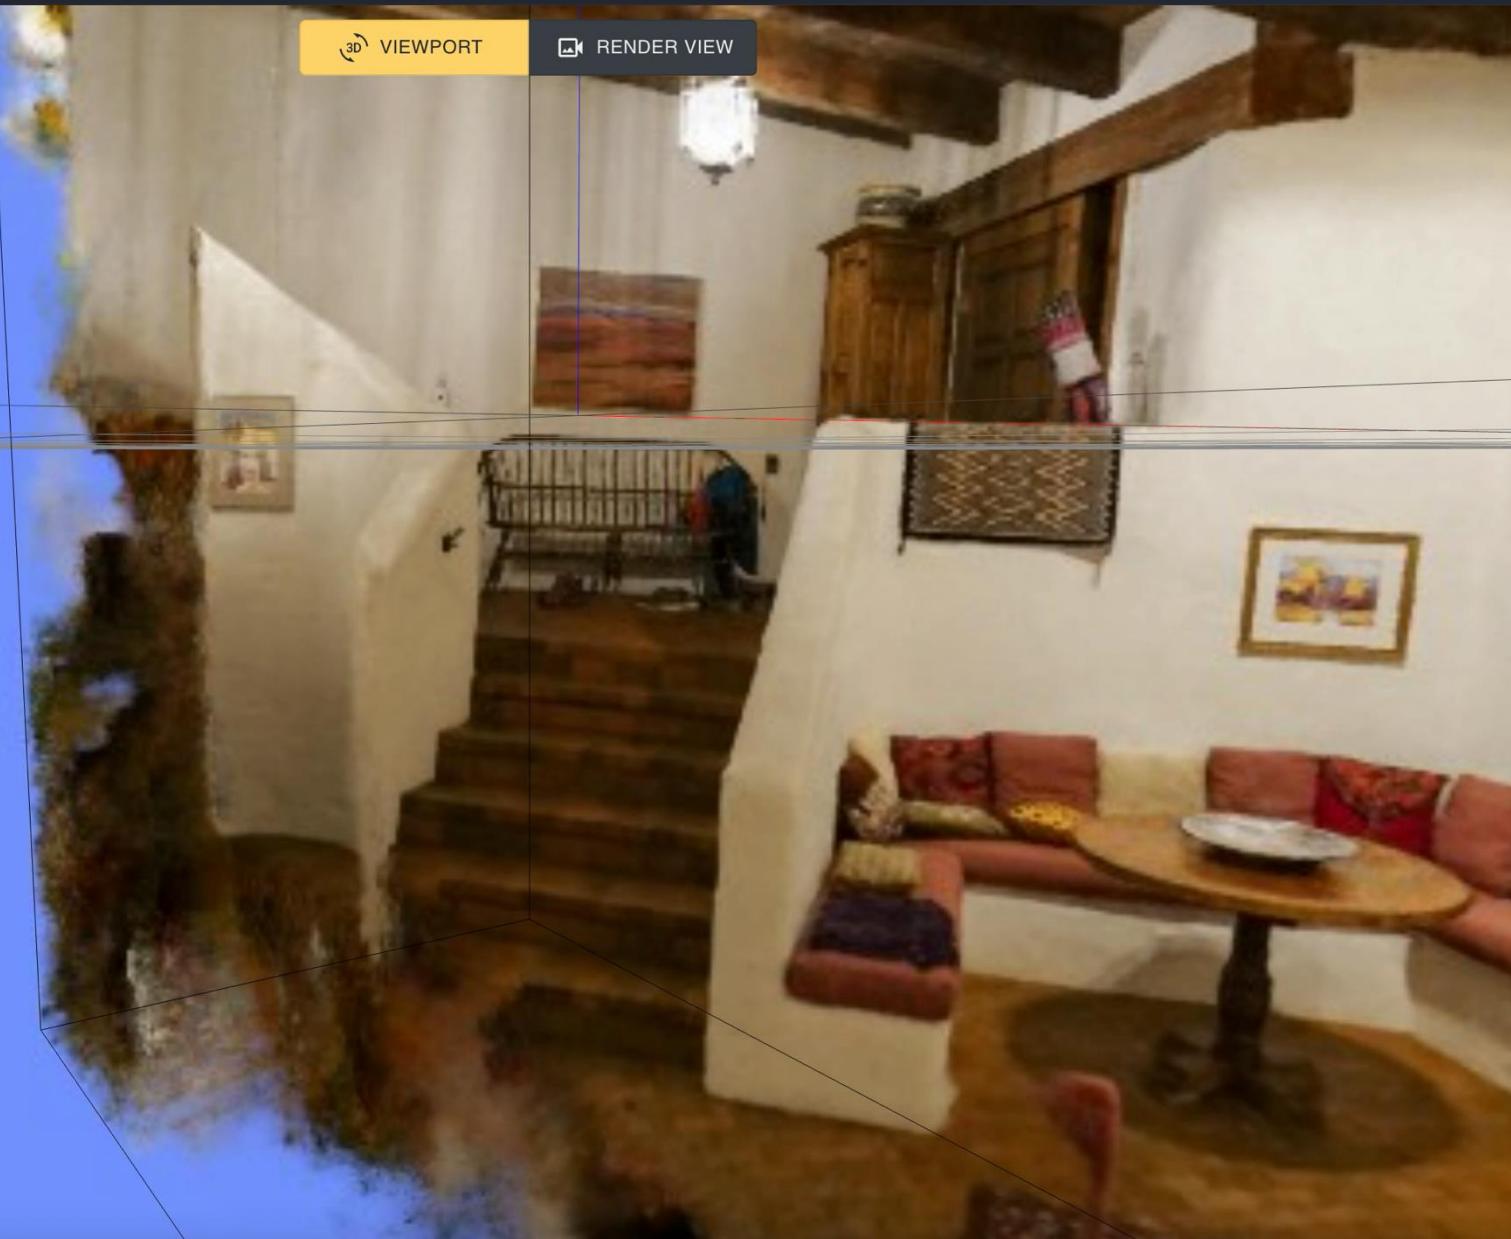
X: 0.05 Y: -1.40 Z: -1.00

Crop Max

X: 1.05 Y: 1.65 Z: 1.00

[GETTING STARTED](#)[GITHUB](#)[DOCUMENTATION](#)[VIEWPORT CONTROLS](#)

nerfstudio



3D VIEWPORT

RENDER VIEW

▶ RESUME TRAINING

 Hide Scene Show Images

Refresh Page

Iteration: 29999

Resolution: 320x512px

CONTROLS

RENDER

SCENE

EXPORT

Train Speed

Balanced

Output Render

rgb

Colormap

default

Train Util

0.85

Max Res

512

Crop Viewport

 Background color

#7091fd

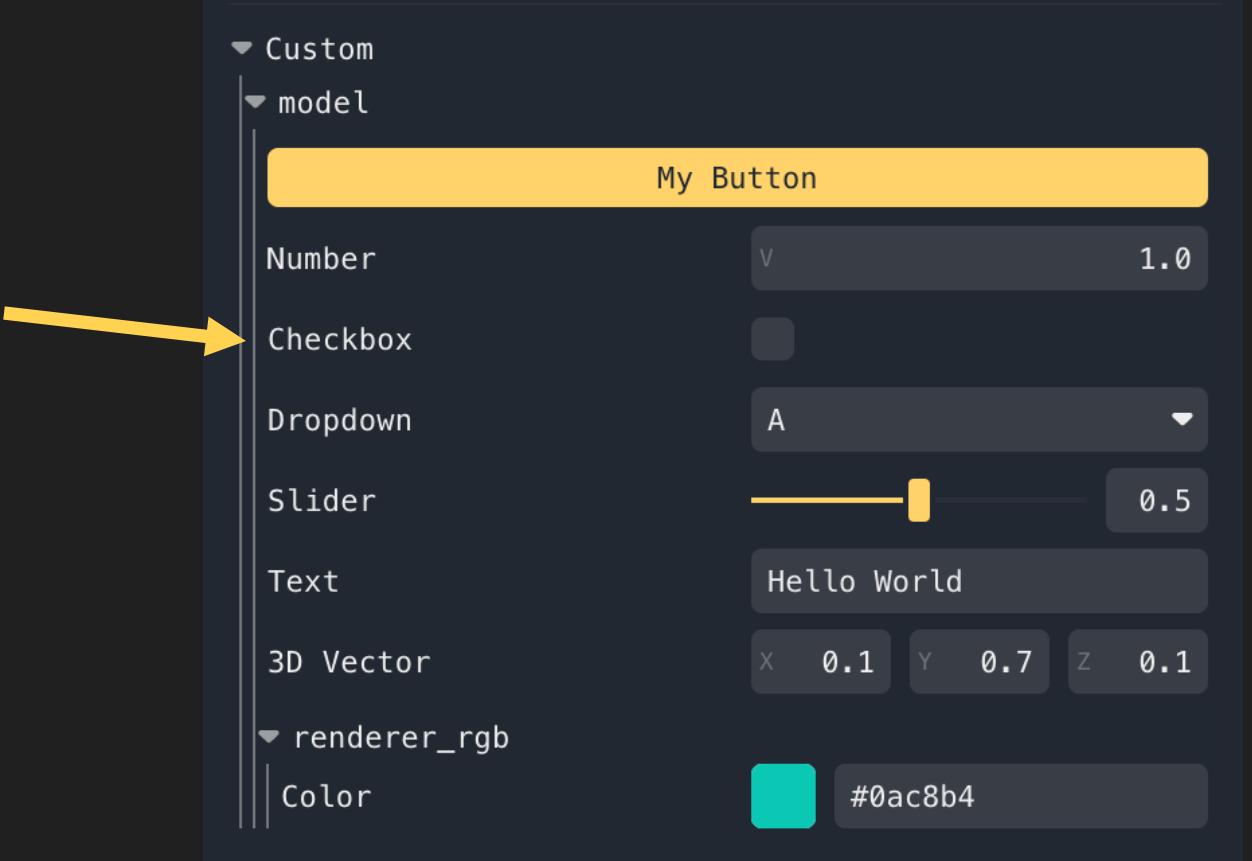
| Crop Min

| Crop Max



Custom Interactivity

```
self.checkbox = ViewerCheckbox(name="Checkbox", default_value=False)  
:  
current_value = self.checkbox.value
```

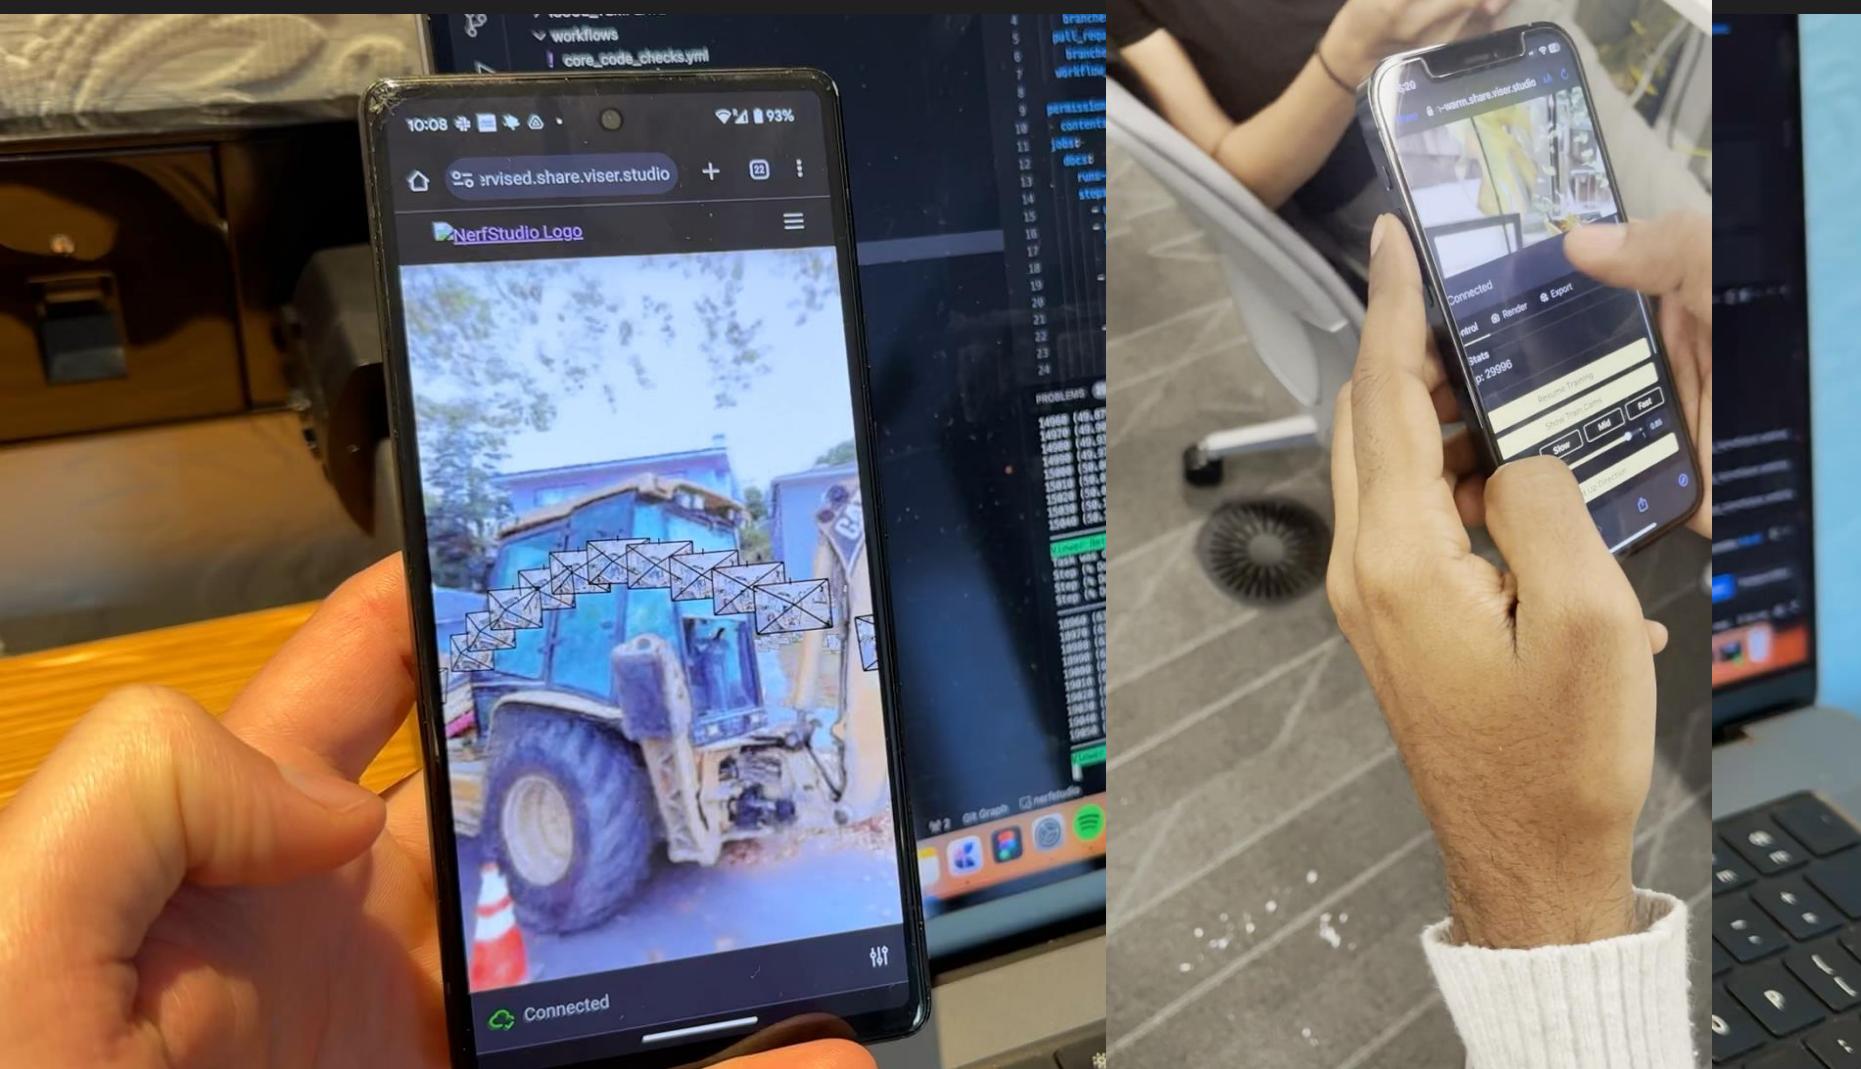




Coming Soon: Viser Integration

Python library for web-based 3D visualization

viser.studio



Shareable Links

Mobile Support



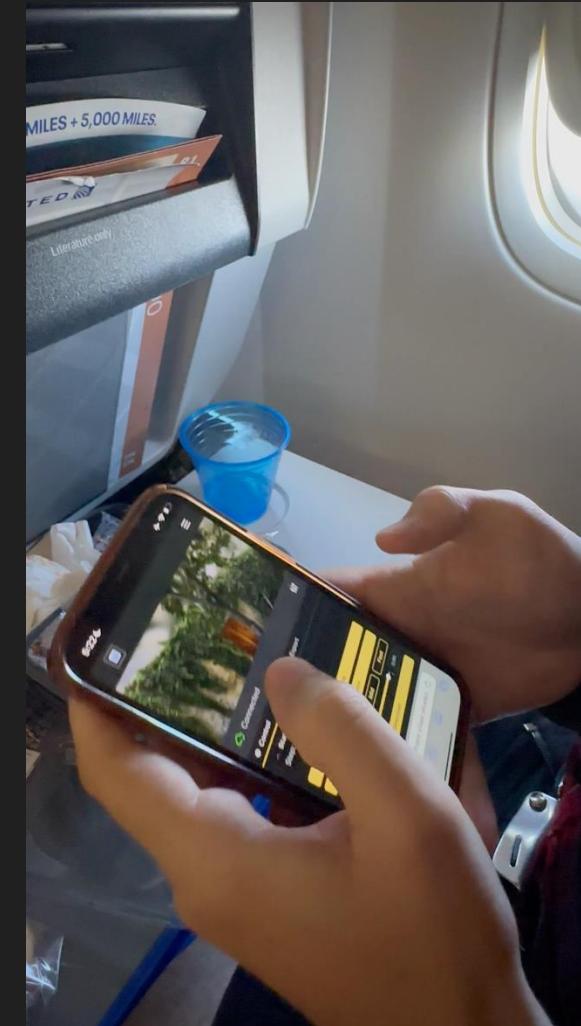
Coming Soon: Viser Integration

Python library for web-based 3D visualization

viser.studio



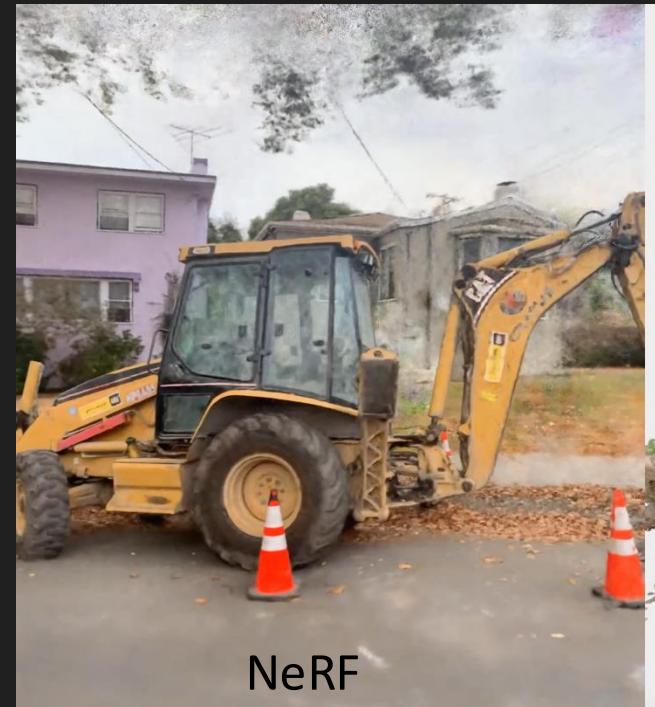
Shareable Links



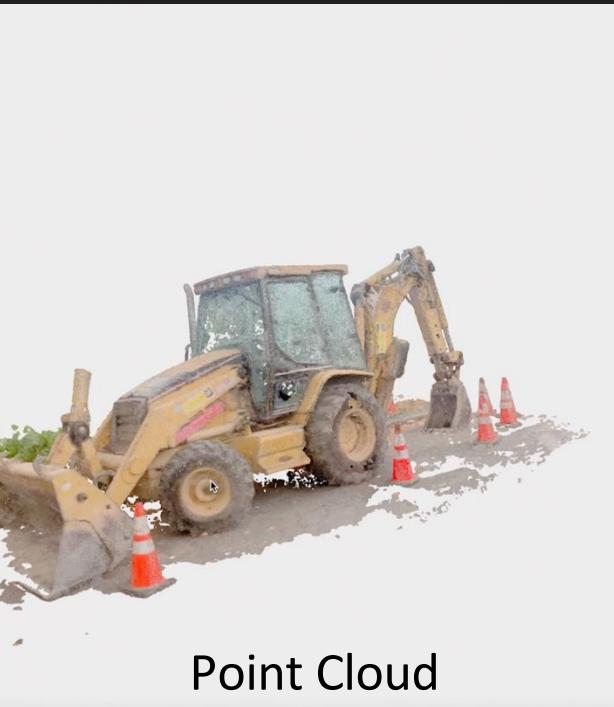
Mobile Support

Export Options

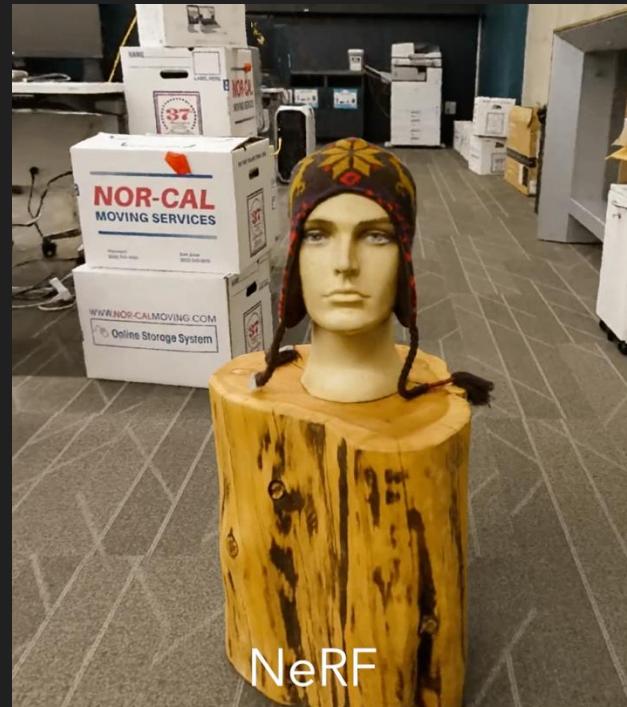
Geometry Conversion



NeRF



Point Cloud



NeRF



Mesh

Camera Effects



VFX: Blender Integration





nerf studio



Grade eterna



Grade eterna

